

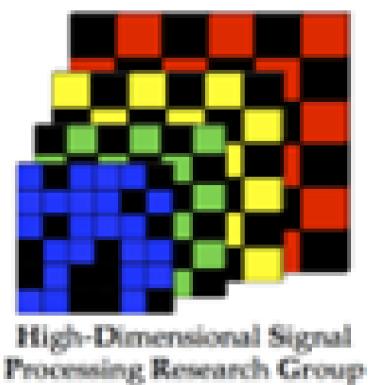
TALLER TEORICO PRACTICO EN ESTRUCTURAS DE DATOS Y APRENDIZAJE PROFUNDO



Fabian Perez - Semillero HDSP



Henry Mantilla - Semillero HDSP



¿Qué entendemos por la palabra
datos?

"Se pueden tener datos sin información pero no se puede tener información sin datos"

VISIÓN POR COMPUTADOR

Classification



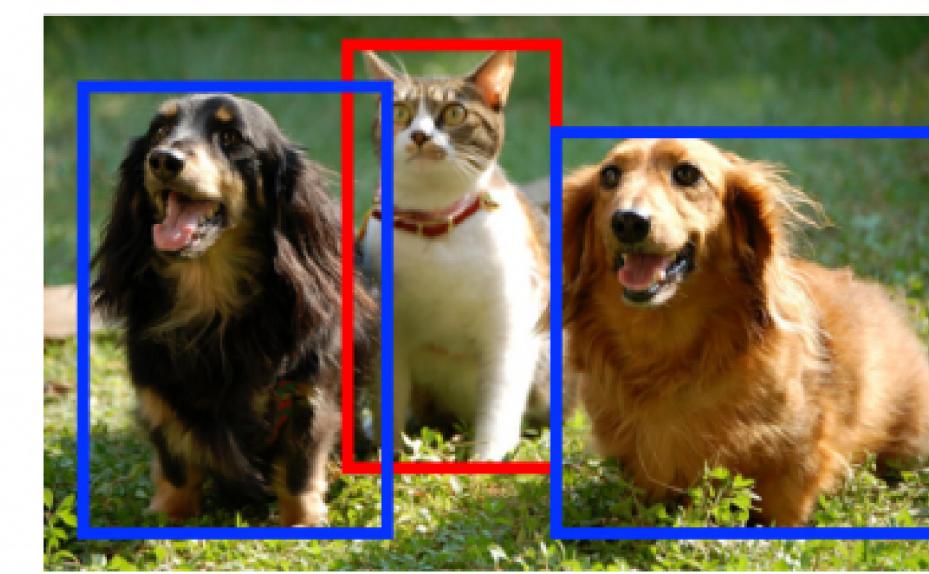
CAT

**Classification
+ Localization**



CAT

Object Detection



CAT, DOG

Instance Segmentation



CAT, DOG

Single object

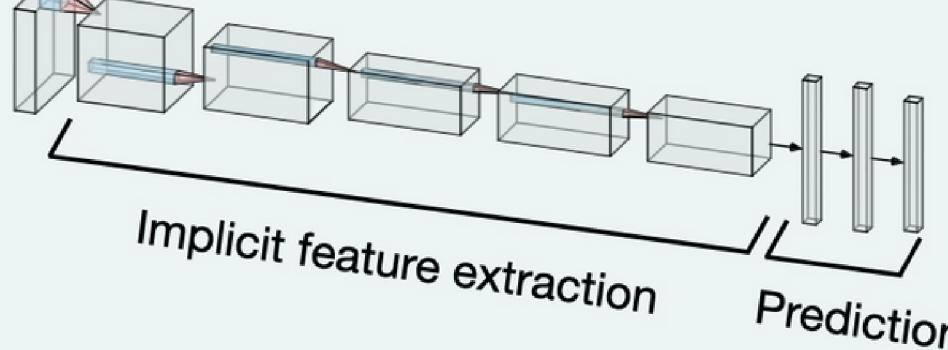
Multiple objects

Datos estructurados y no estructurados

Unstructured data



Deep learning



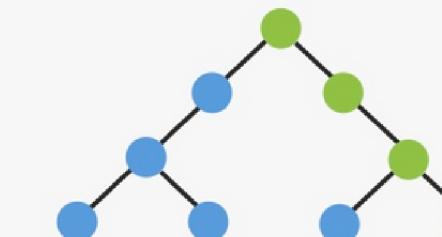
**“Manual”
feature extraction**



Structured data

	Sepal length	Sepal width	Petal length	Petal width
Flower 1	5.1	3.5	1.4	0.3
Flower 2	4.9	3.0	1.4	0.2
Flower 3	5.9	3.0	5.1	1.8
...

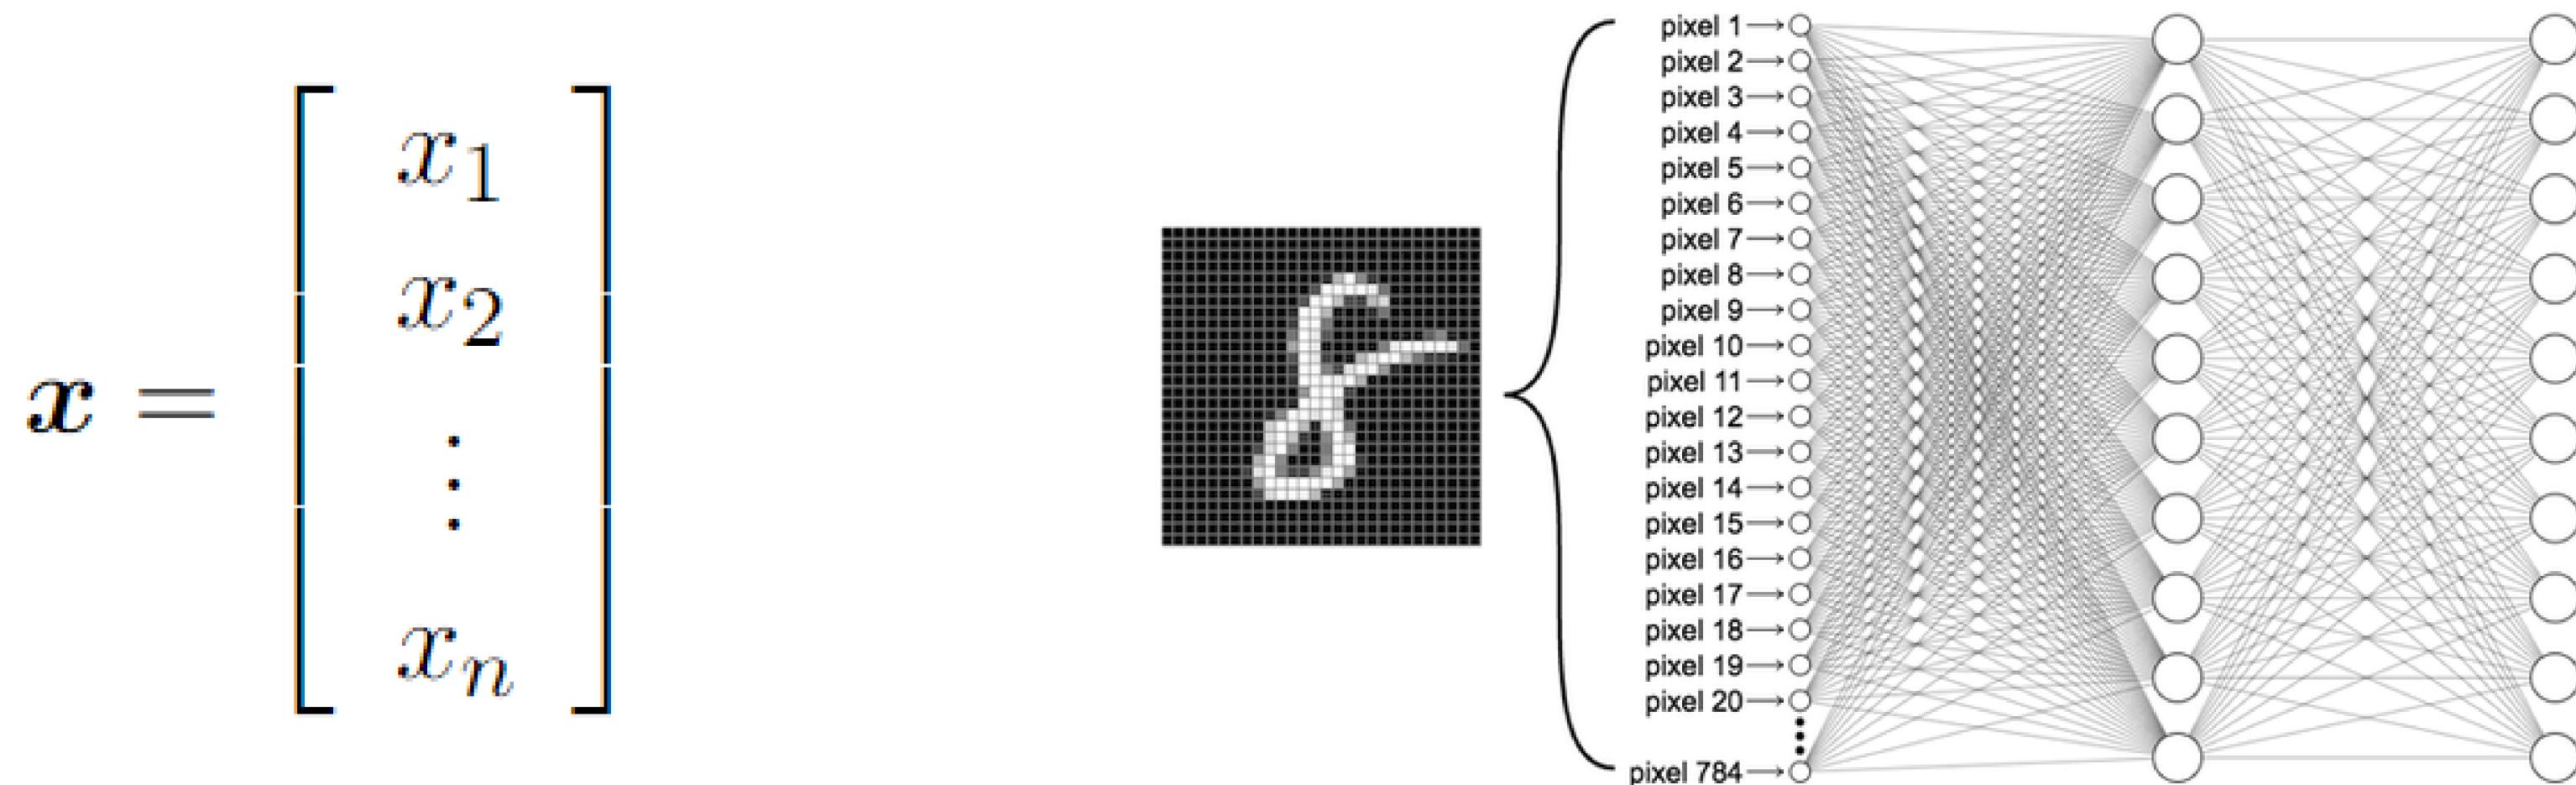
Conventional
machine
learning



Estructuras de datos usadas en aprendizaje profundo

Vectores

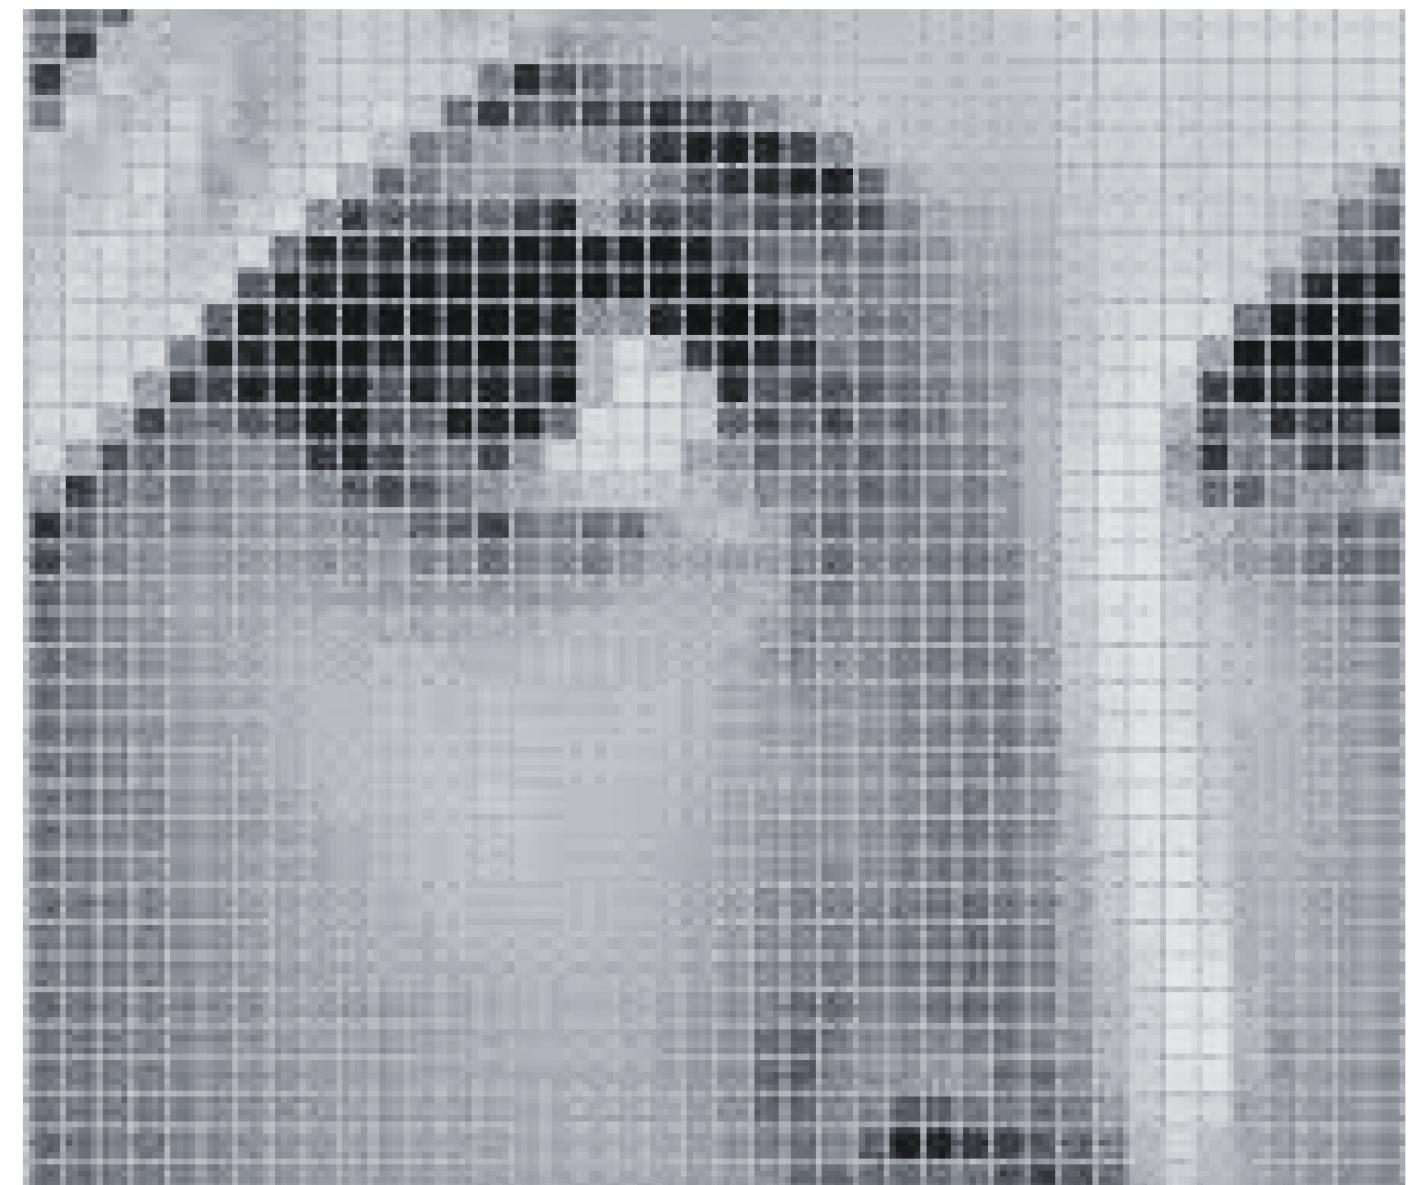
Un vector es un arreglo de valores numéricos, cada componente del vector es identificada mediante su indice. El numero de componentes corresponde a la dimensionalidad del vector.



Matrices

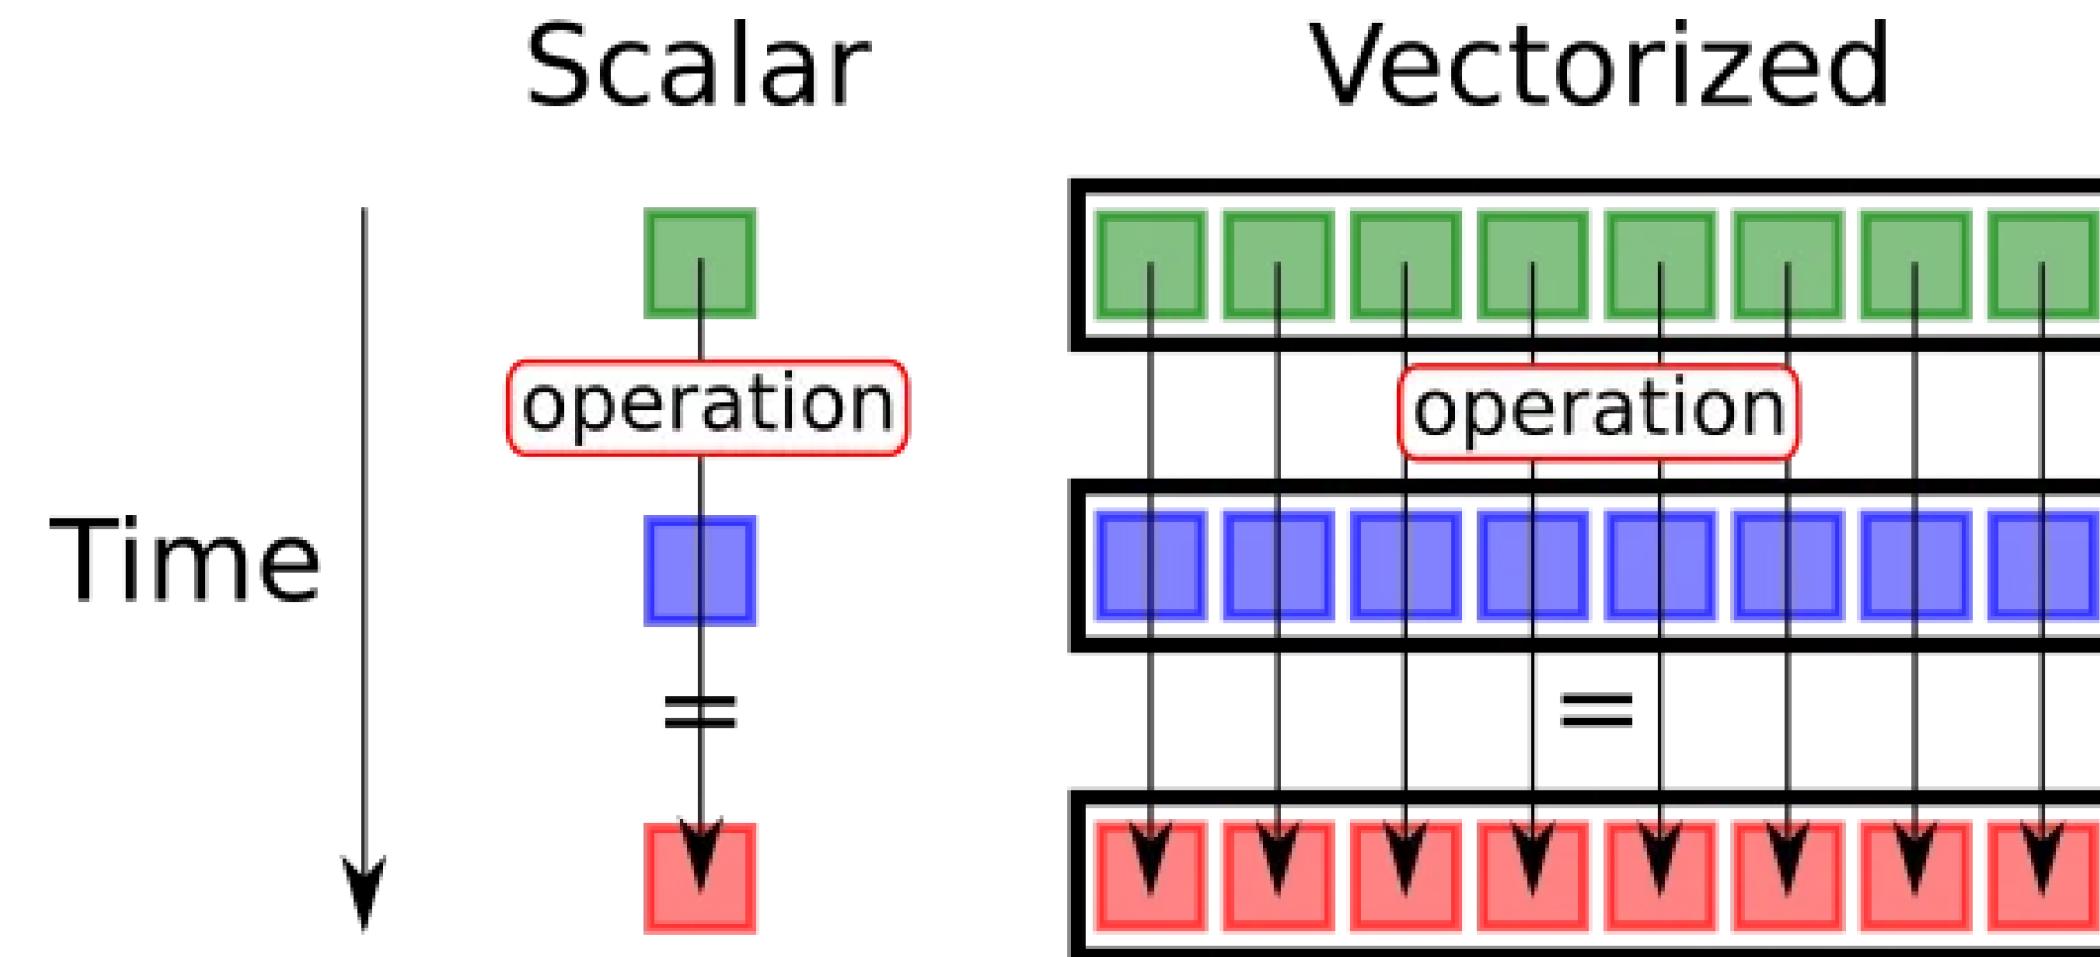
Una matriz es un arreglo bidimensional de valores numéricos, cada elemento es identificado mediante dos índices.

$$A = \begin{bmatrix} a_{11} & a_{12} & \cdots & a_{1n} \\ a_{21} & a_{22} & \cdots & a_{2n} \\ \vdots & \vdots & & \vdots \\ a_{m1} & a_{m2} & \cdots & a_{mn} \end{bmatrix}$$



Vectorización

La vectorización es el proceso de convertir un algoritmo que opera sobre un solo valor a la vez a operar en un conjunto de valores a la vez.





```
import time
w = np.random.rand(1000000)
x = np.random.rand(1000000)
tic = time.time()
c = np.dot(w,x)

toc = time.time()
print('Vectorized version: '+ str(1000*(toc-tic))+'ms')
c = 0

tic = time.time()
for i in range(1000000):
    c +=w[i]*x[i]

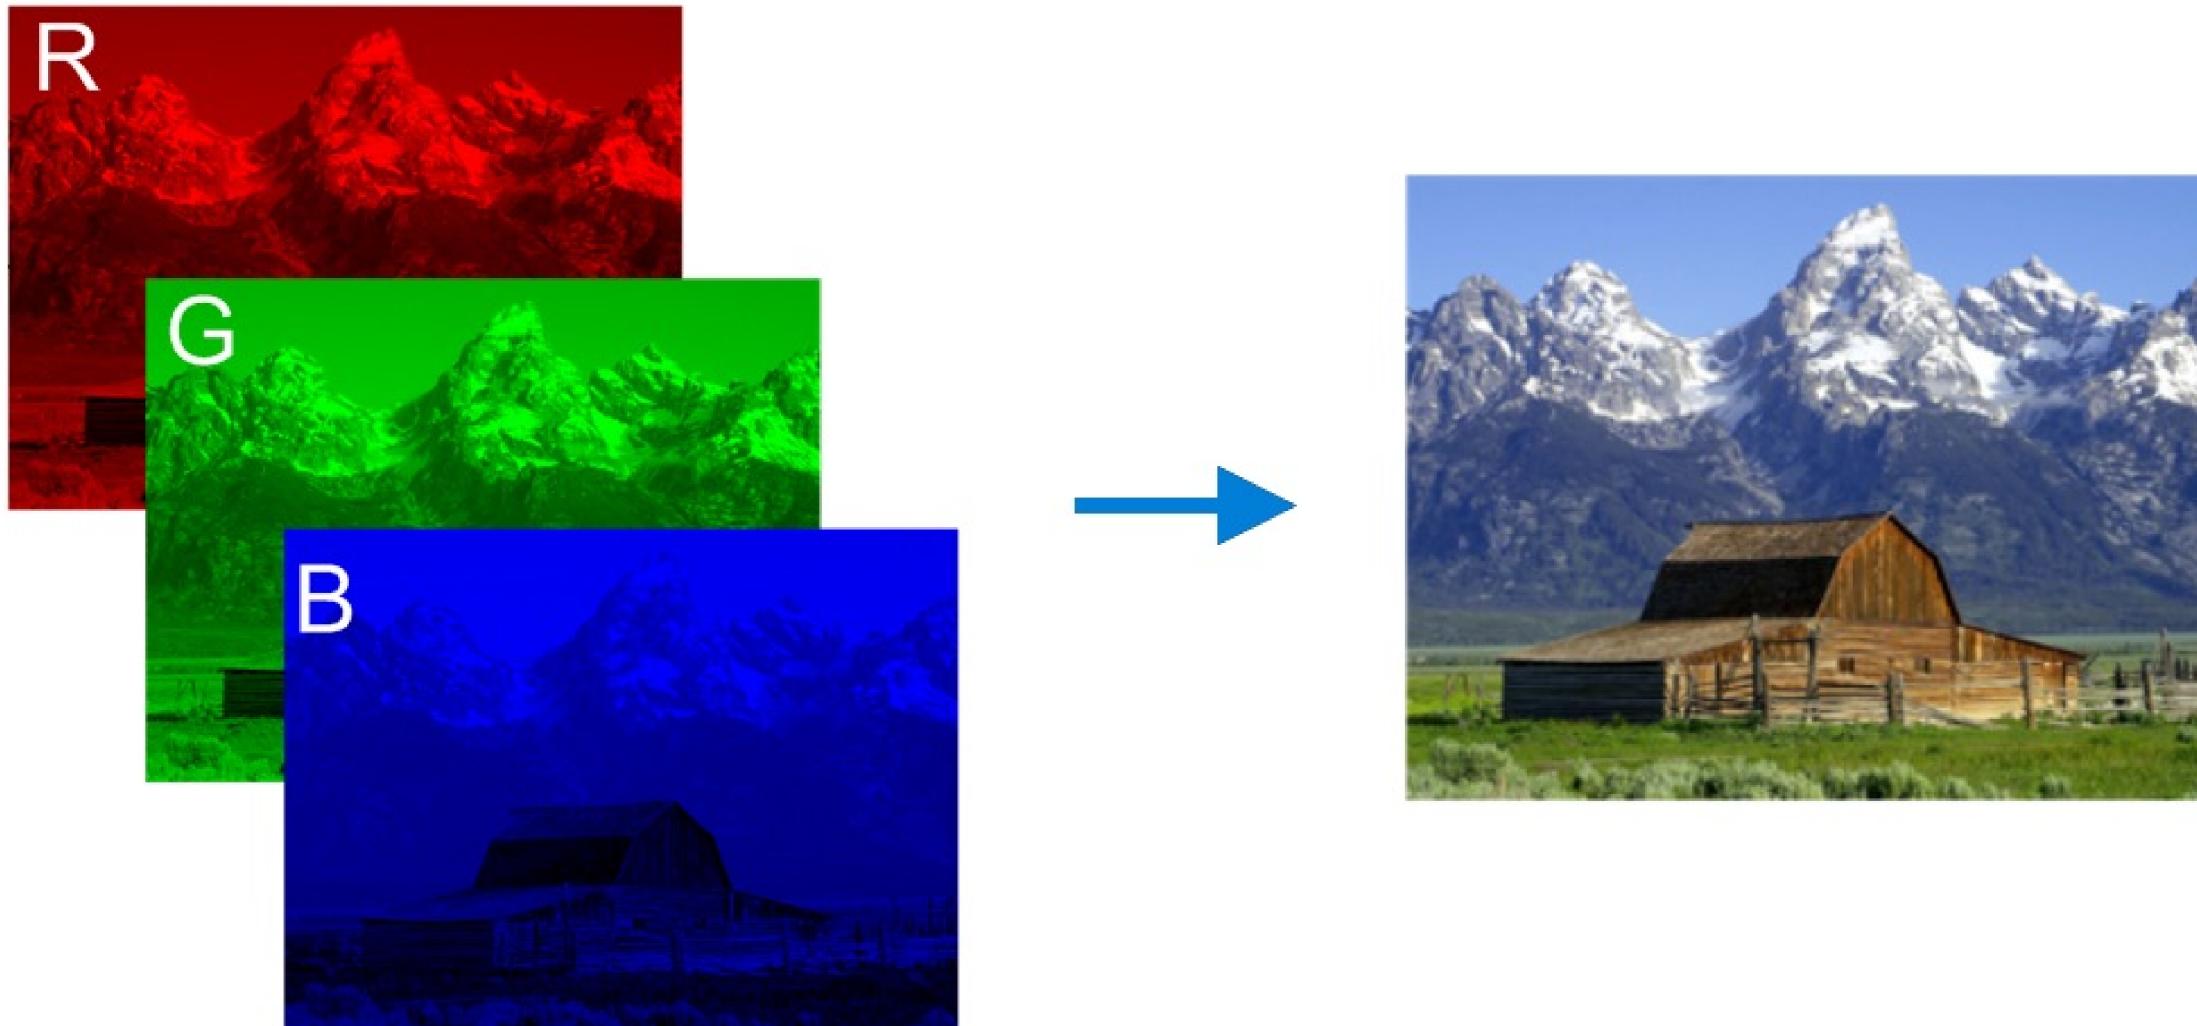
toc = time.time()
print("For loop: ",str(1000*(toc-tic))+'ms')
```

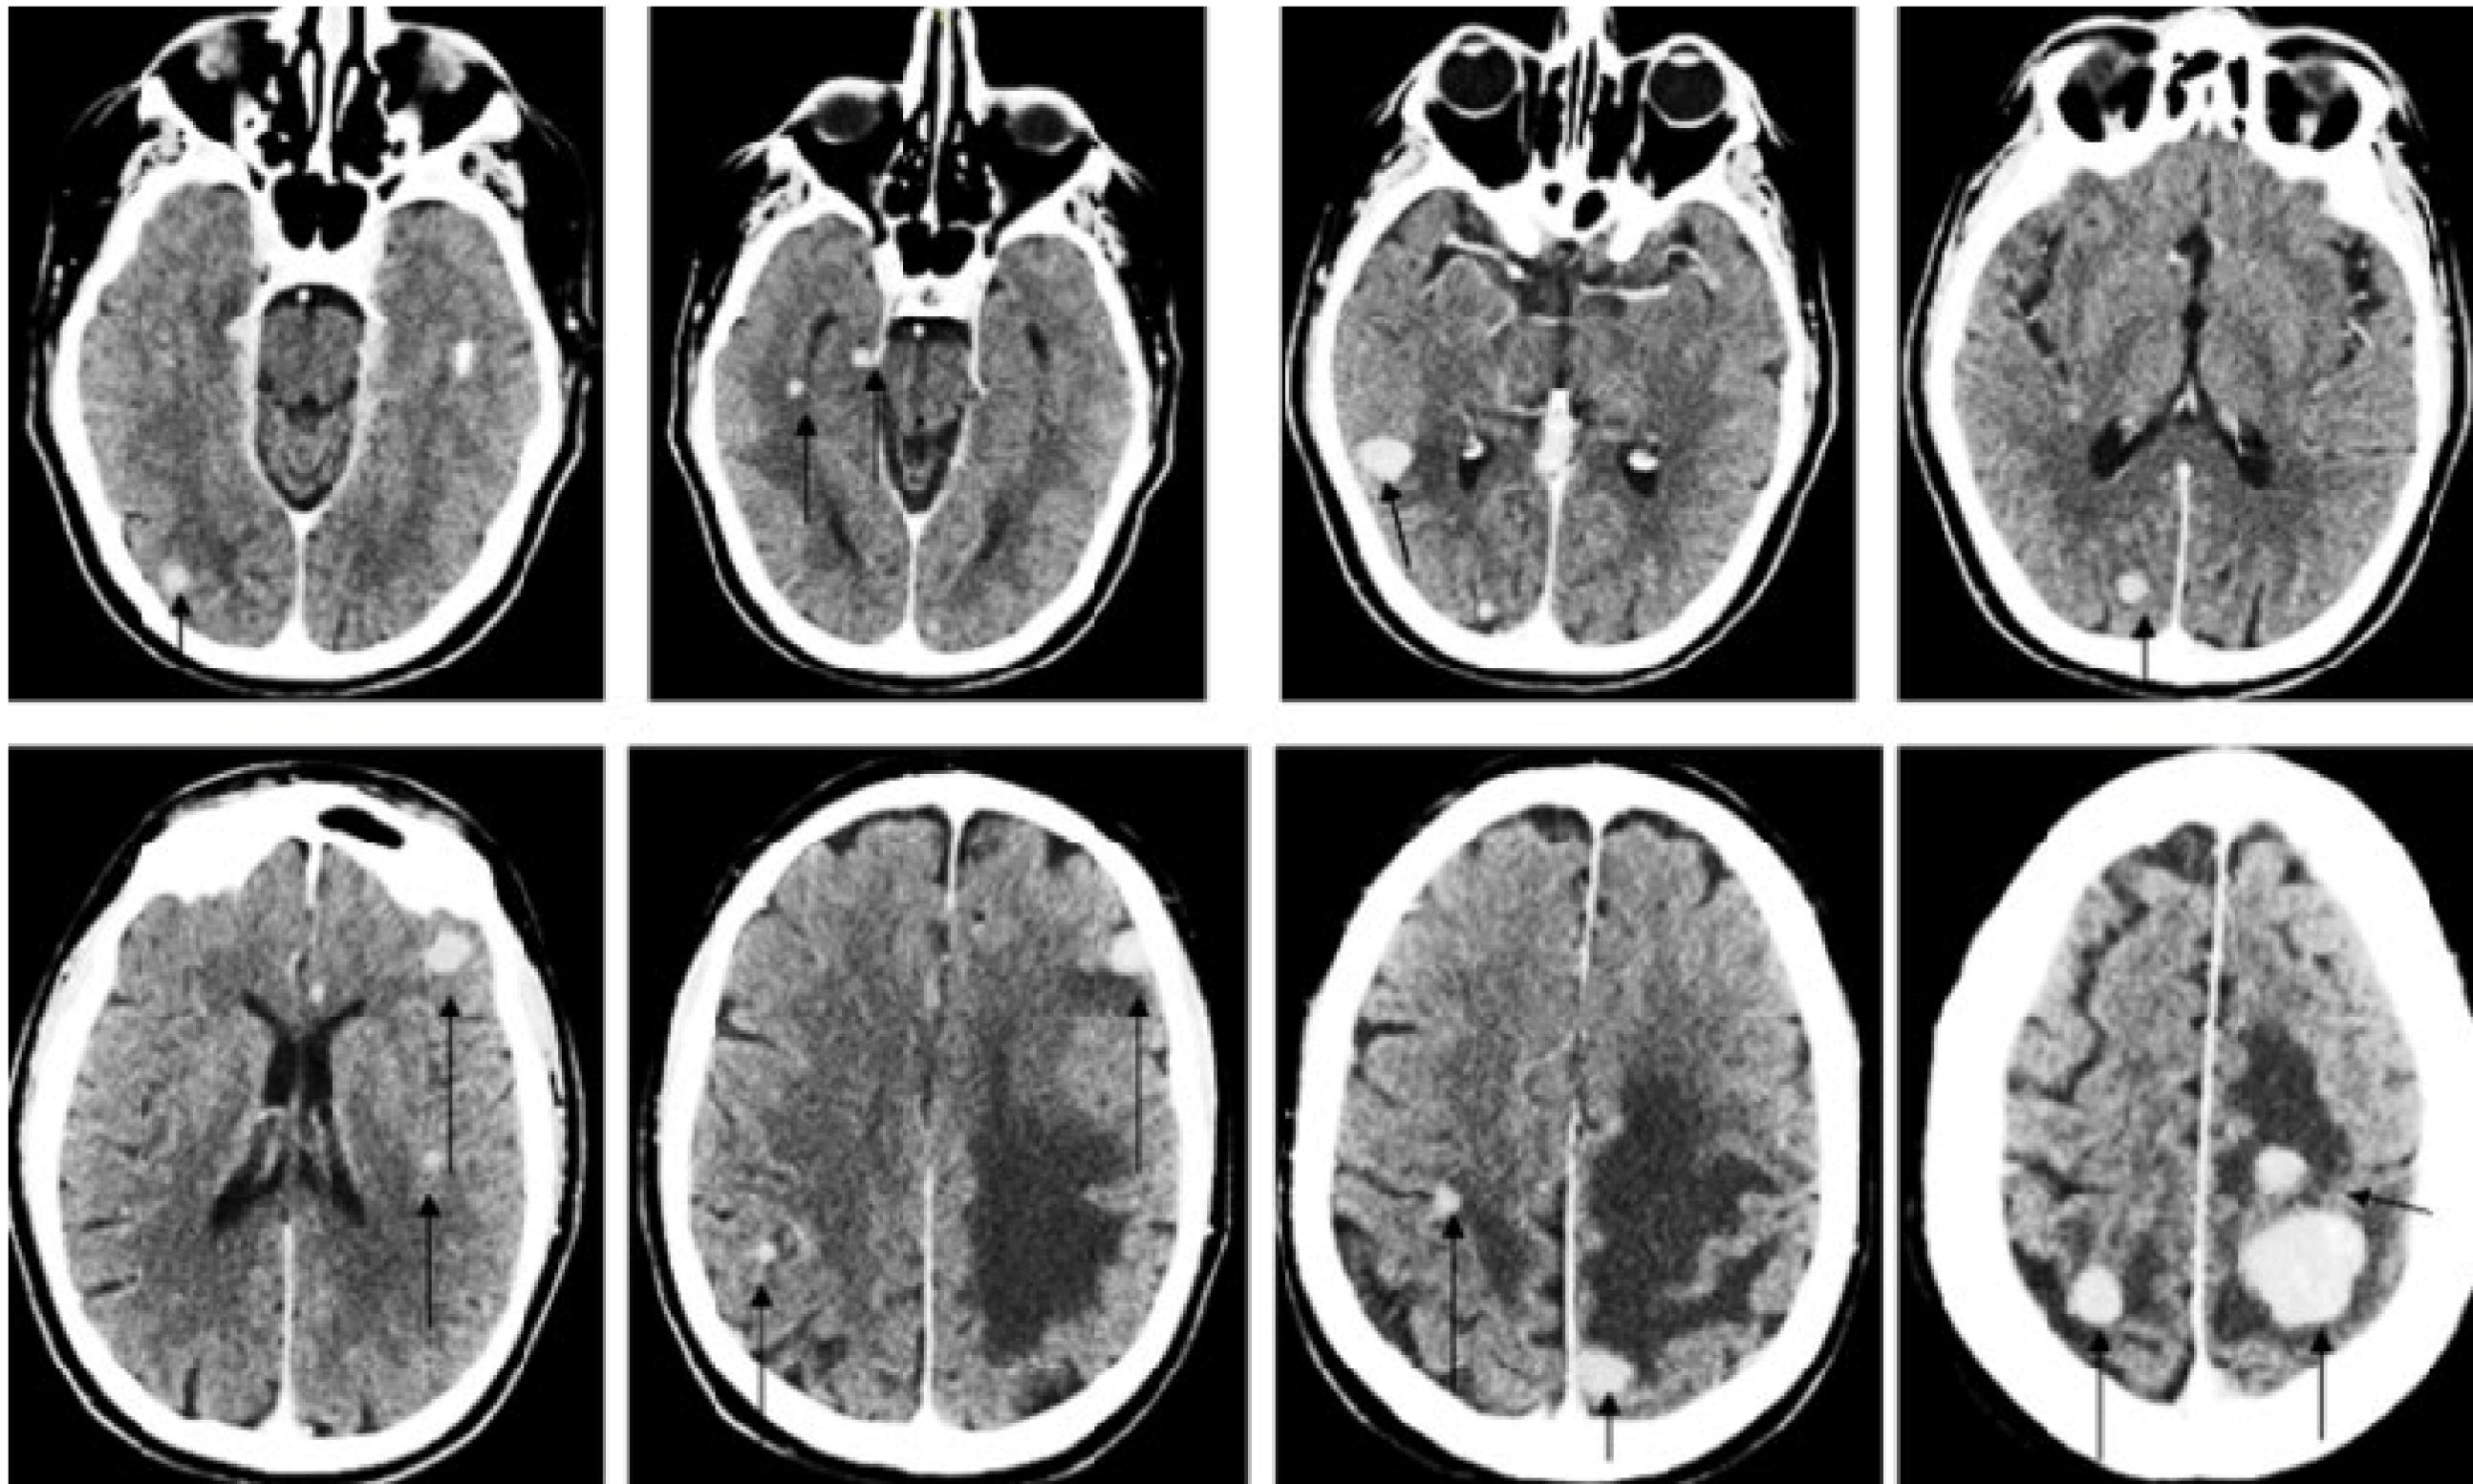
Tensores

En aprendizaje profundo los tensores se emplean como estructura de datos, estos permiten generalizar matrices de más de dos dimensiones, podemos verlos como arreglos multidimensionales.

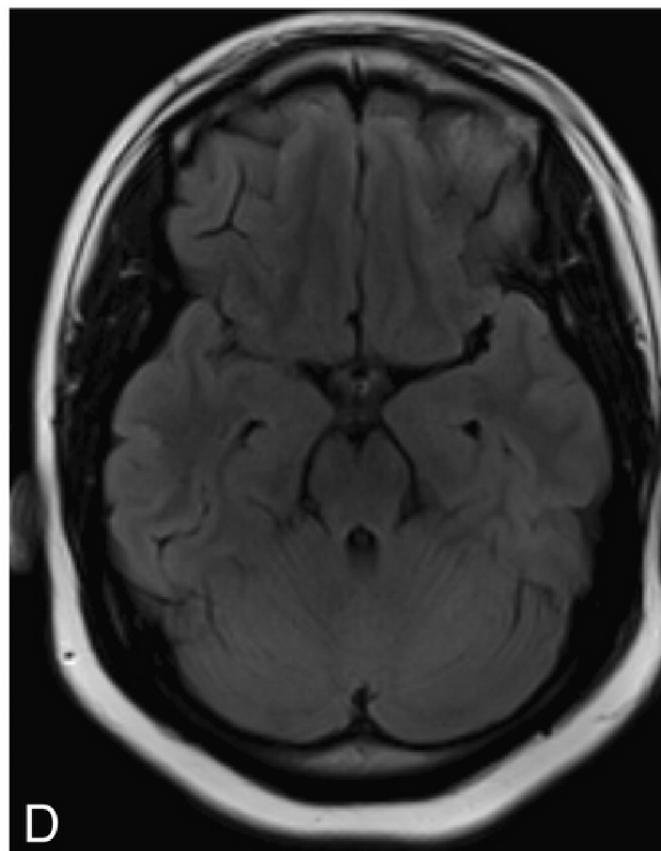
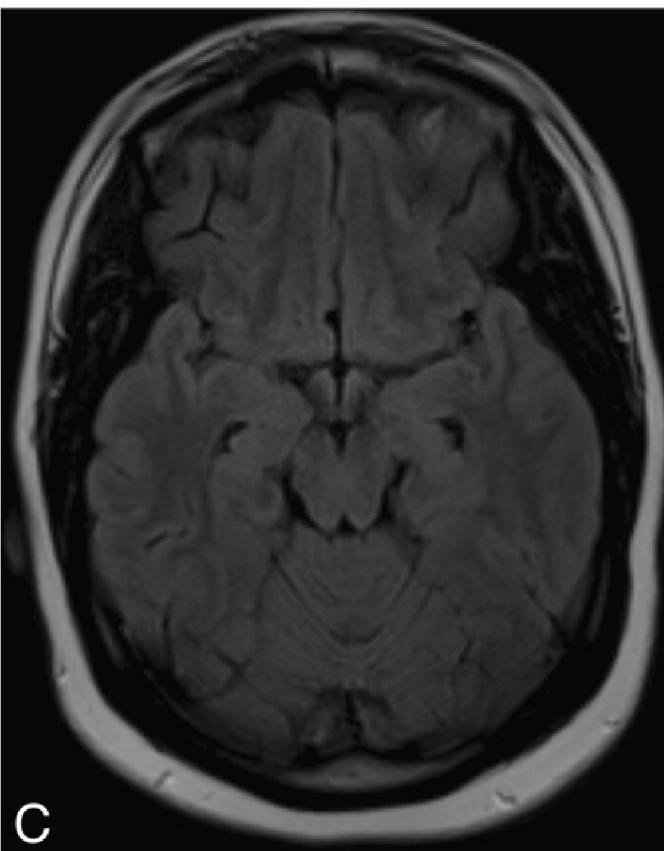
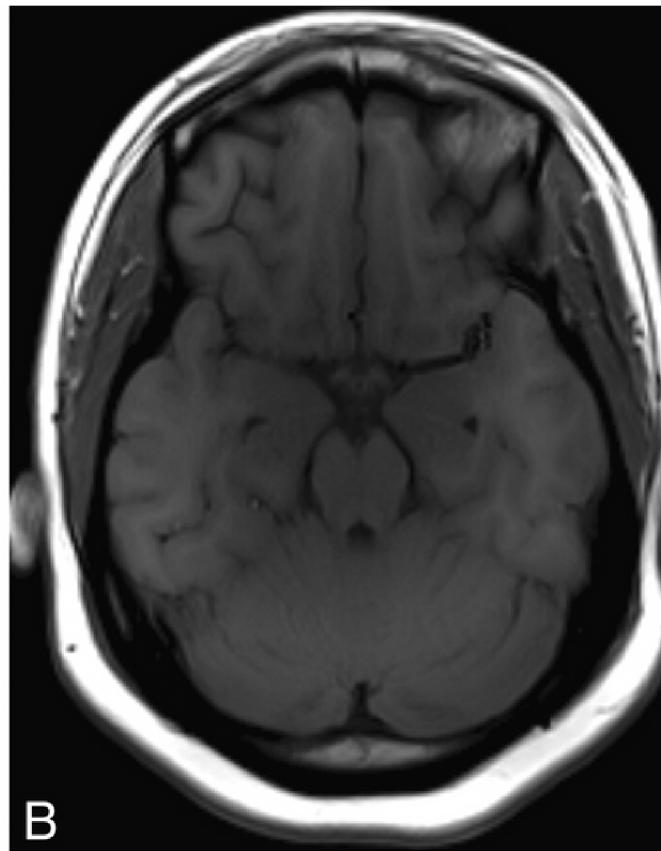
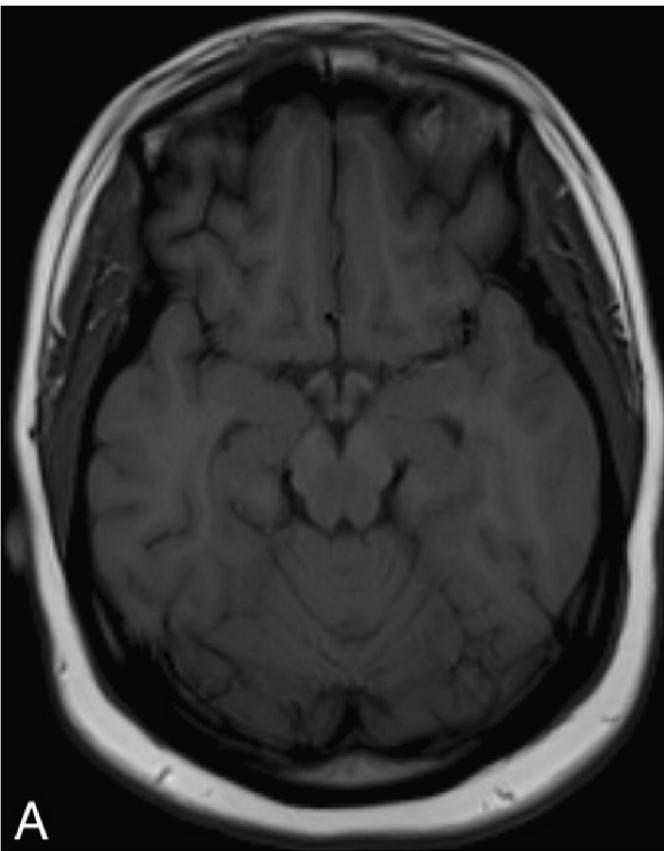


Ejemplos de tensores de rango 3





Computerized Tomography



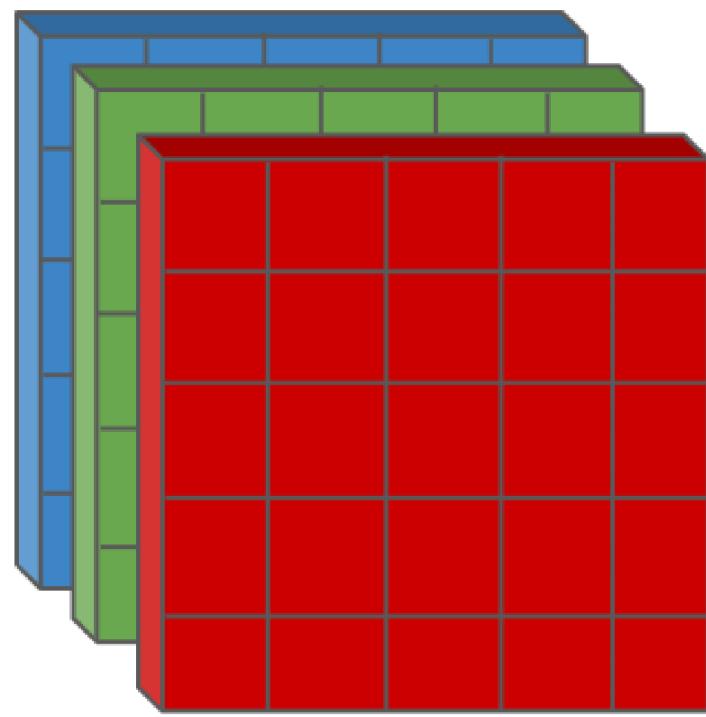
Magnetic Resonance Imaging



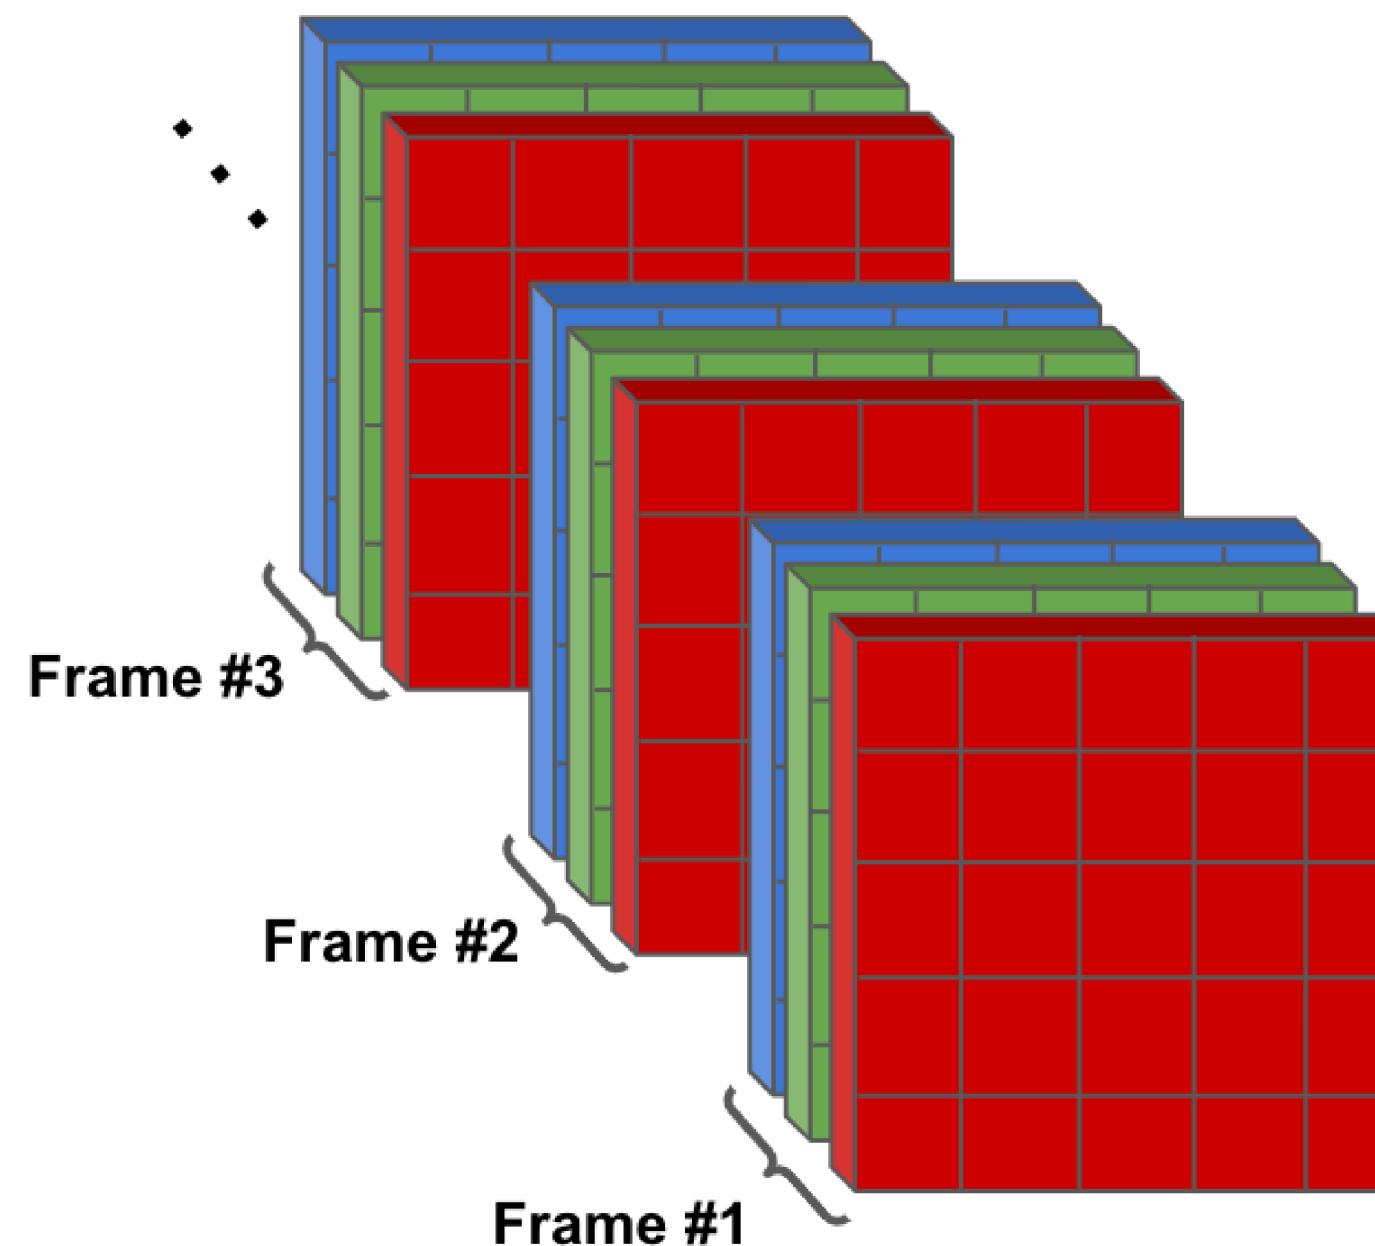
X-ray Imaging

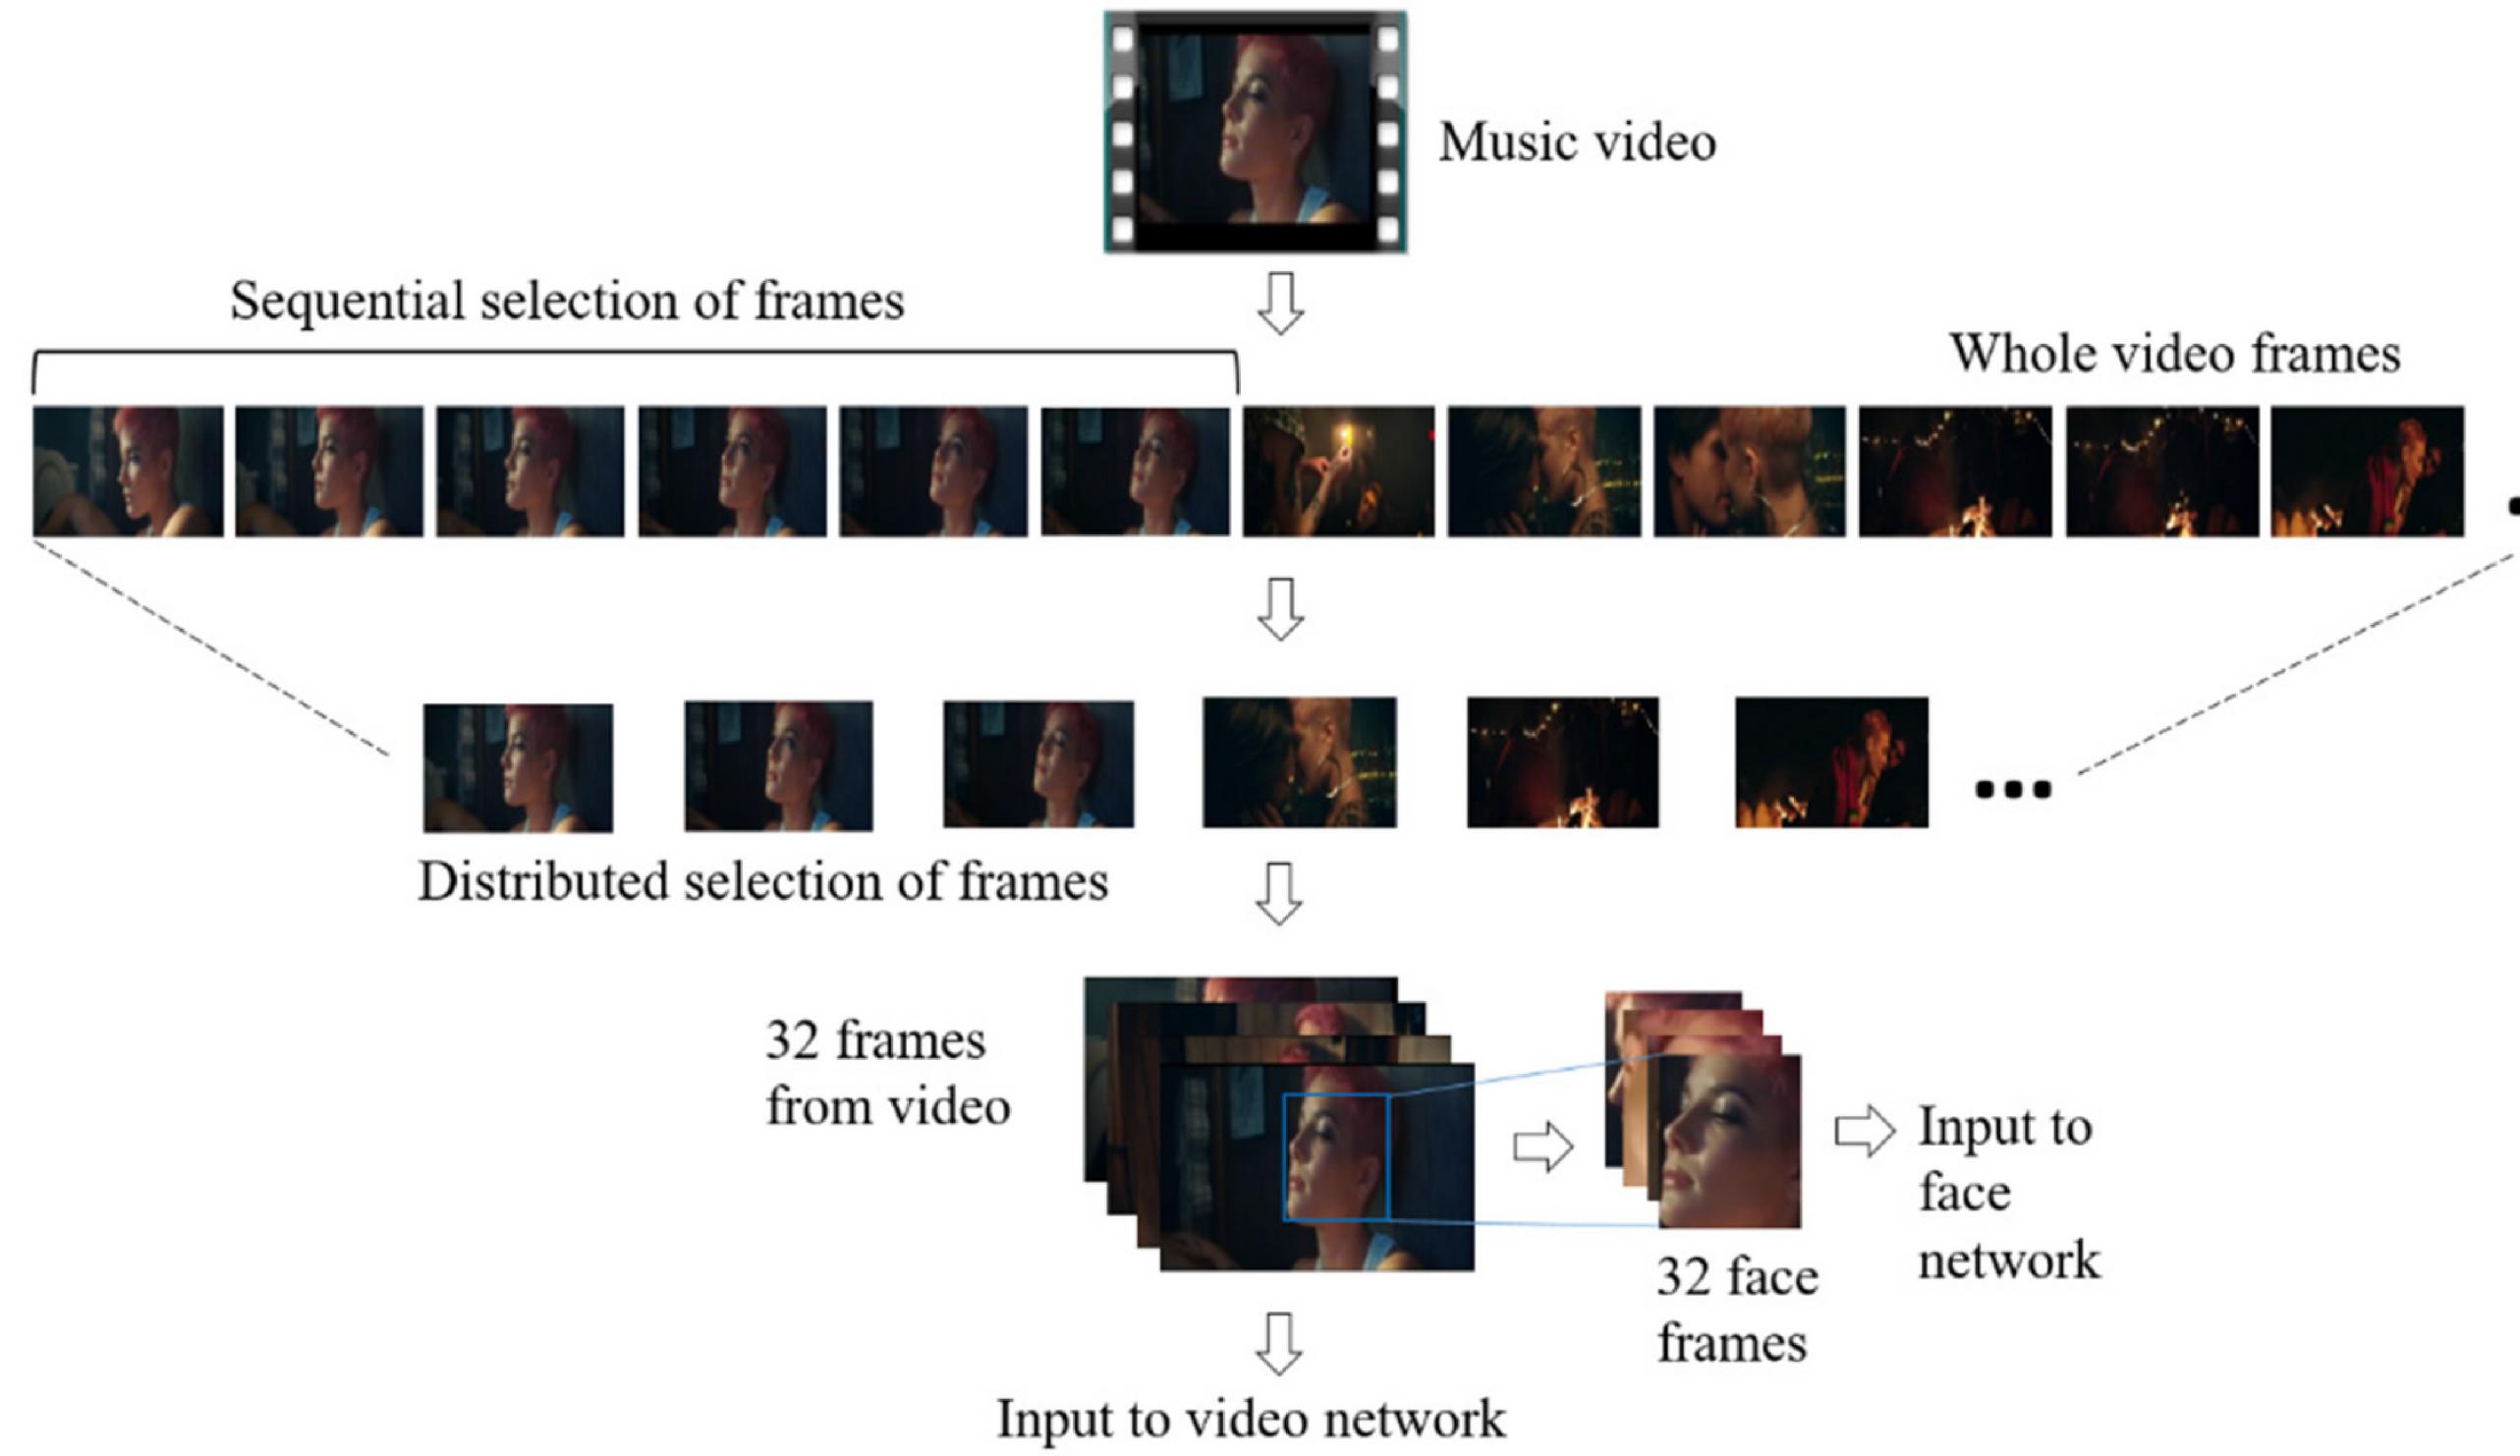
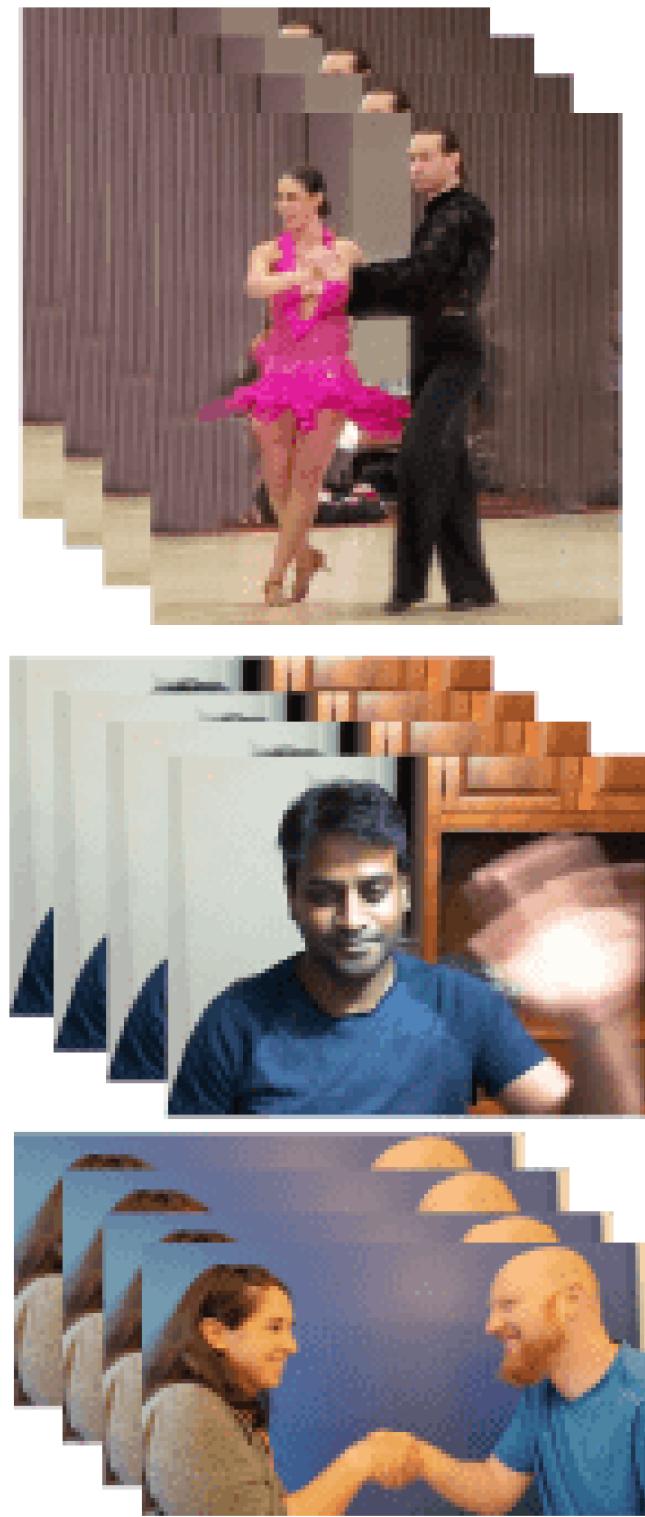
Ejemplos de tensores 4D

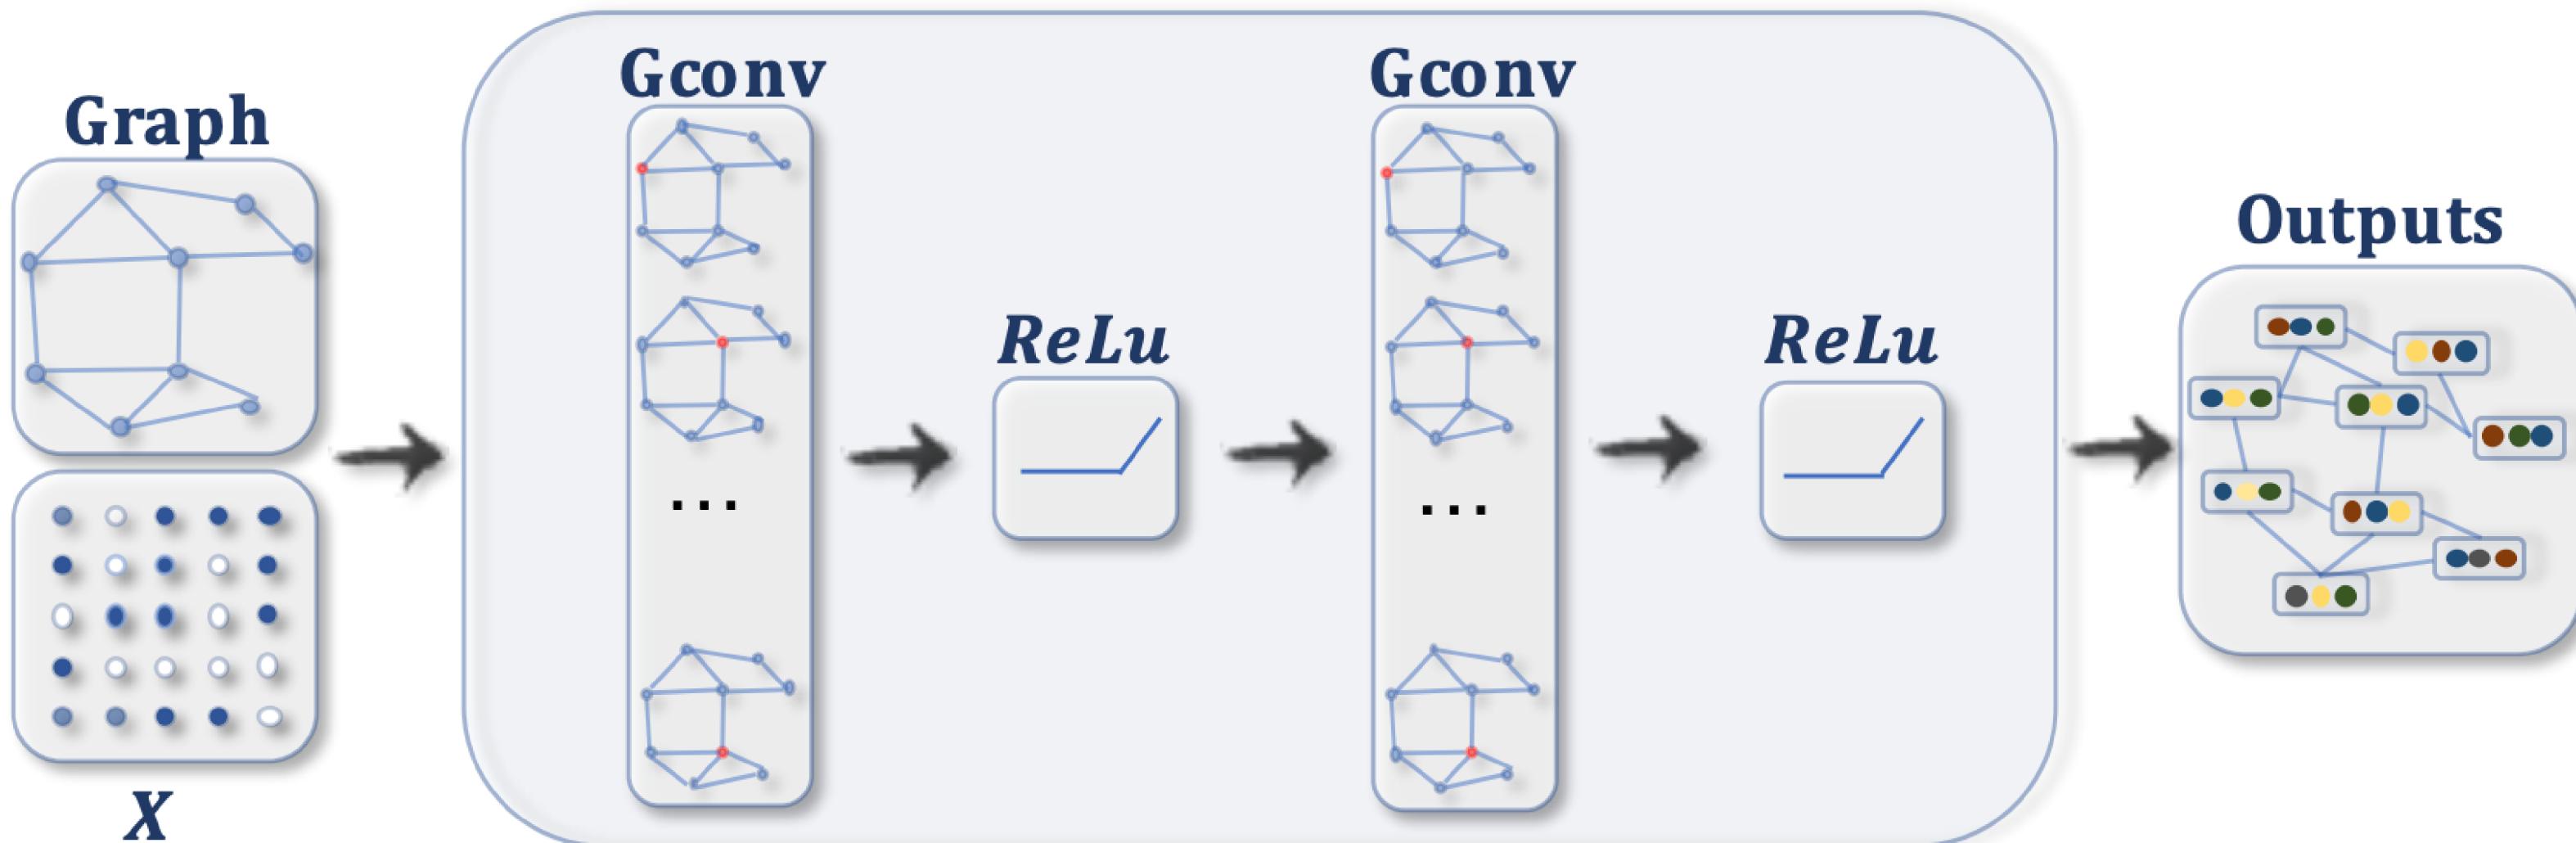
Single Image



Video Clip

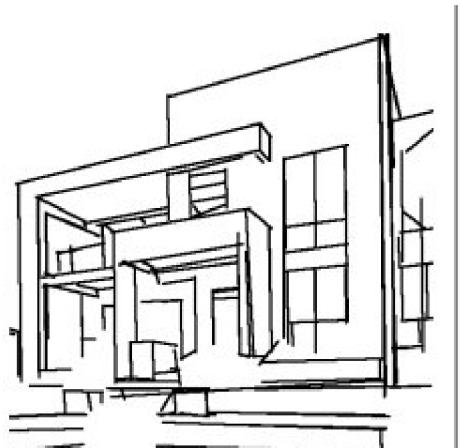






Deep Learning

no limits



"a modern house with windows"



Crea una base de datos llamada "Tienda en línea" que contenga las siguientes tablas:

"Productos": con los campos "ID", "Nombre", "Descripción", "Precio", "Stock".

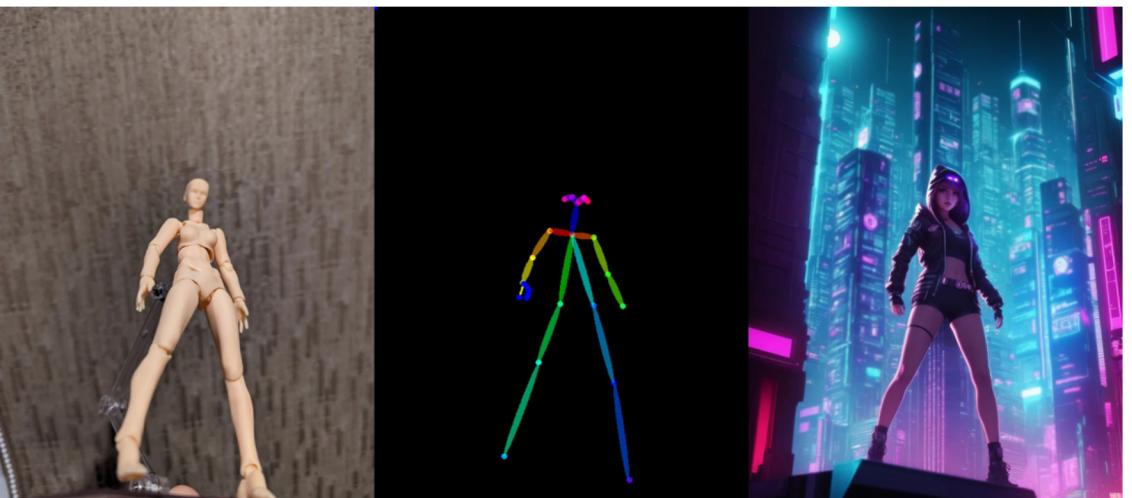
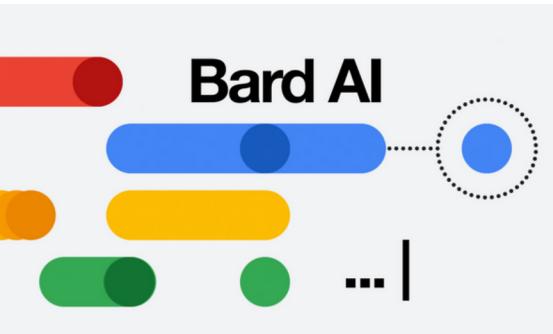
"Clientes": con los campos "ID", "Nombre", "Apellido", "Dirección", "Correo electrónico", "Contraseña".

"Pedidos": con los campos "ID", "Fecha", "ID del cliente", "Total del pedido".

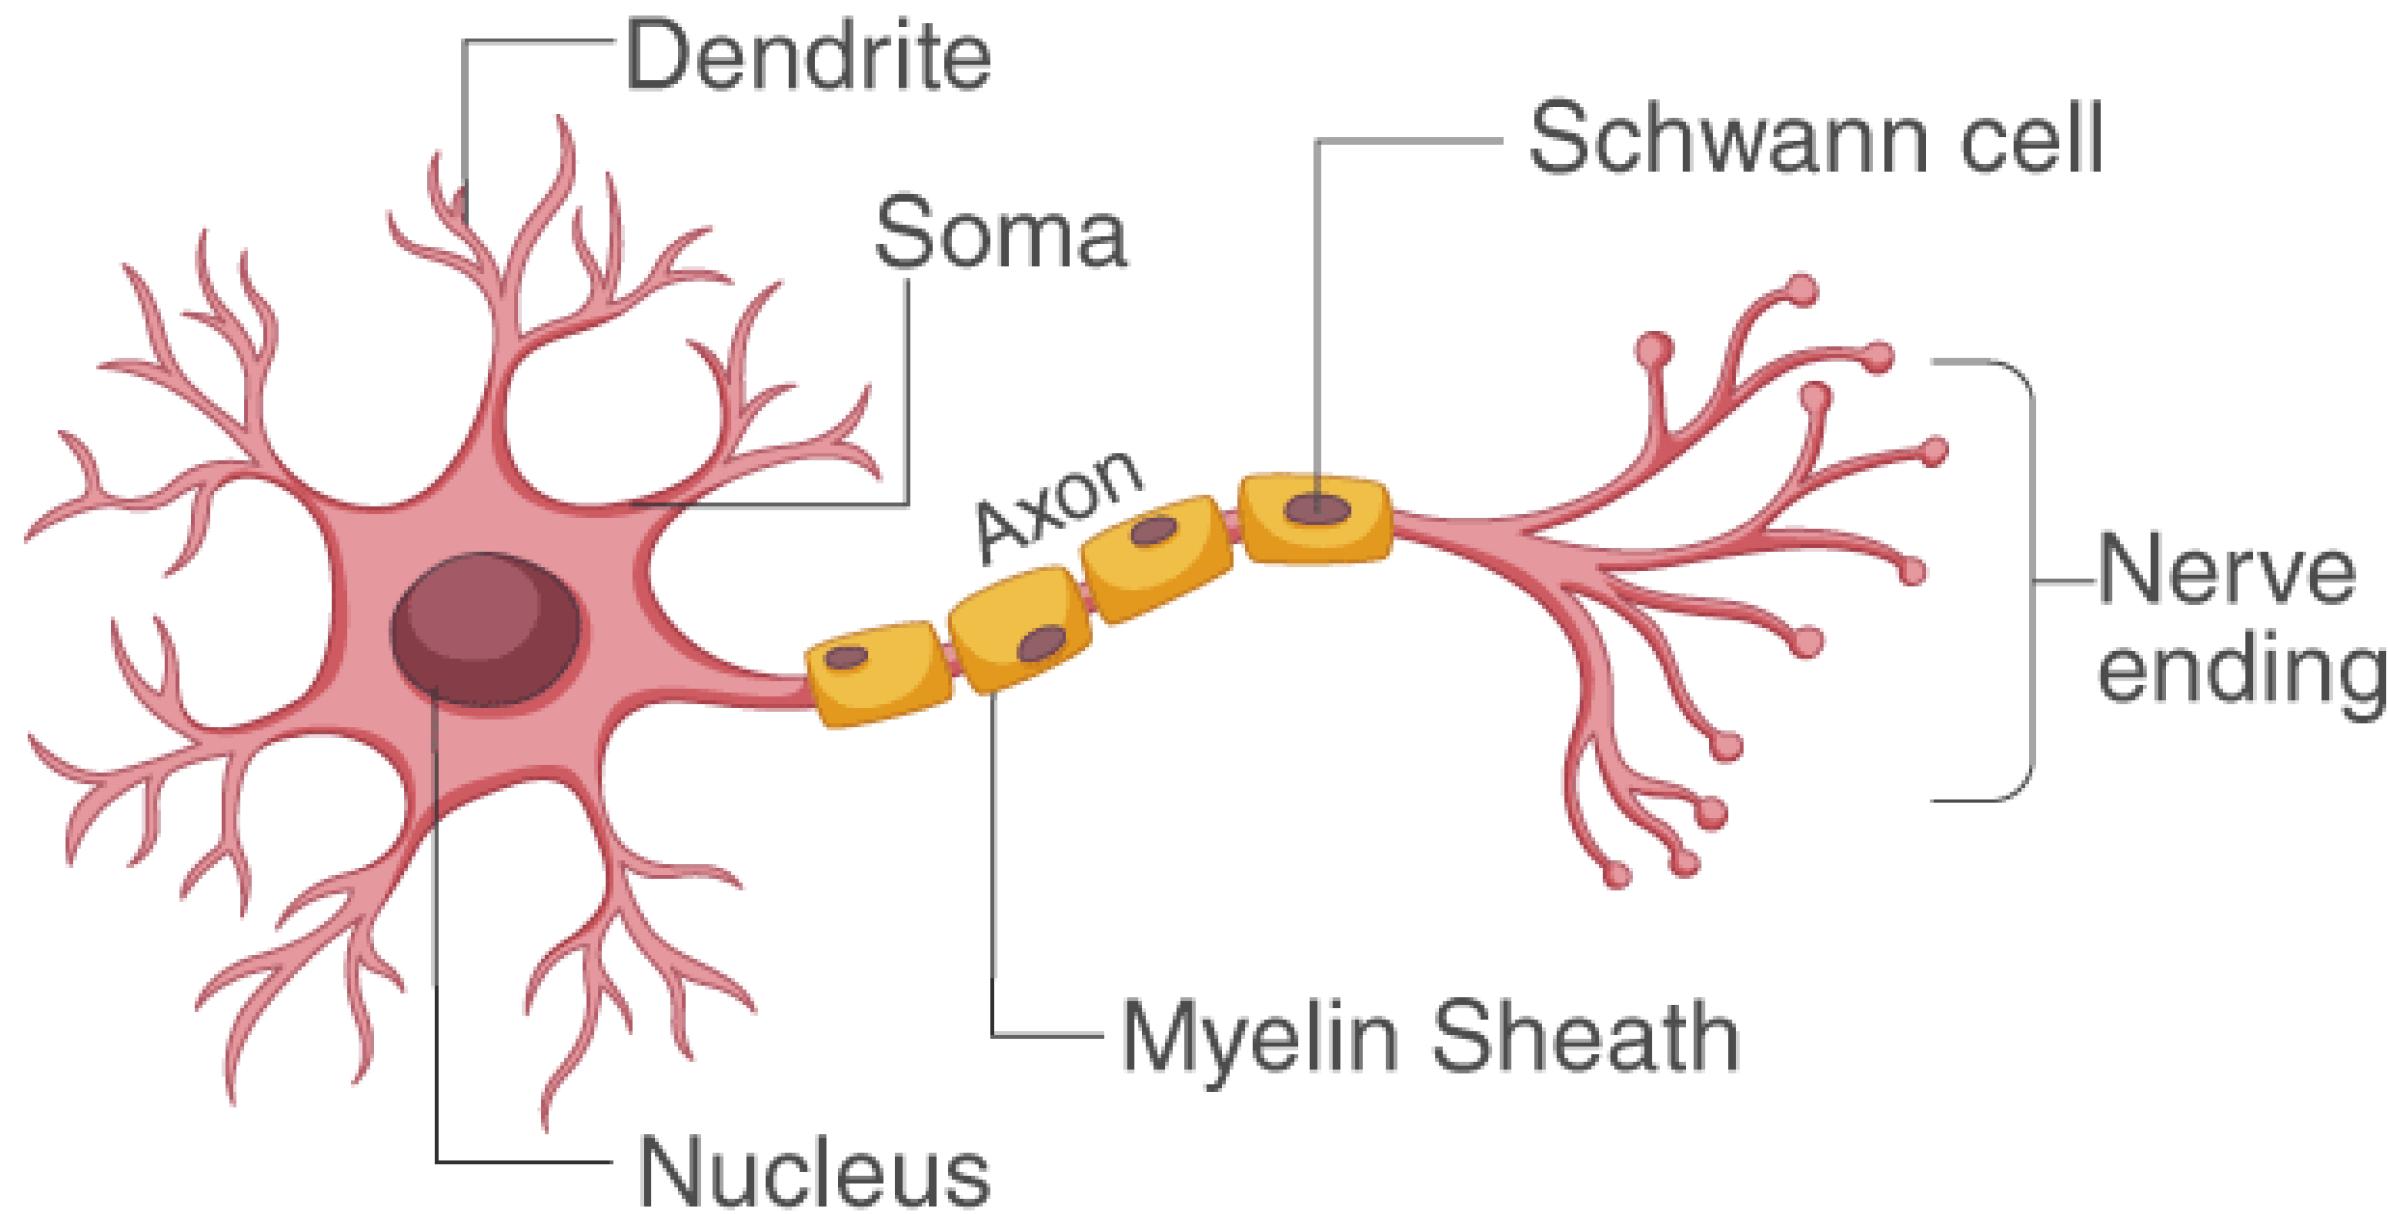
Aquí tienes la creación de la base de datos "Tienda en línea" con las tablas "Productos", "Clientes" y "Pedidos":

```
sql Copy code
CREATE DATABASE `Tienda en línea`;
USE `Tienda en línea`;

CREATE TABLE `Productos` (
  `ID` INT(11) NOT NULL AUTO_INCREMENT,
  `Nombre` VARCHAR(255) NOT NULL,
  `Descripción` TEXT NOT NULL,
  `Precio` DECIMAL(10,2) NOT NULL
```



STRUCTURE OF NEURON

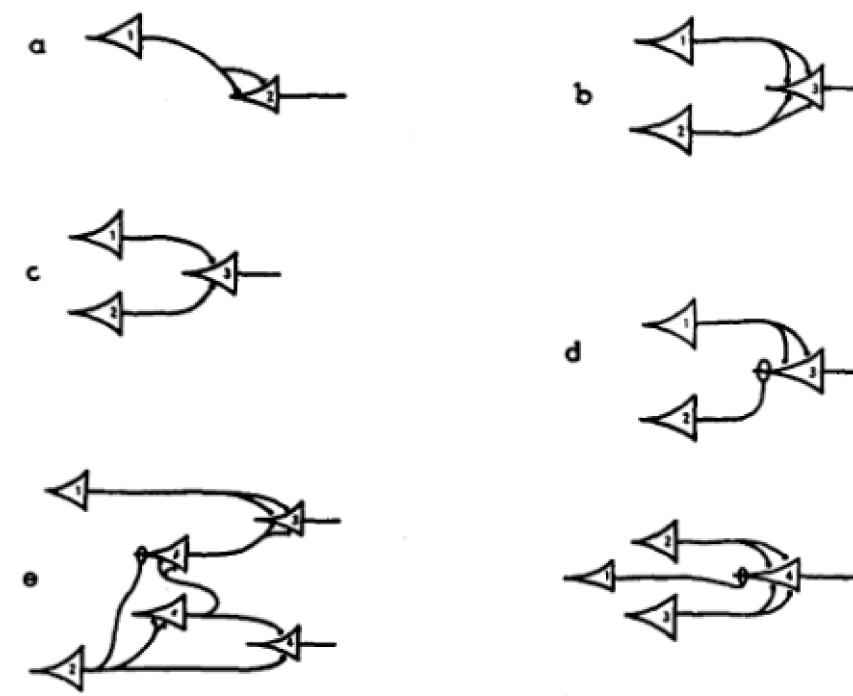


Artificial neural networks (ANN) are inspired
in biological neural networks

A LOGICAL CALCULUS OF THE IDEAS IMMANENT IN NERVOUS ACTIVITY

WARREN S. McCULLOCH and WALTER H. PITTS

originally published in: Bulletin of Mathematical Biophysics, vol. 5, 1943, p. 115-133



Artificial neural networks (ANN) are inspired
in biological neural networks

Diagram of a perceptron

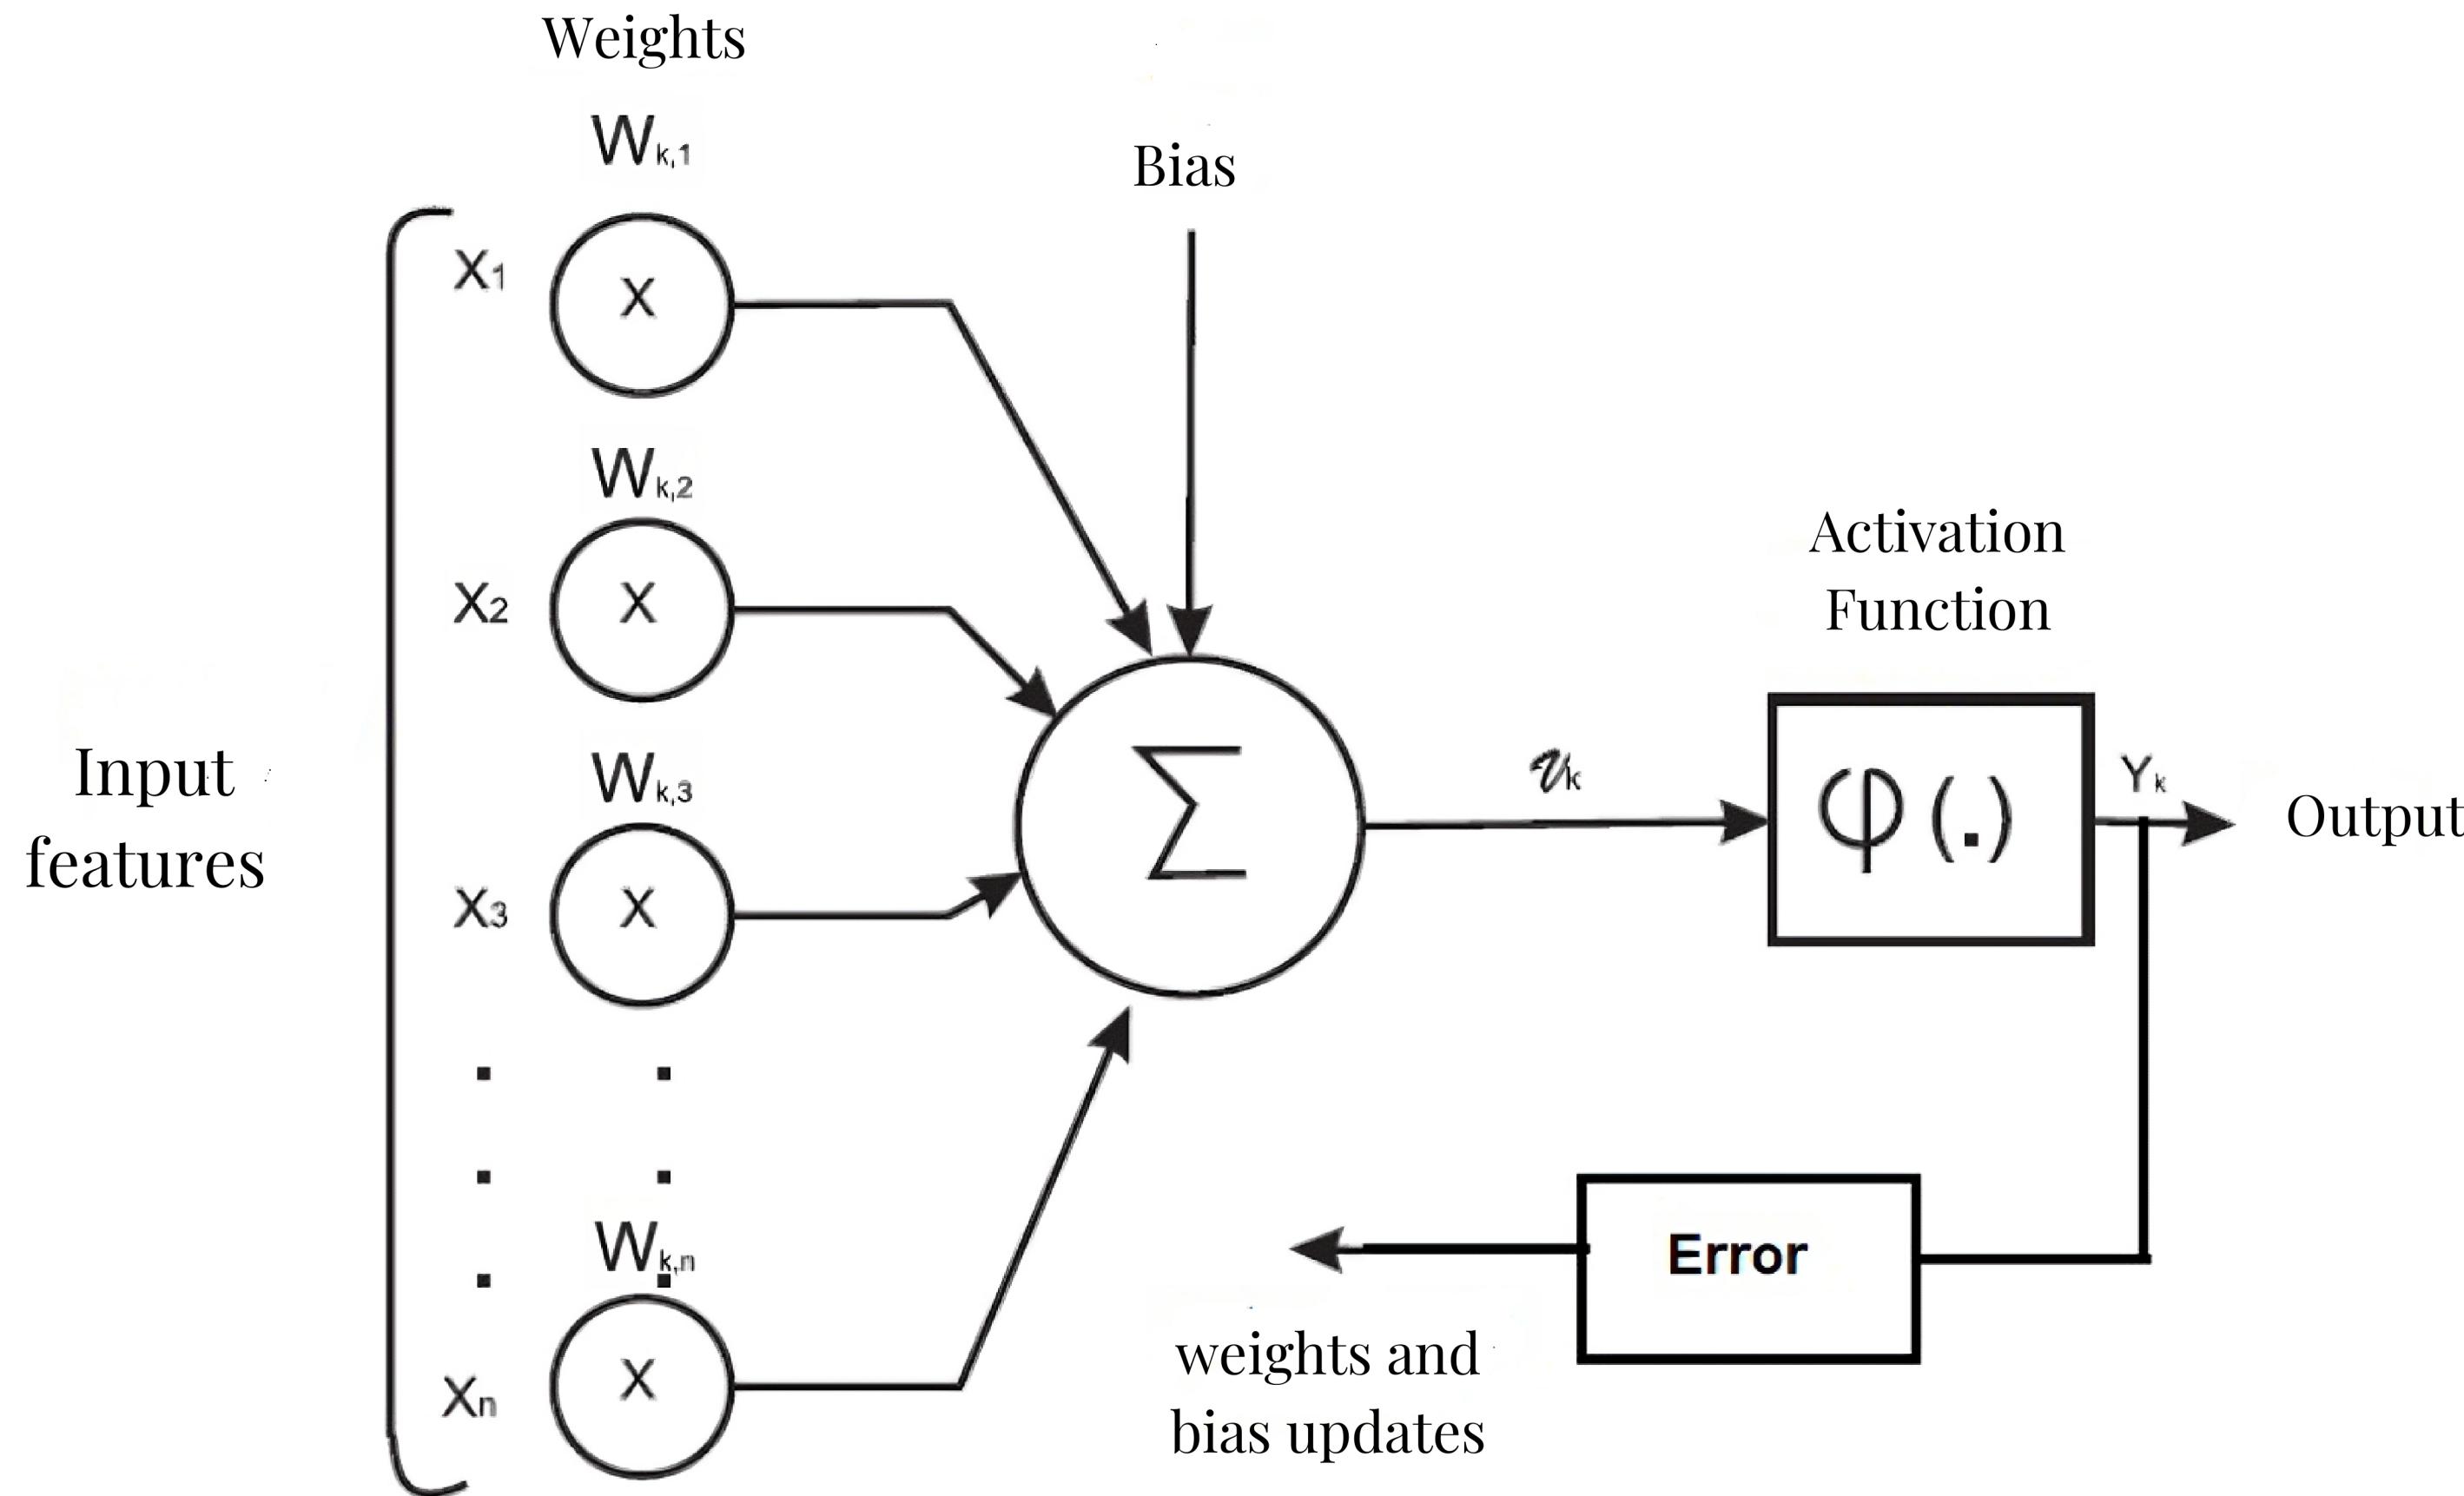
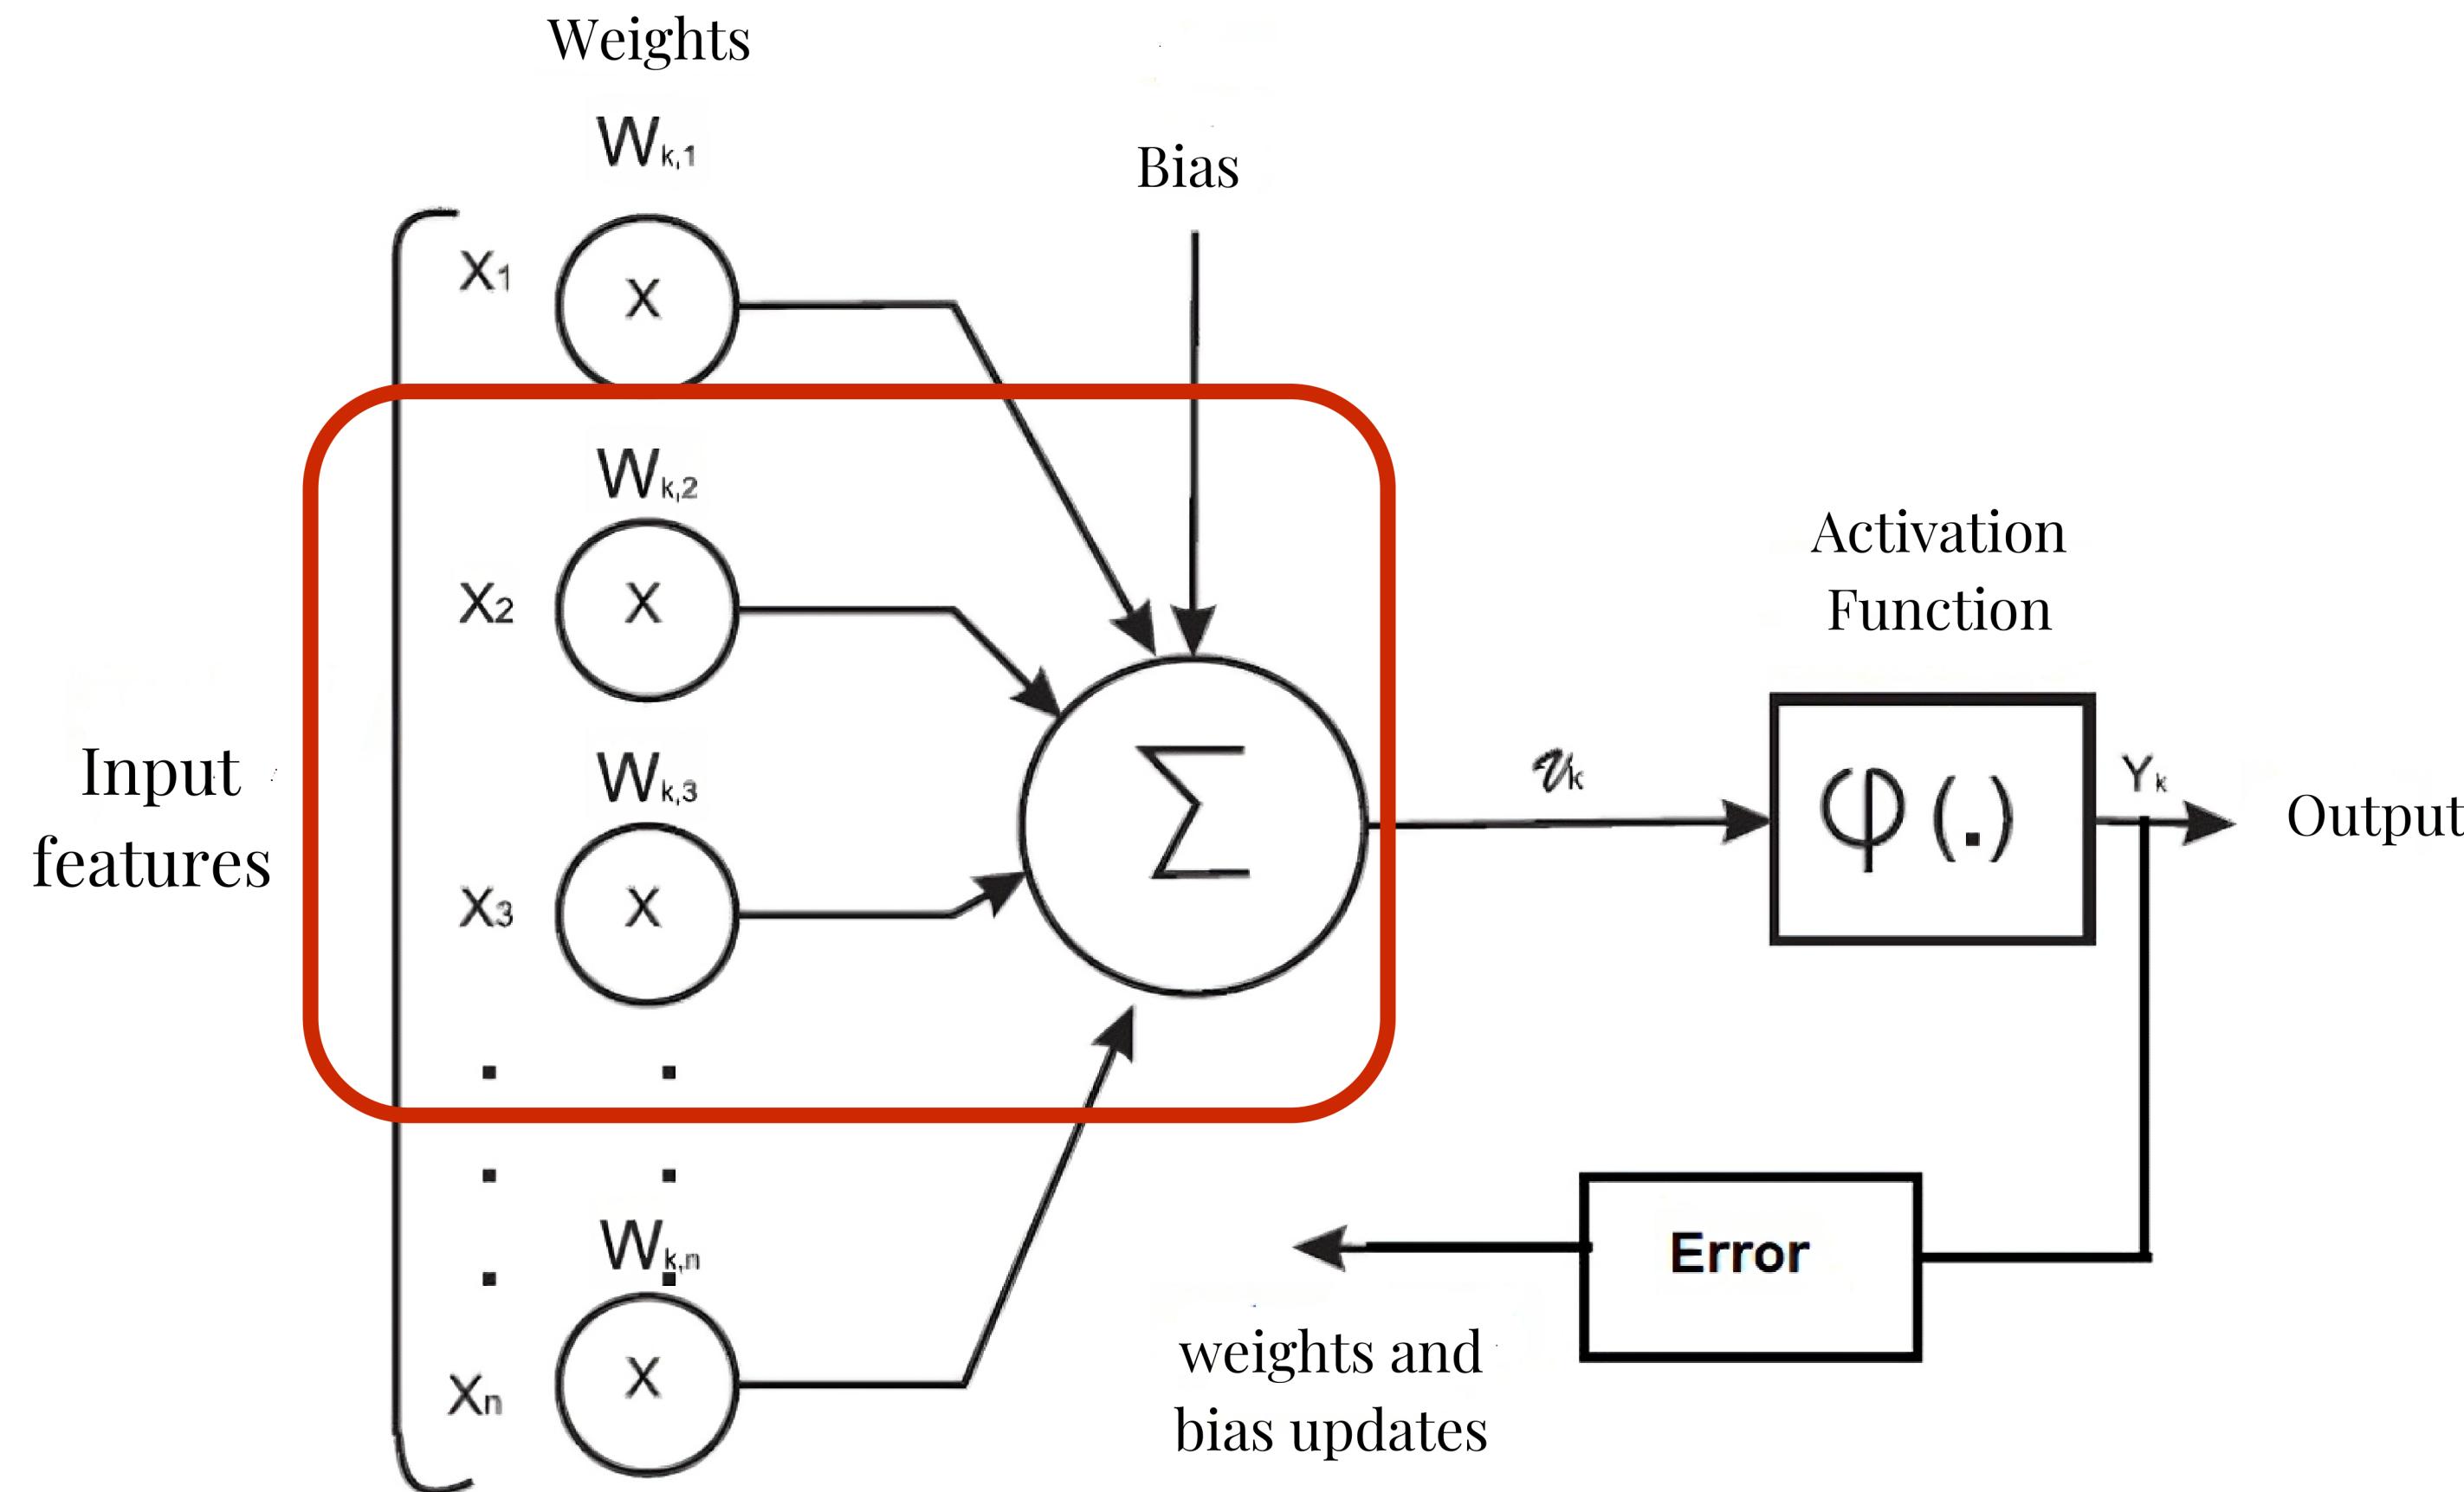
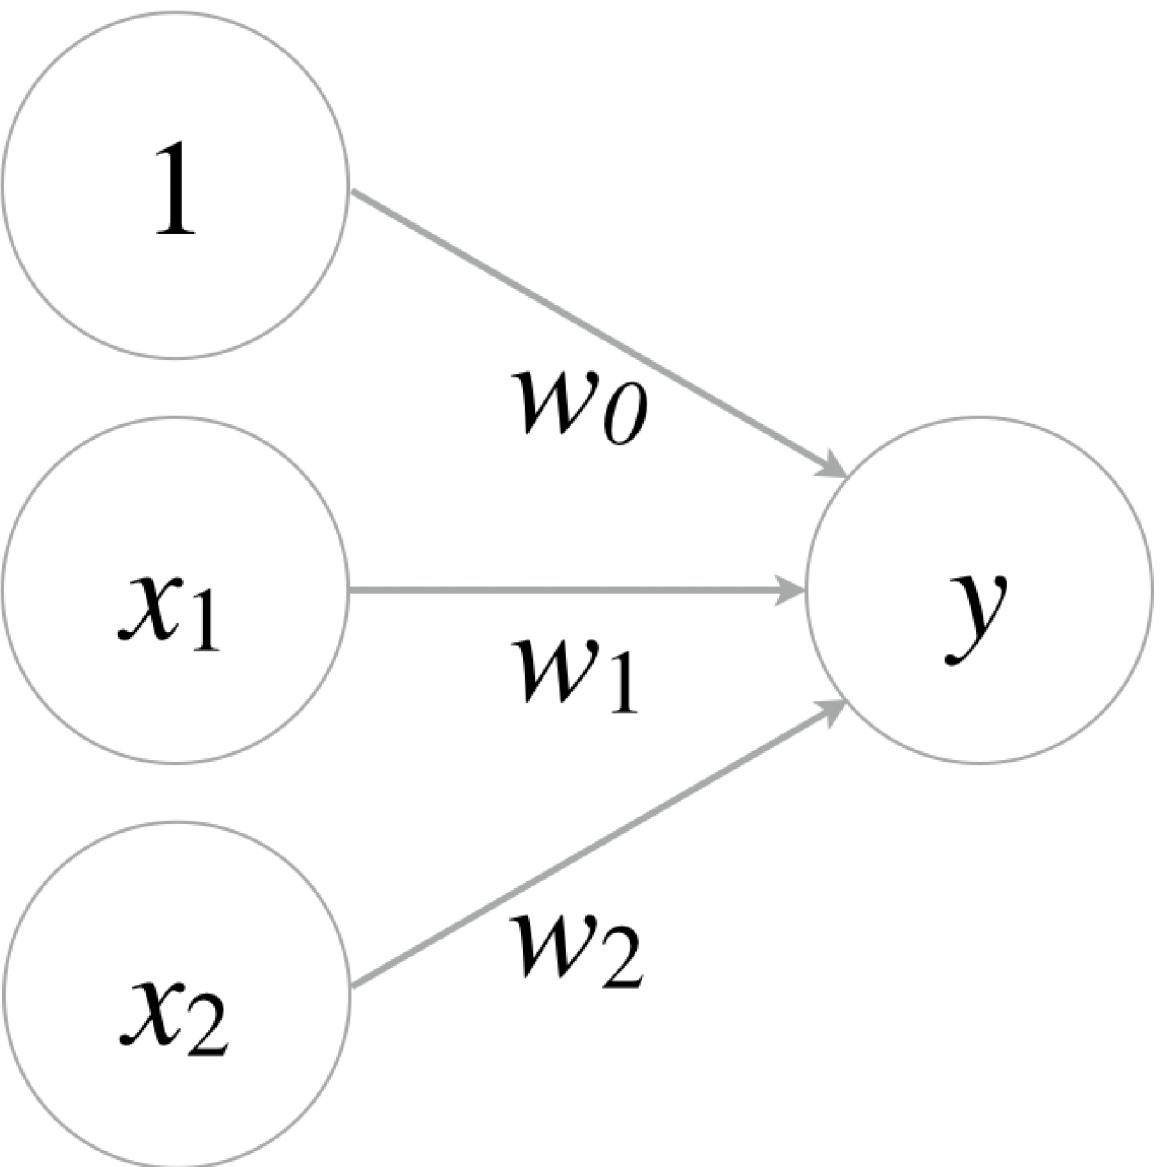


Diagram of a multilayer perceptron

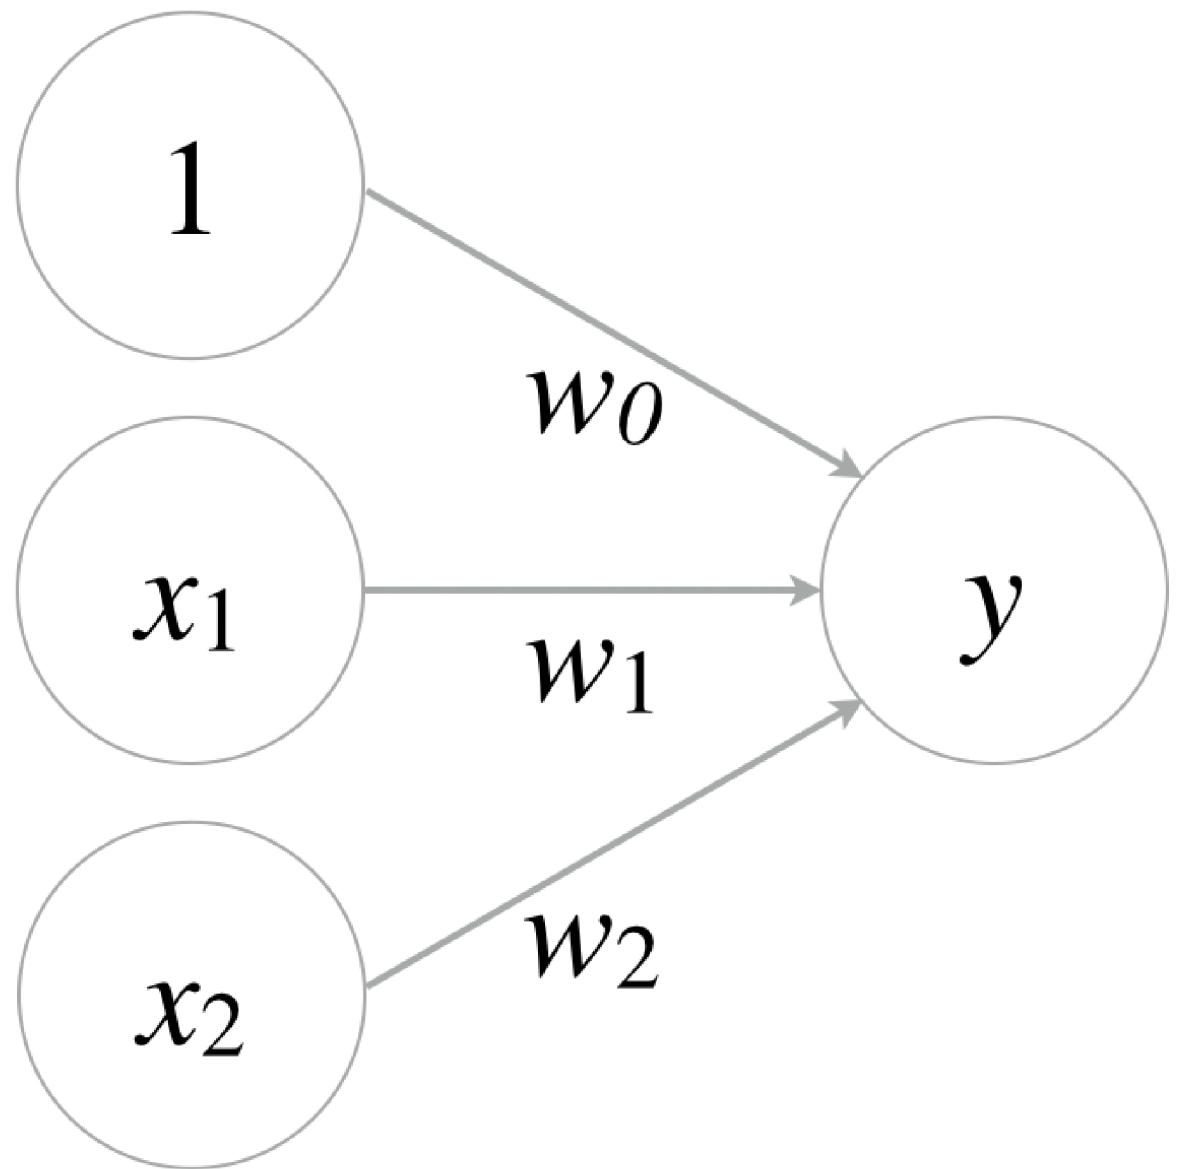


The perceptron



$$y = \sum_{i=0}^n x_n \cdot w_n$$

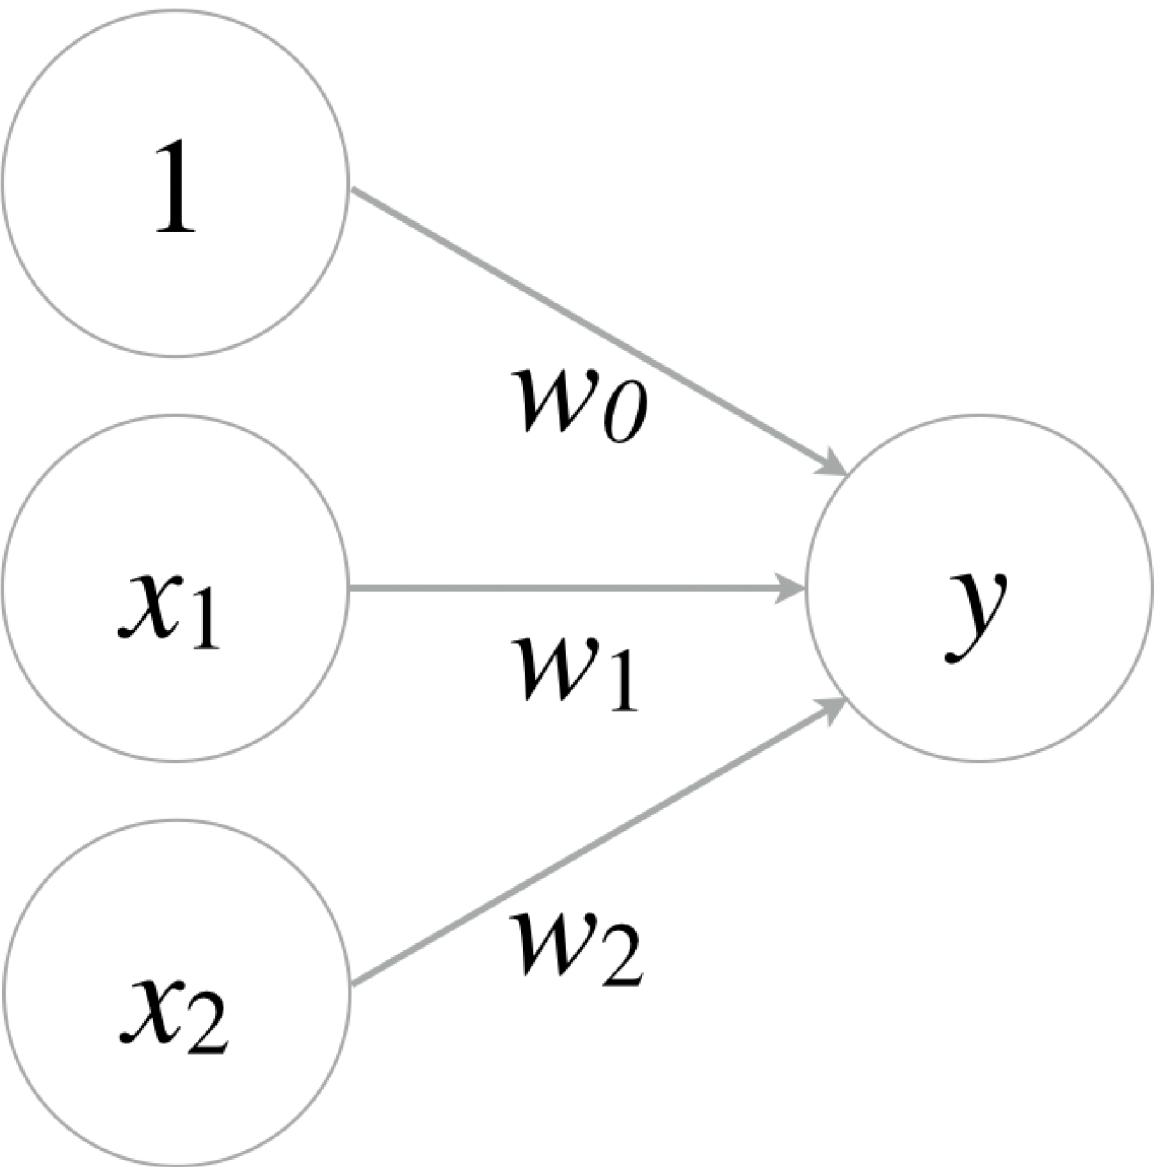
The perceptron



$$y = \sum_{i=0}^n x_n \cdot w_n$$

$$\hat{y} = w_0 + \sum_{i=1}^n x_n \cdot w_n$$

The perceptron



This can be simplified as:

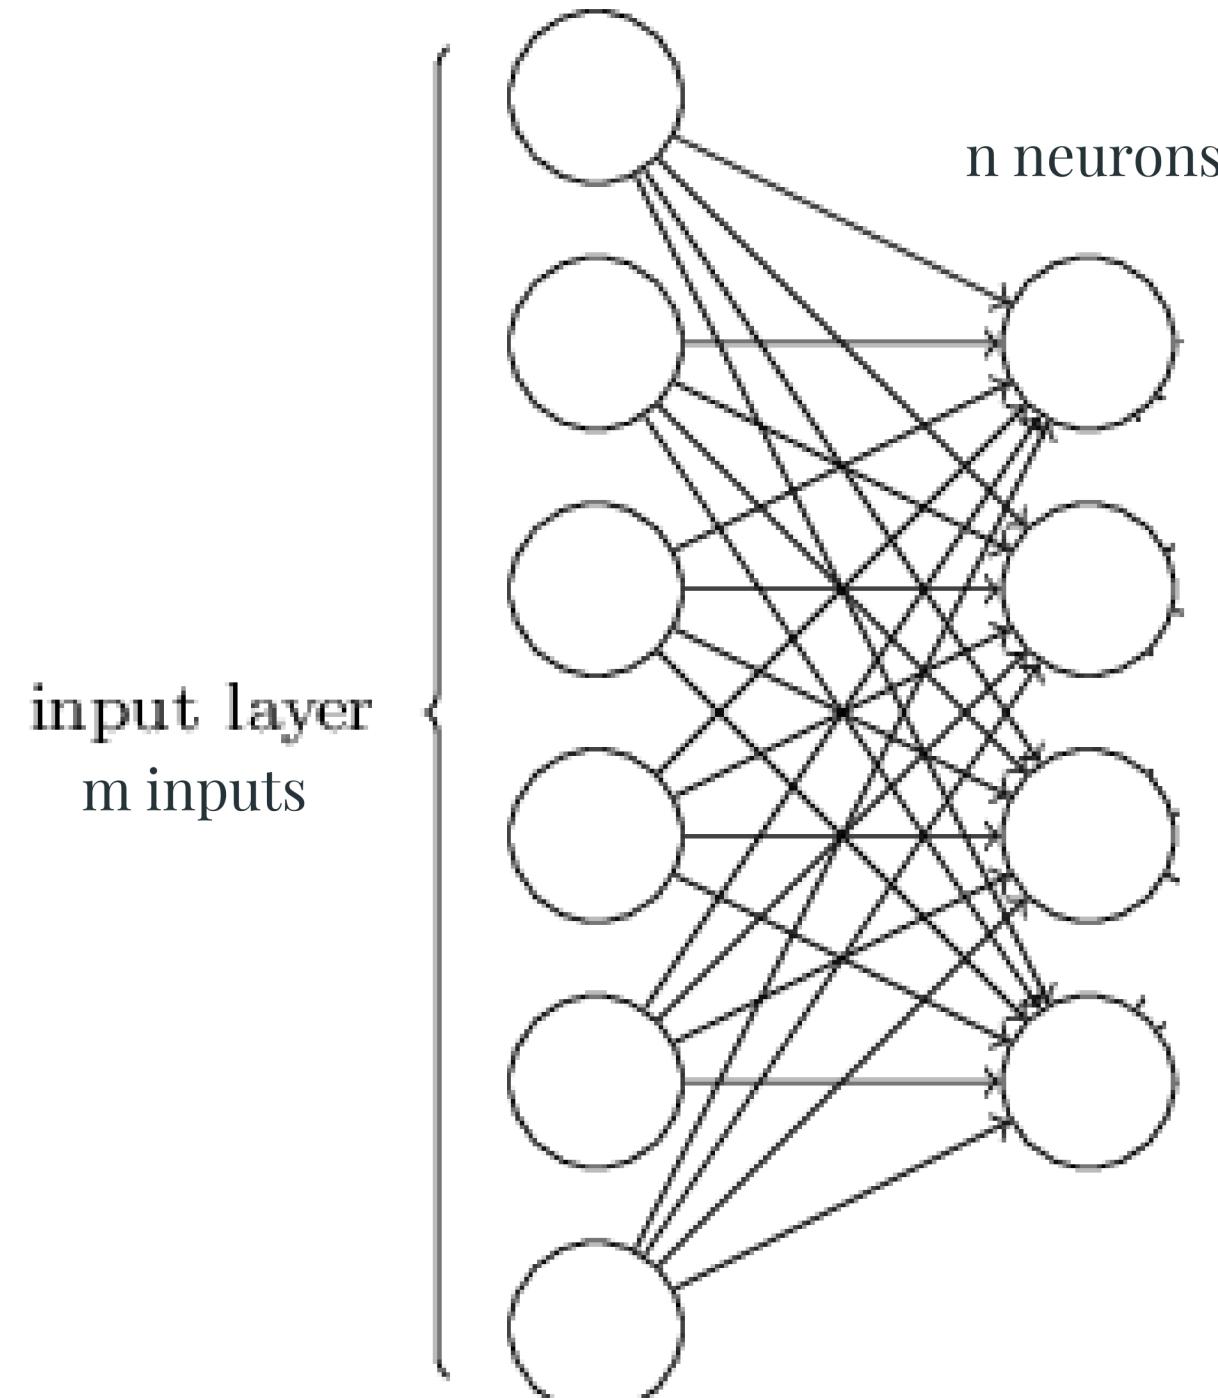
$$\hat{y} = \mathbf{w}^T \mathbf{x} + \mathbf{b}$$

where:

$$\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix}$$

$$\mathbf{w} = \begin{bmatrix} w_1 \\ w_2 \\ \vdots \\ w_n \end{bmatrix}$$

The perceptron



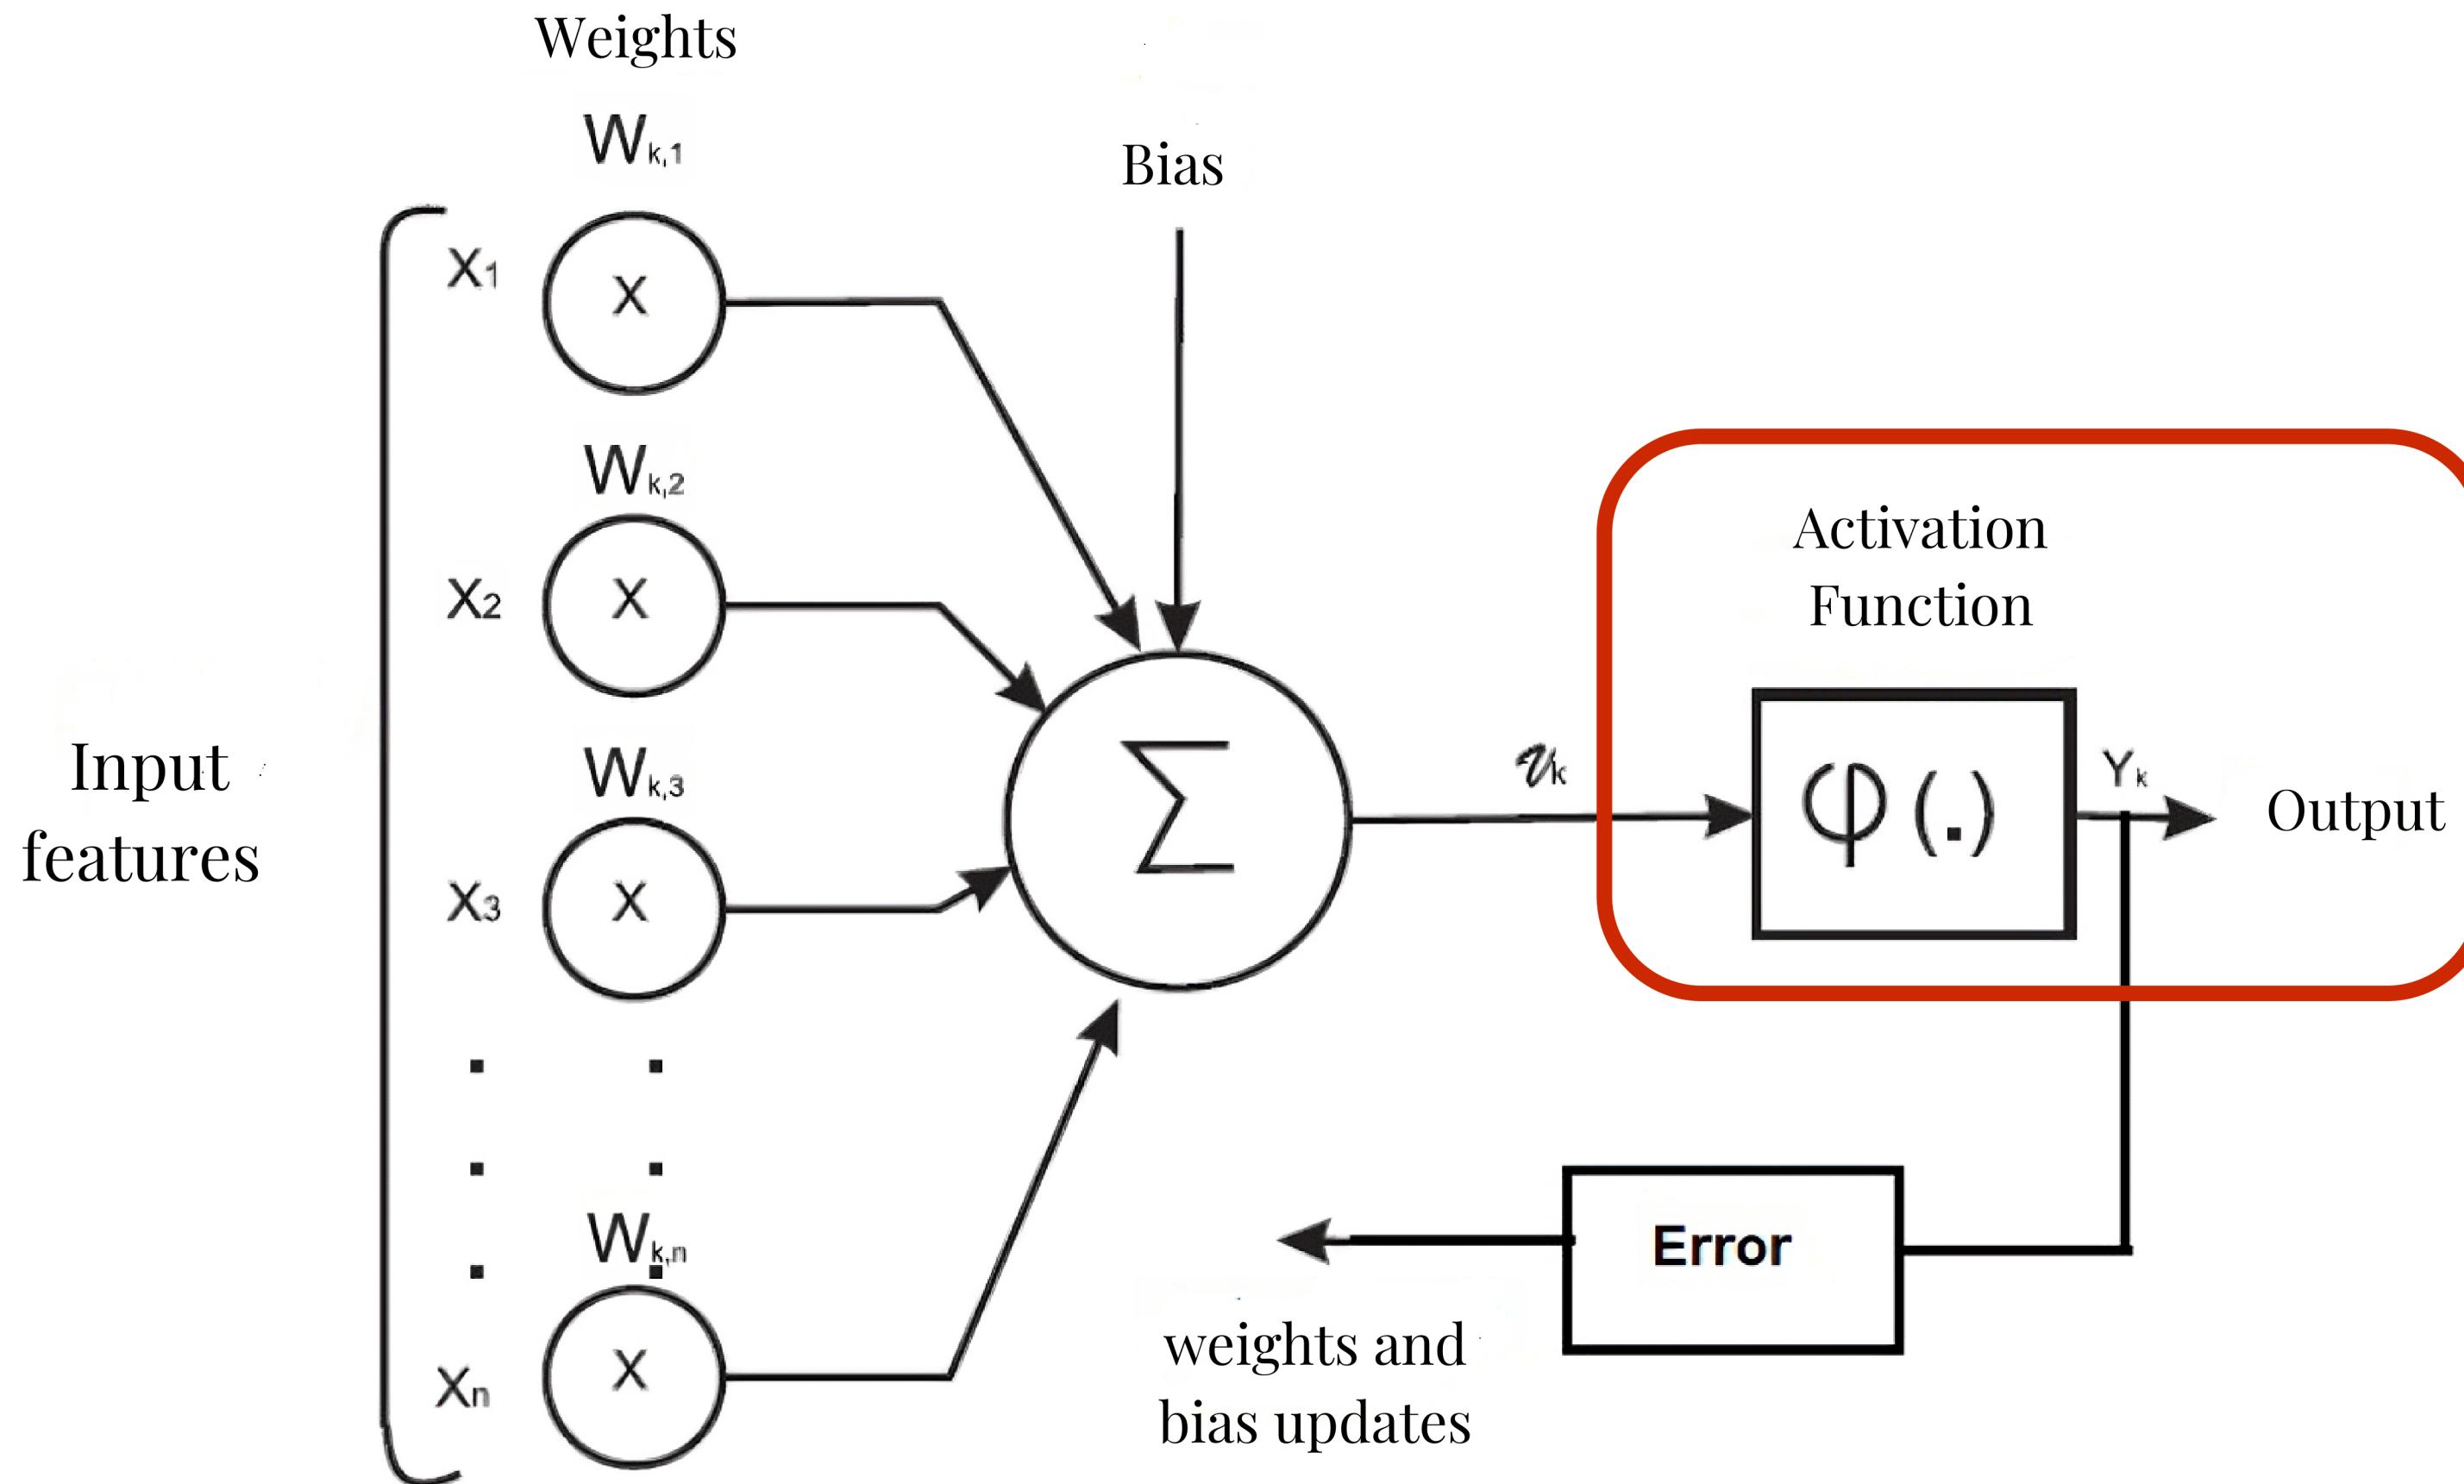
This can be generalized as:

$$X = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_m \end{bmatrix} \quad W = \begin{bmatrix} w_{11} & w_{12} & w_{13} & \cdots & w_{1n} \\ w_{21} & w_{22} & w_{23} & \cdots & w_{2n} \\ \vdots & & & & \vdots \\ w_{m1} & w_{m2} & w_{m3} & \cdots & w_{mn} \end{bmatrix}$$

and obtain;

$$\hat{\mathbf{y}} = \mathbf{W}^T \mathbf{x} + \mathbf{b}$$

Diagram of a multilayer perceptron



Activation Functions

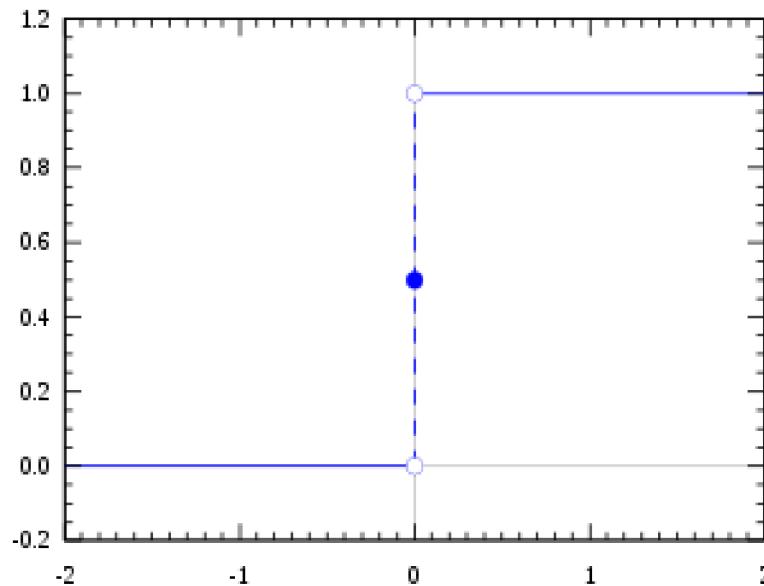
can be expressed as applying a function to the output:

$$\hat{\mathbf{y}} = g(\mathbf{W}^T \mathbf{x} + \mathbf{b})$$

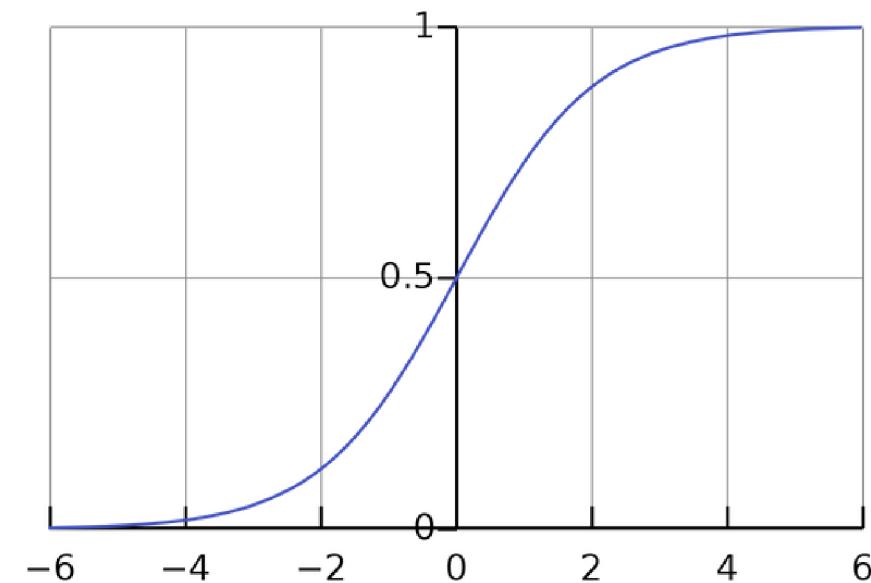
any function can be applied, but a **non-linear function** is expected to be applied

Common Activation Functions

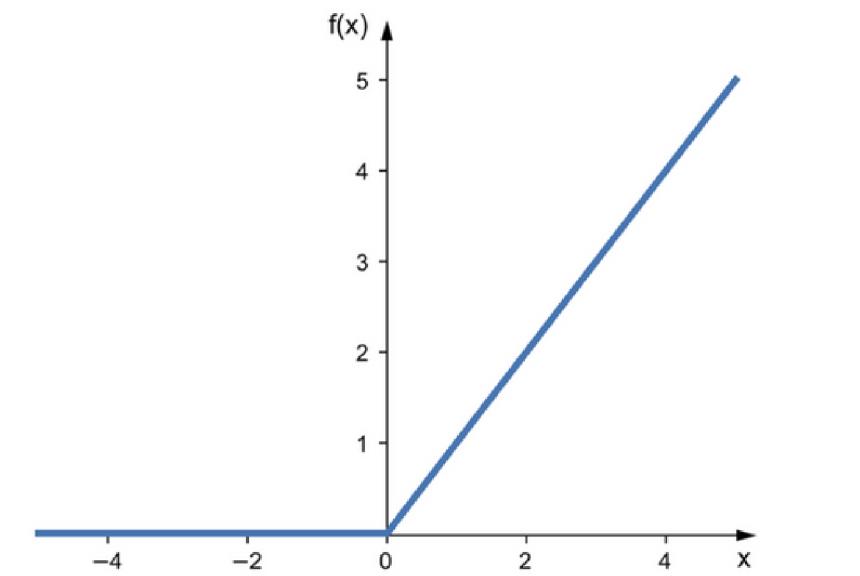
Heaviside



Sigmoid



RELU

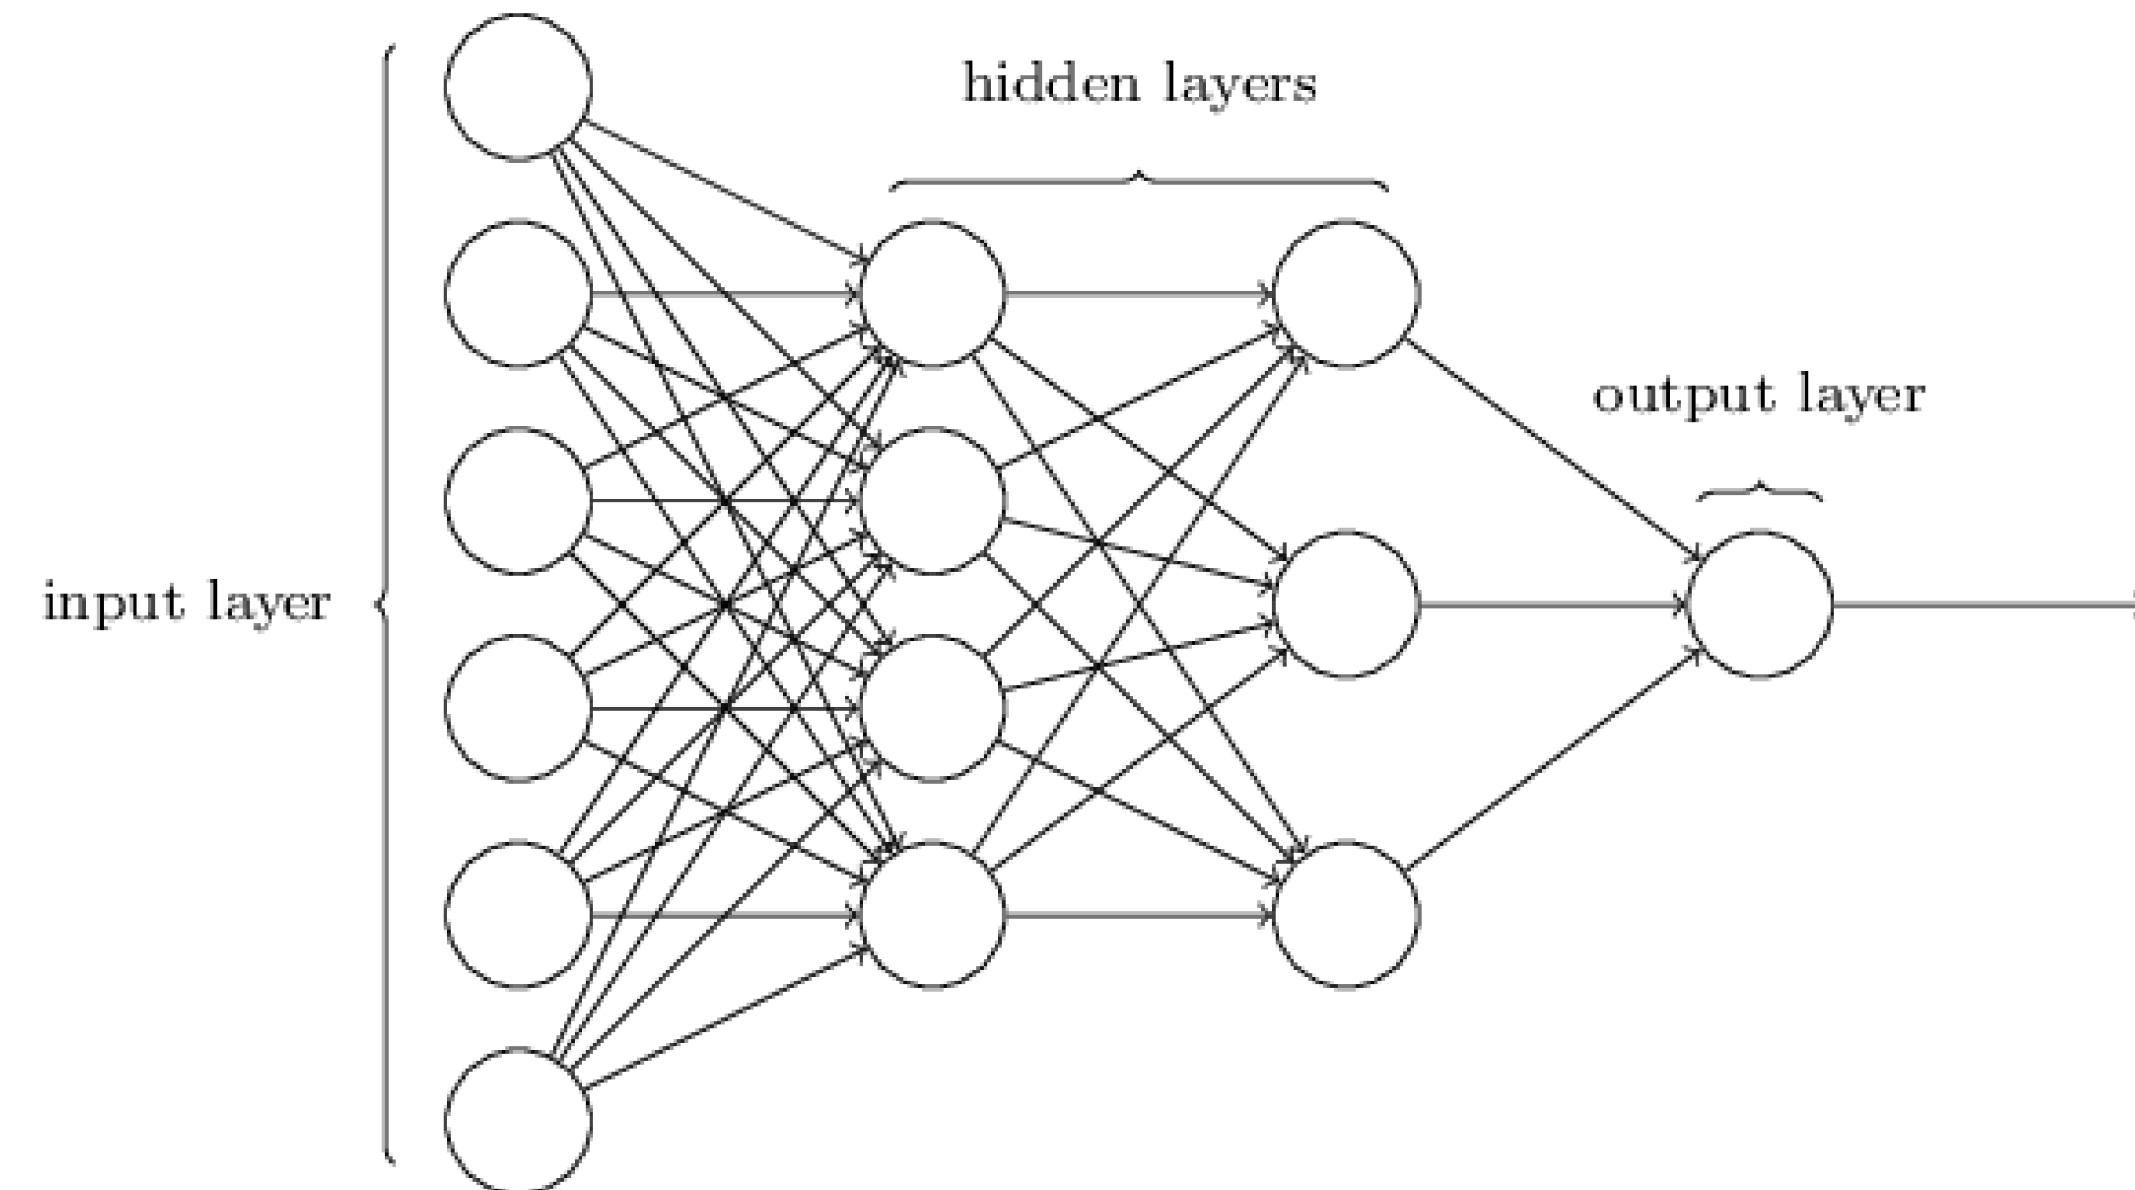


$$g(x) = \begin{cases} 0 & \text{if } x > 0 \\ 1 & \text{if } x \leq 0 \end{cases}$$

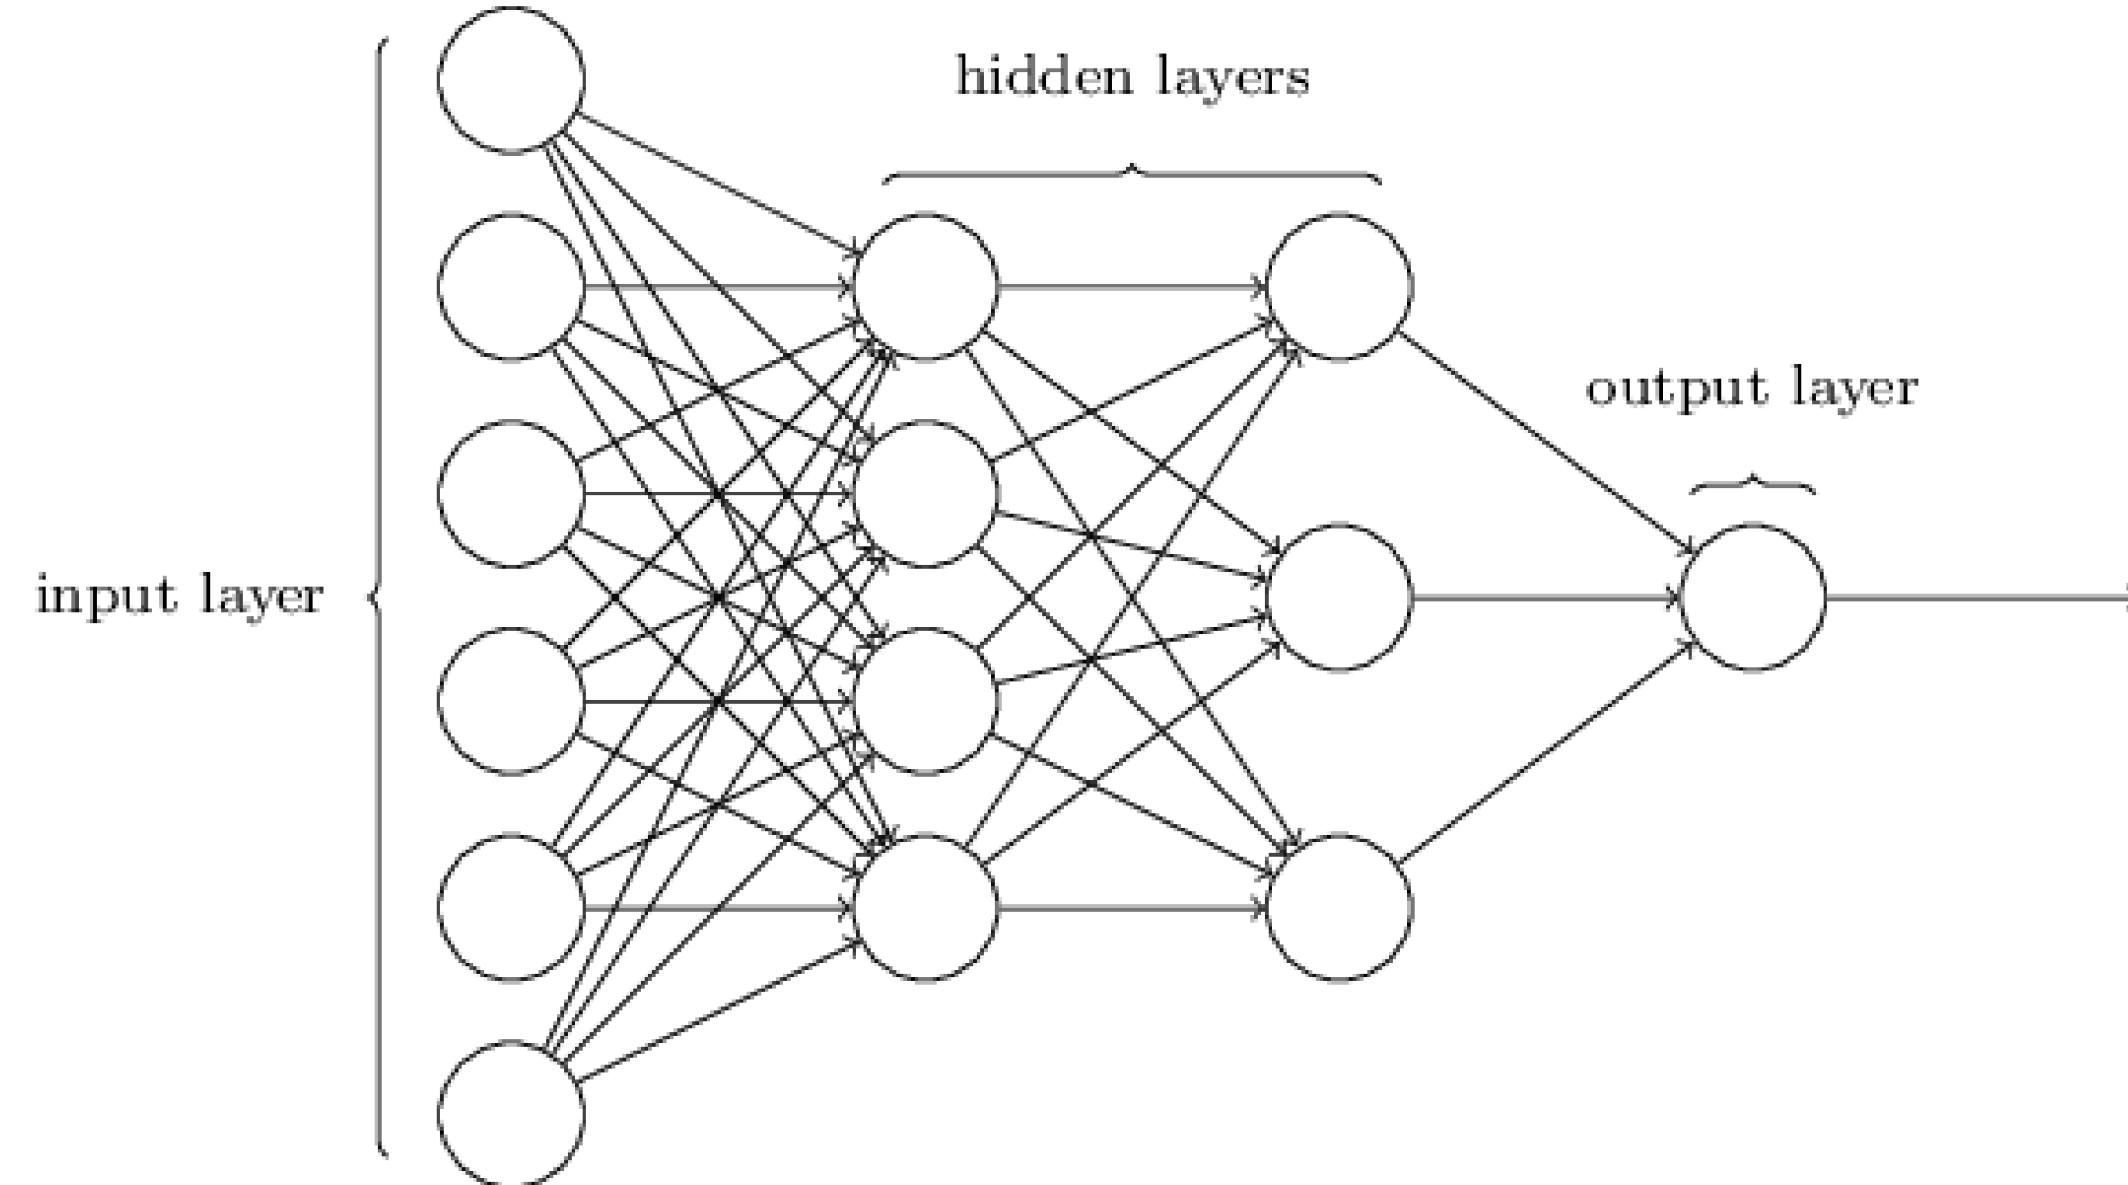
$$g(x) = \frac{1}{1 + e^{-x}}$$

$$g(x) = \max(0, x)$$

The perceptron

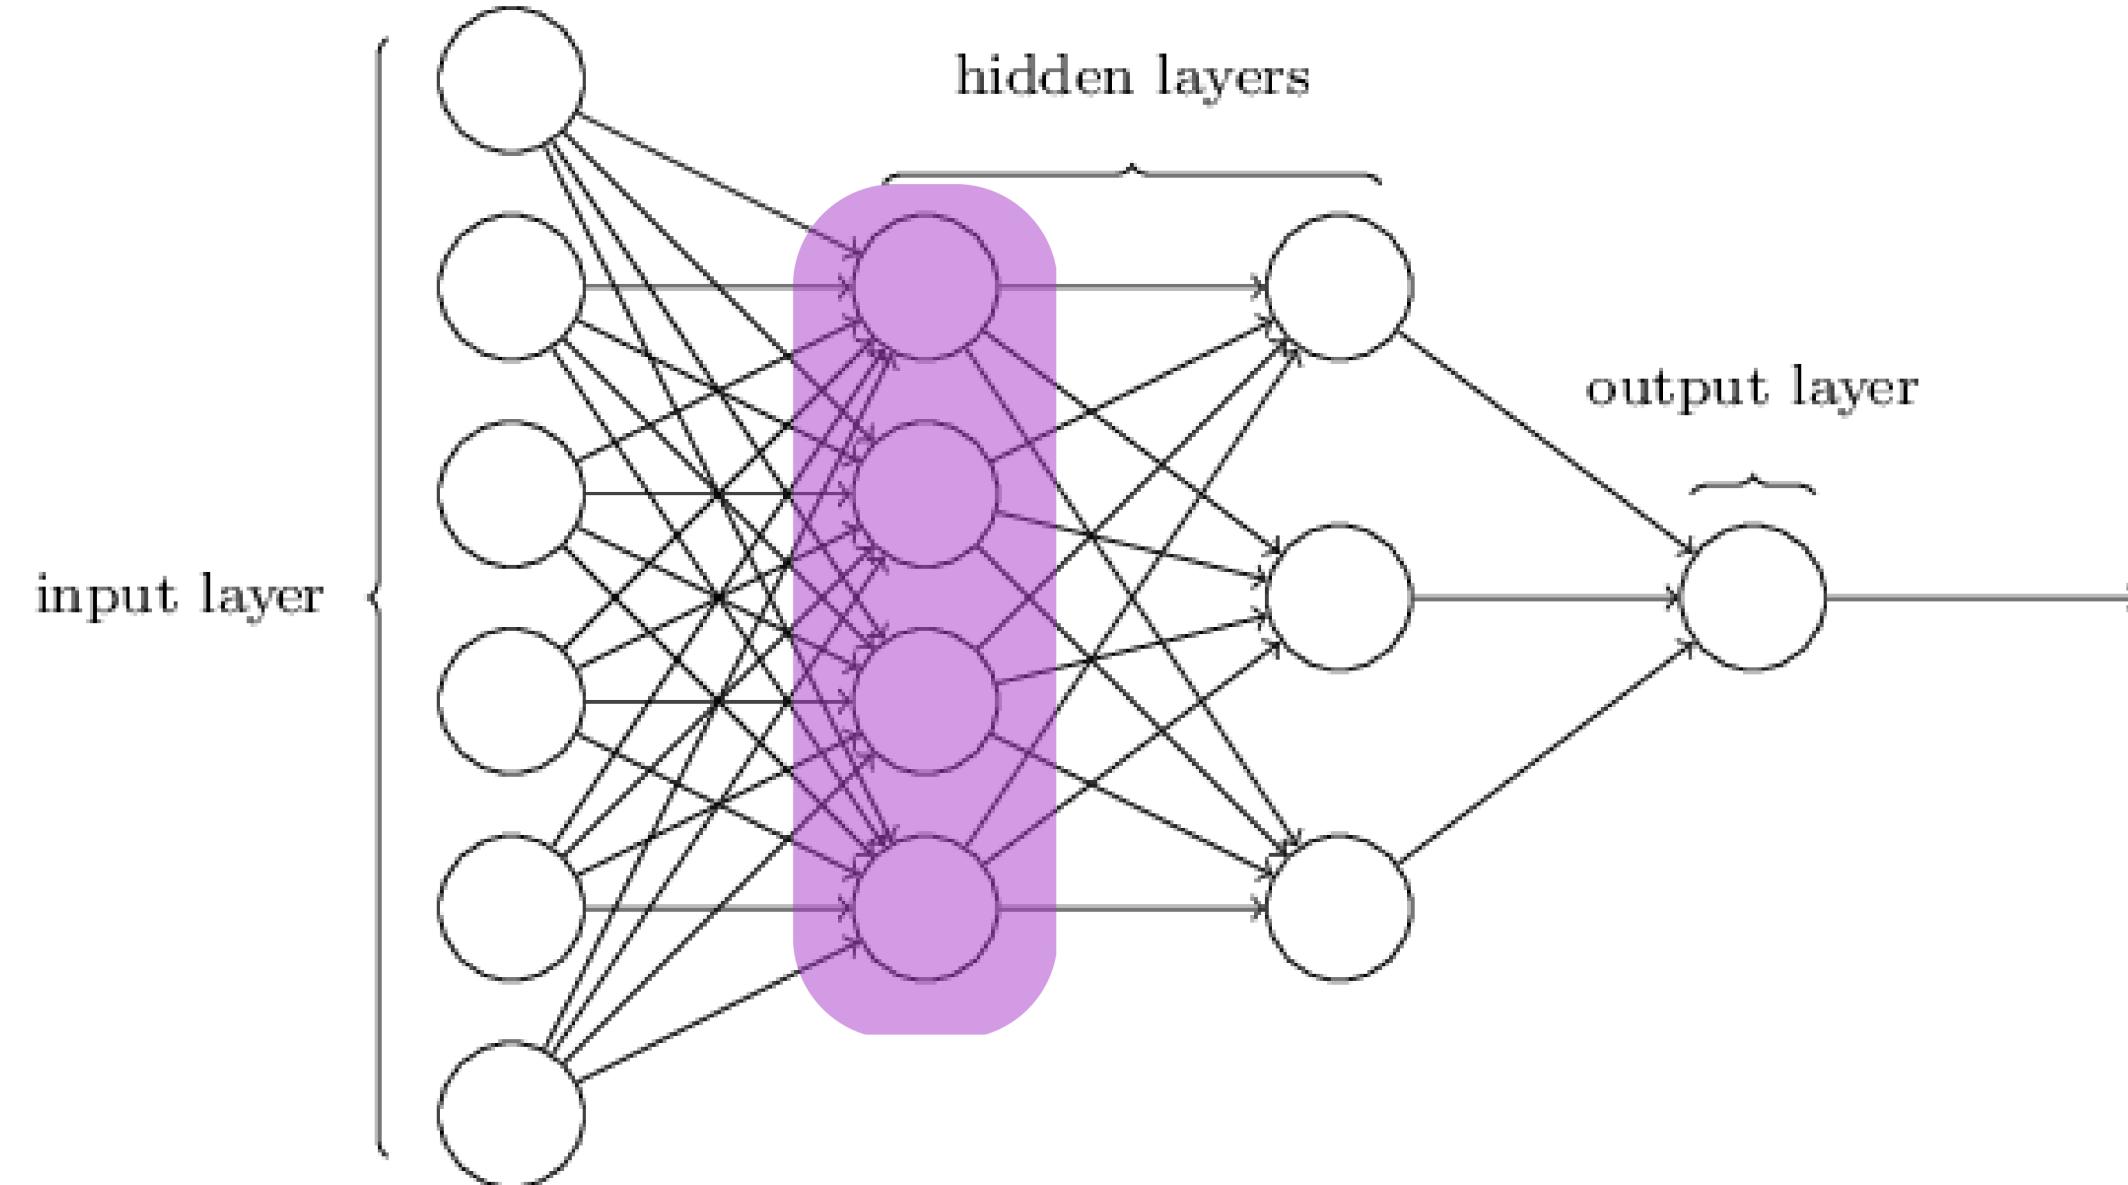


The perceptron



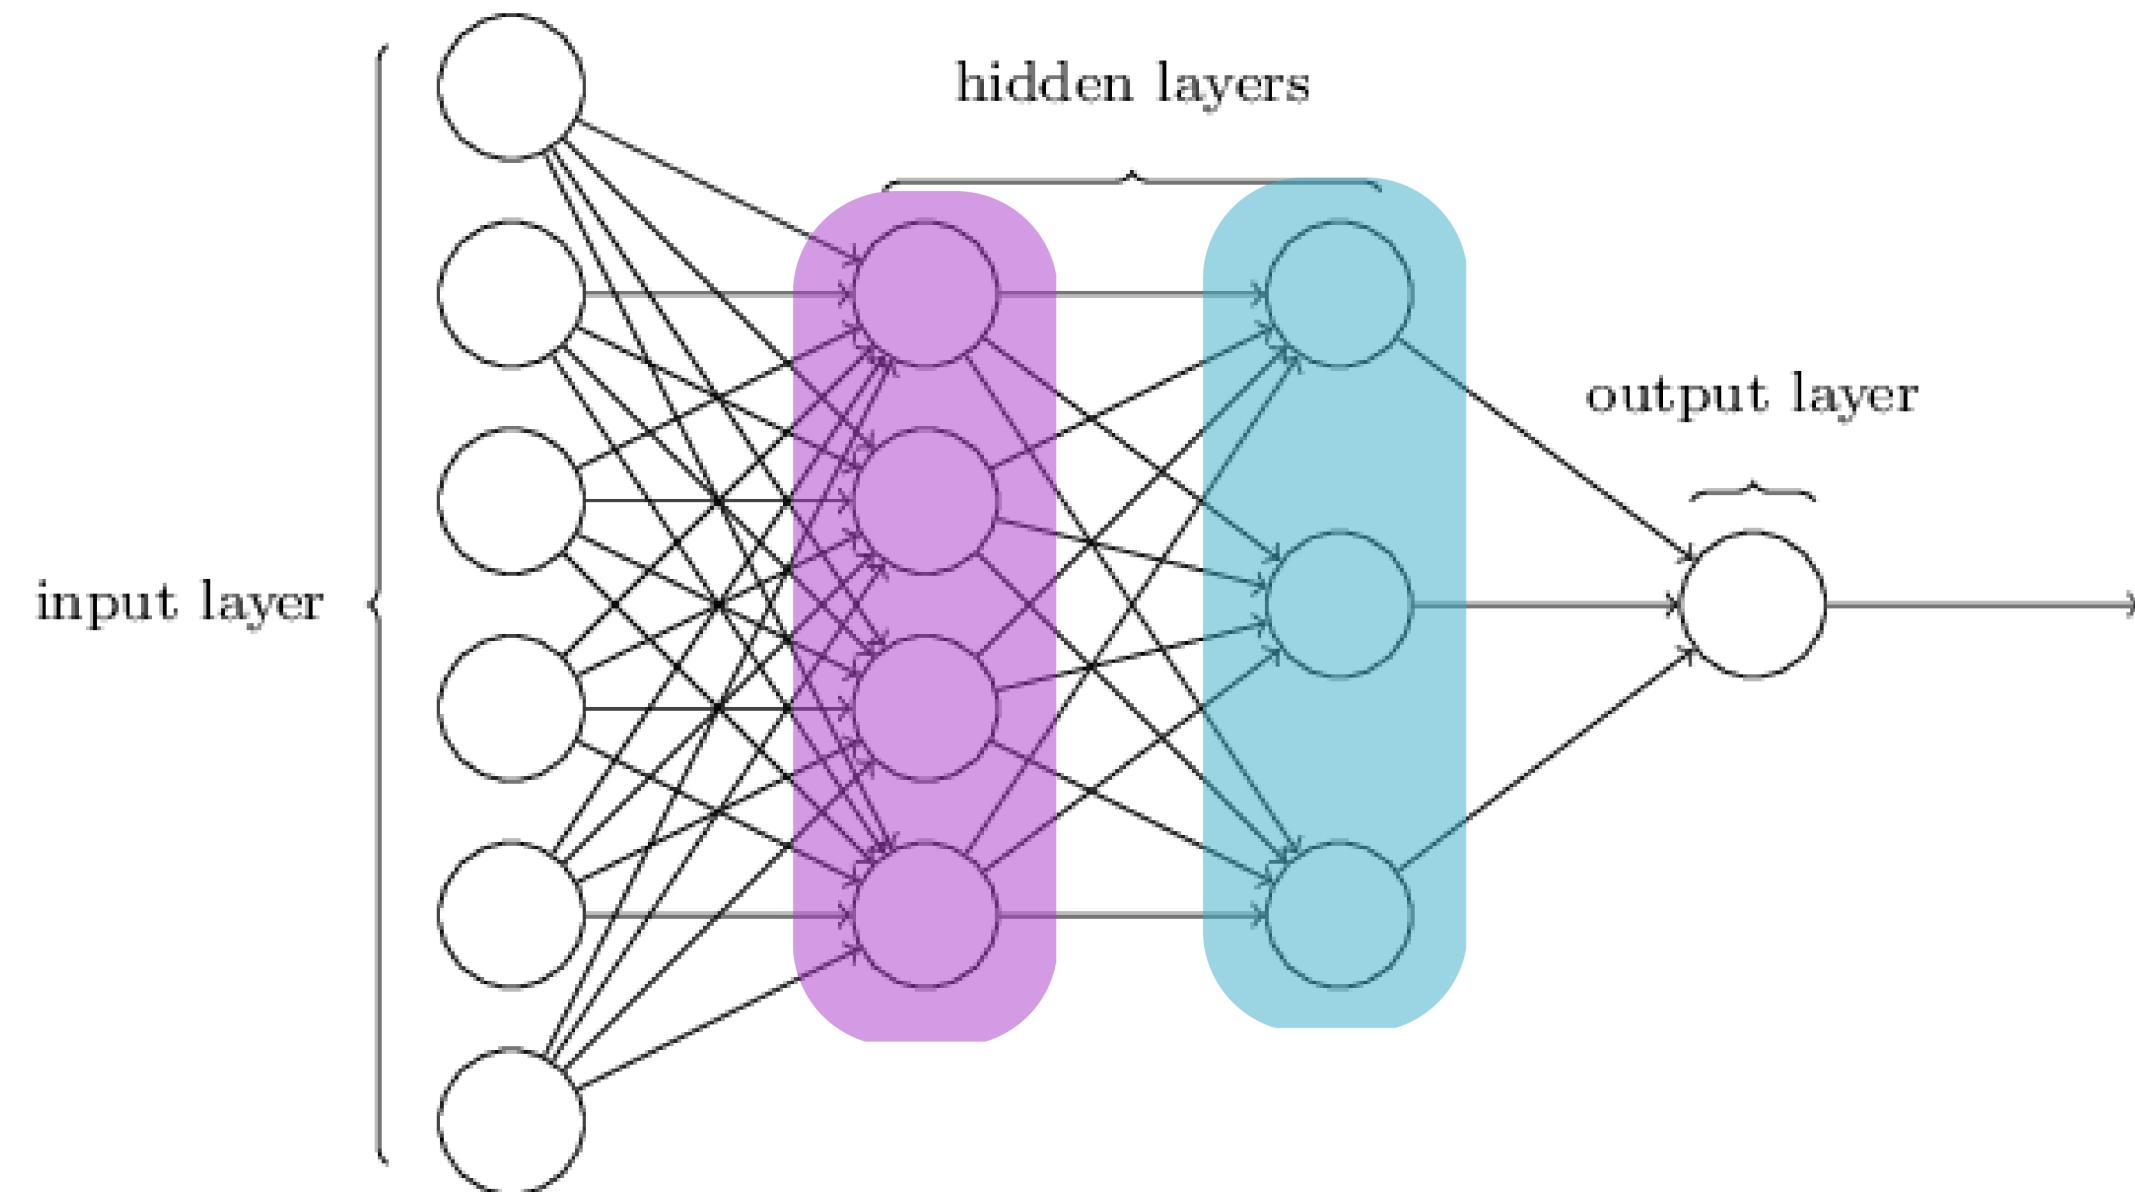
$$\hat{y} = W_3 \cdot g(W_2 \cdot g(W_1 \cdot X + b_1) + b_2) + b_3$$

The perceptron



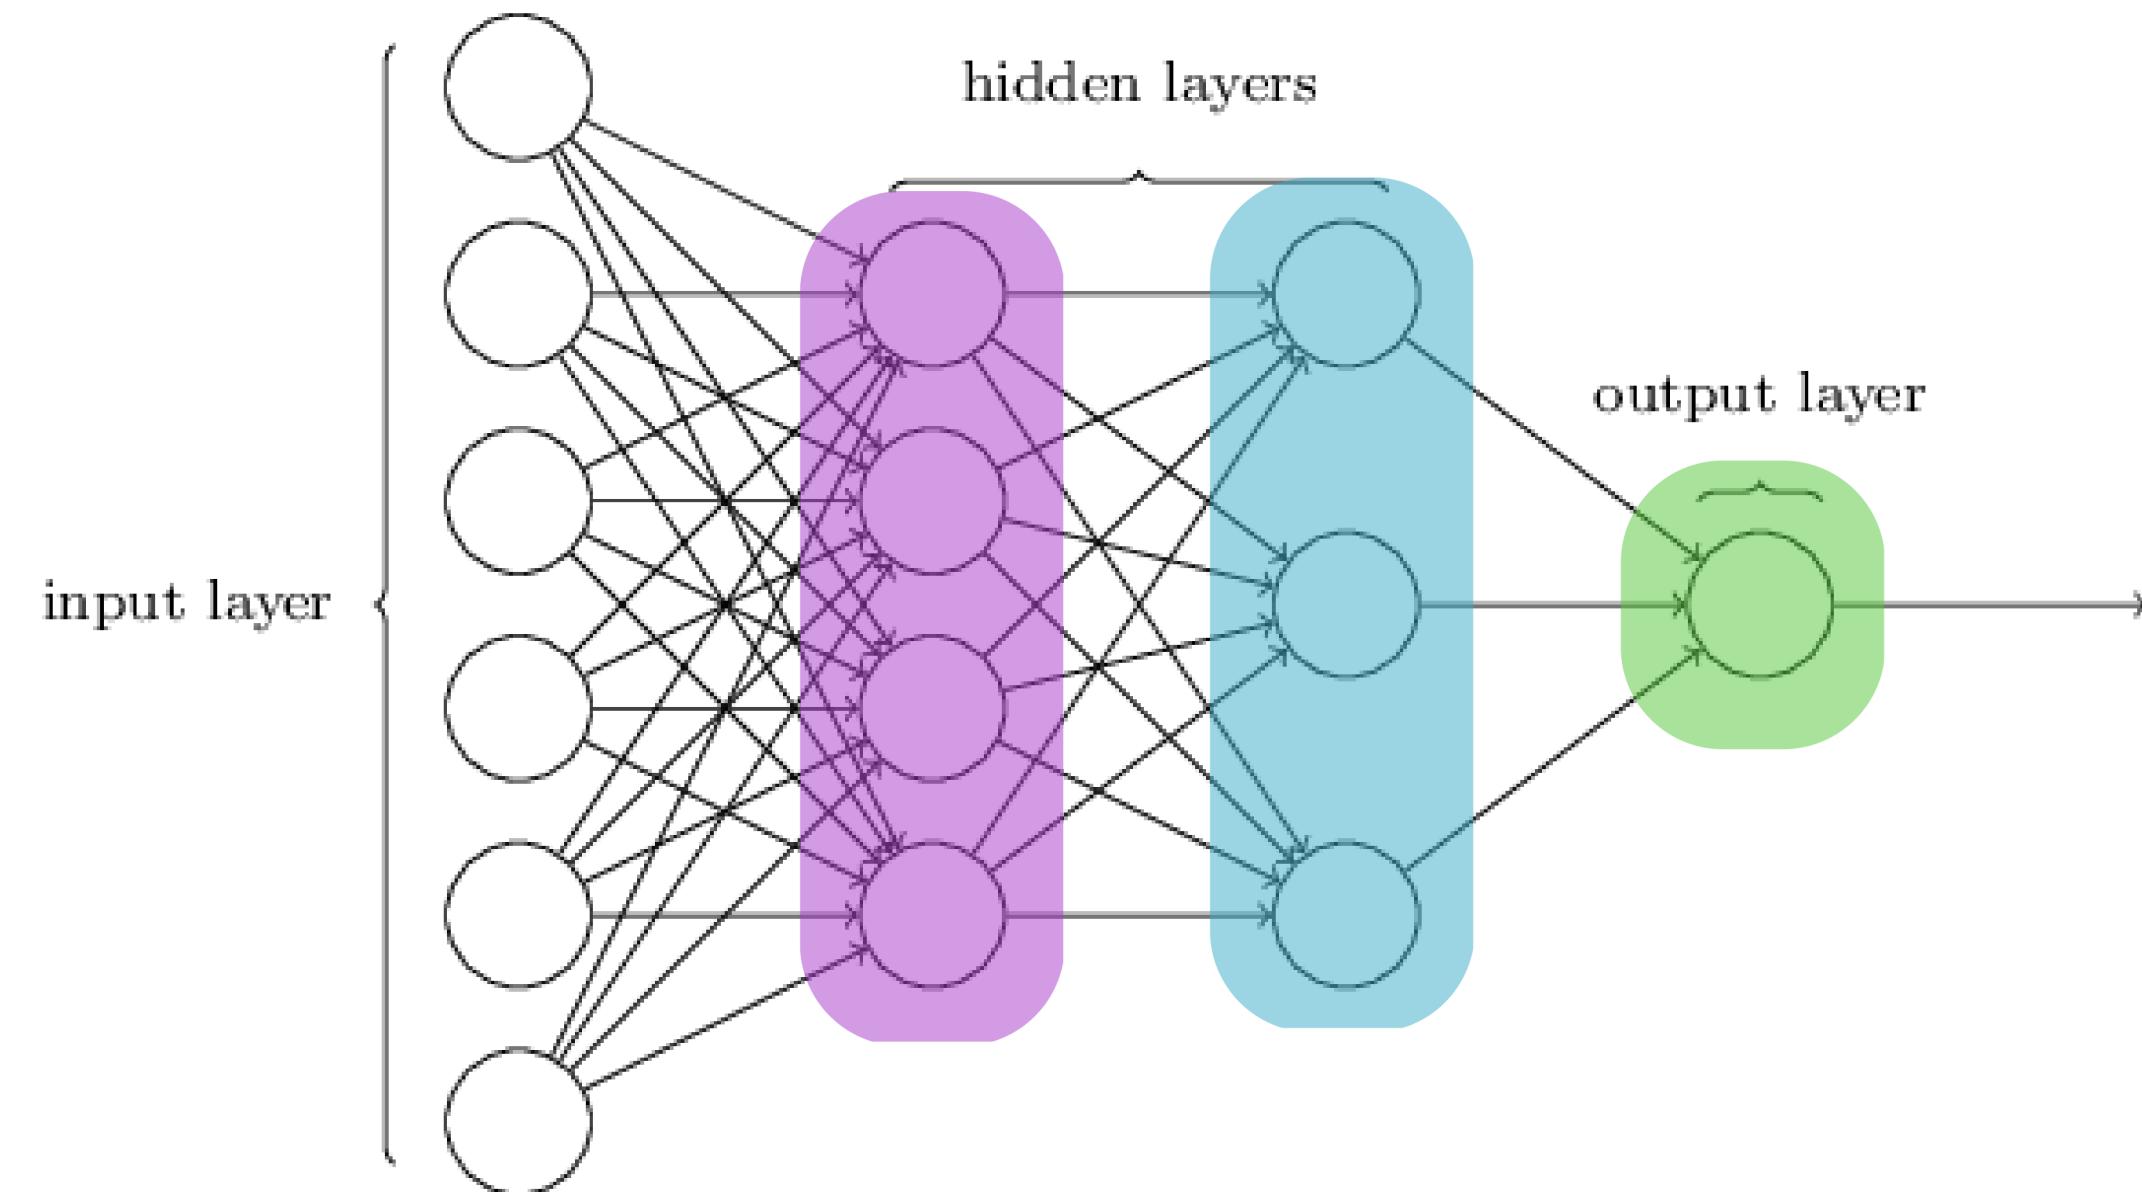
$$\hat{y} = W_3 \cdot g(W_2 \cdot g(W_1 \cdot X + b_1) + b_2) + b_3$$

The perceptron



$$\hat{y} = W_3 \cdot g(W_2 \cdot g(W_1 \cdot X + b_1) + b_2) + b_3$$

The perceptron



$$\hat{y} = W_3 \cdot g(W_2 \cdot g(W_1 \cdot X + b_1) + b_2) + b_3$$



```
# forward-pass of a 3-layer neural network:

f = lambda x: 1.0/(1.0 + np.exp(-x)) # we will use sigmoid
X = np.random.randn(1,3) # random input vector (3x1)

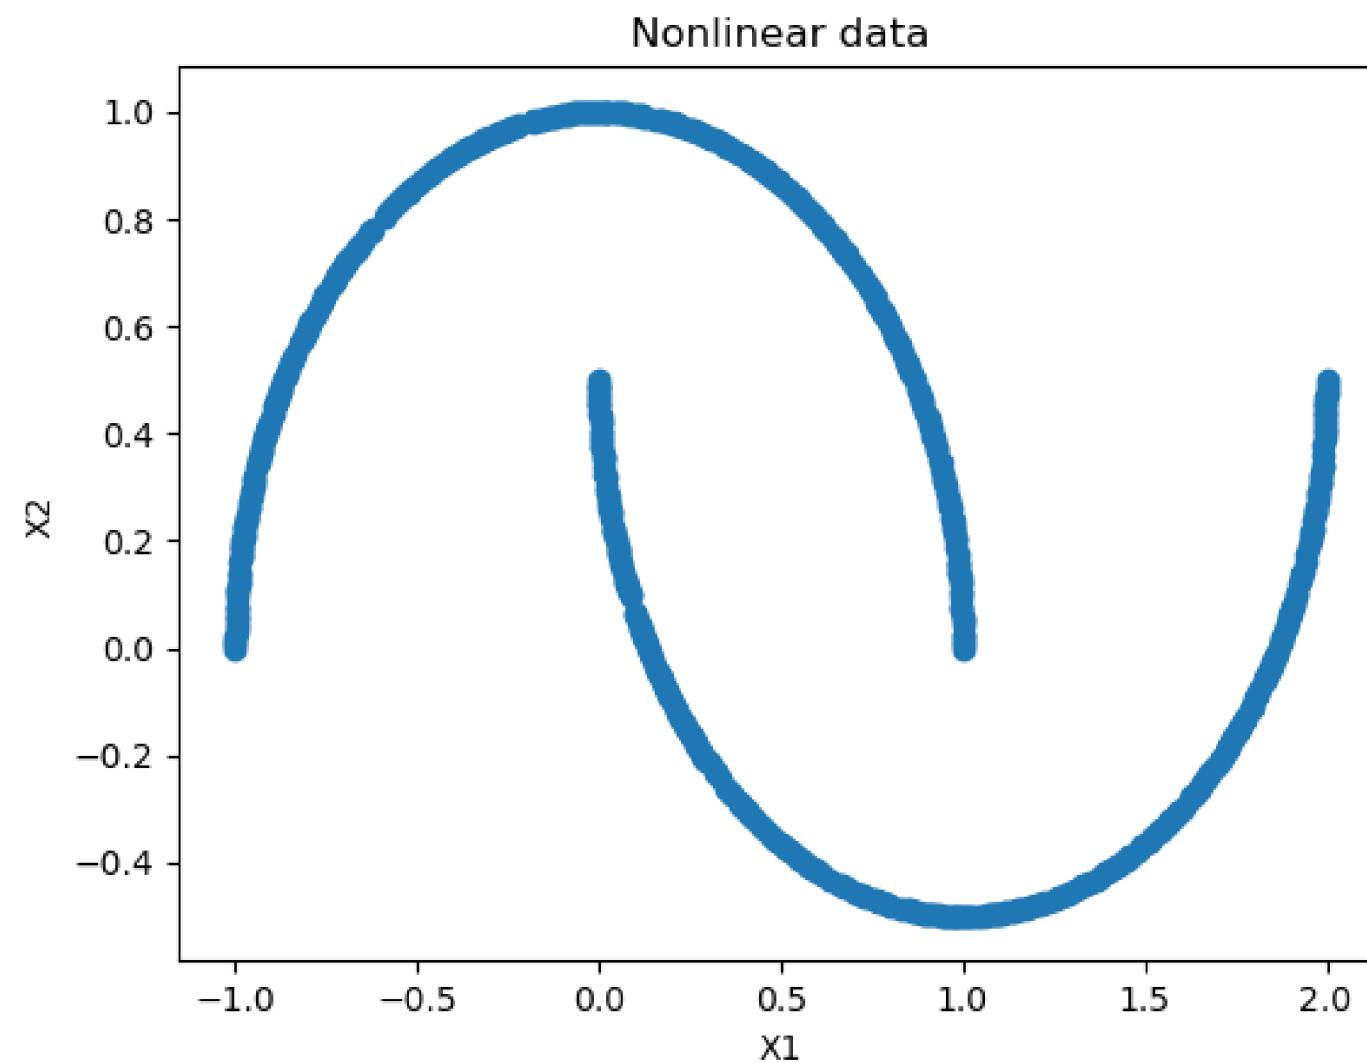
W1 = np.random.randn(3,8)
W2 = np.random.randn(8,8)
W3 = np.random.randn(8,1)
b1,b2,b3 = 0.7,2.5,9

h1 = f(np.dot(X,W1) + b1)
h2 = f(np.dot(h1,W2) + b2)

out = f(np.dot(h2,W3) + b3)
```

Activation Functions

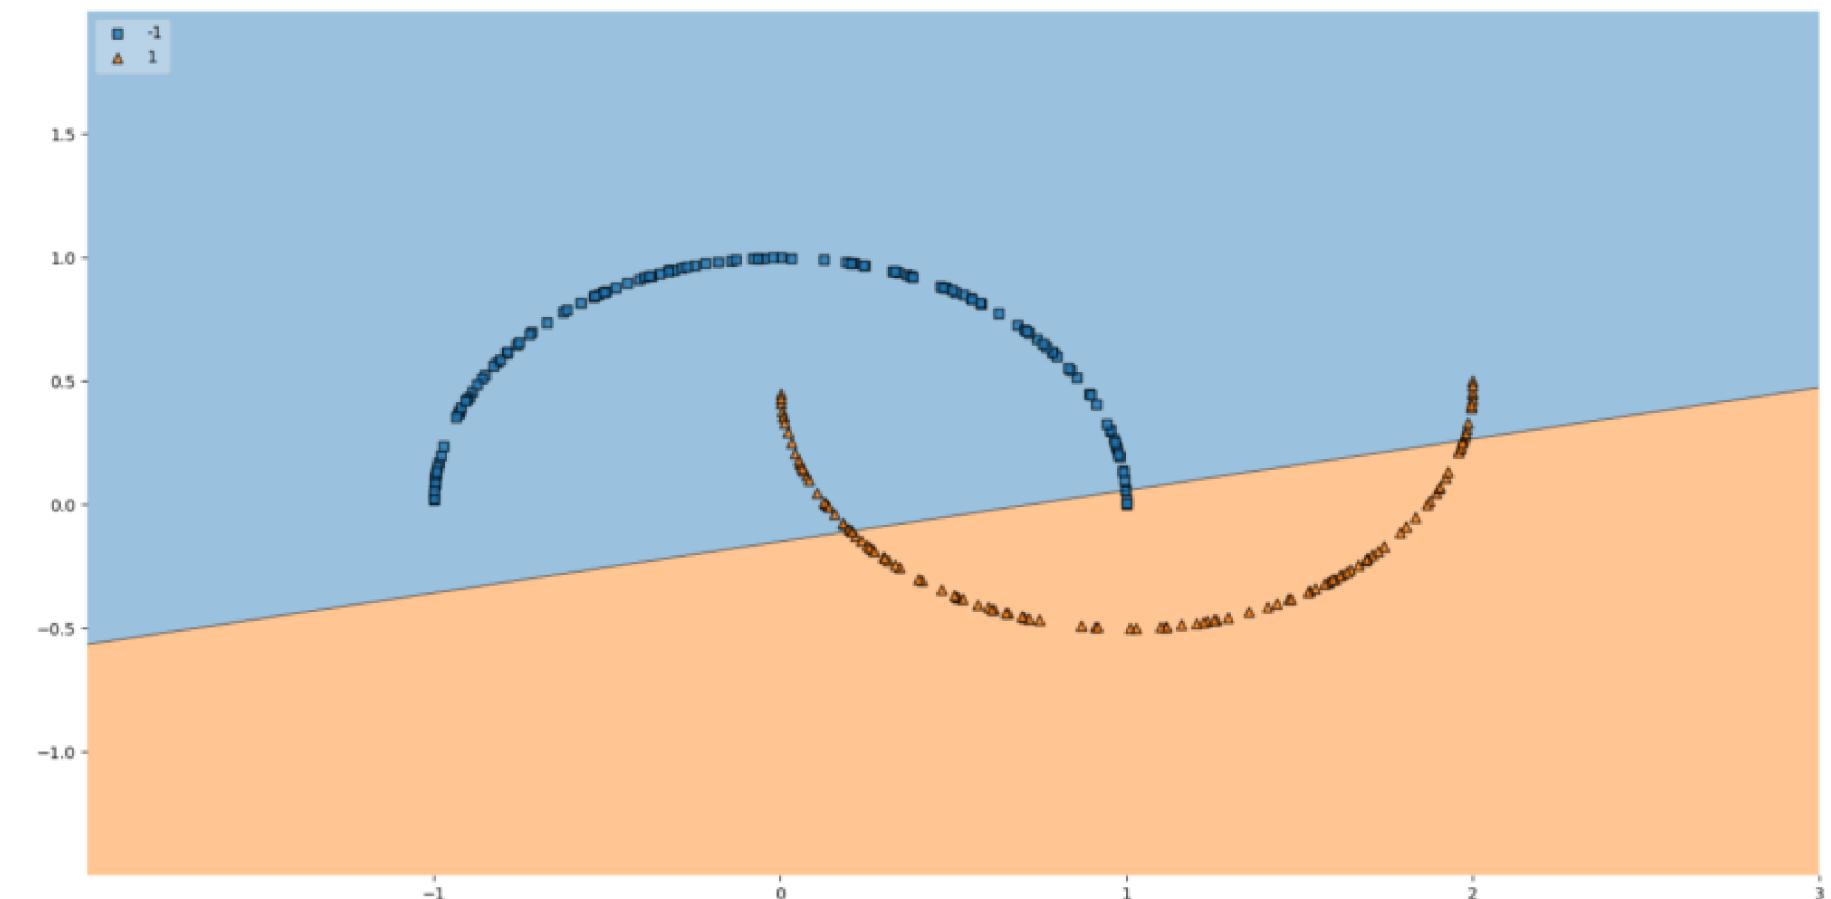
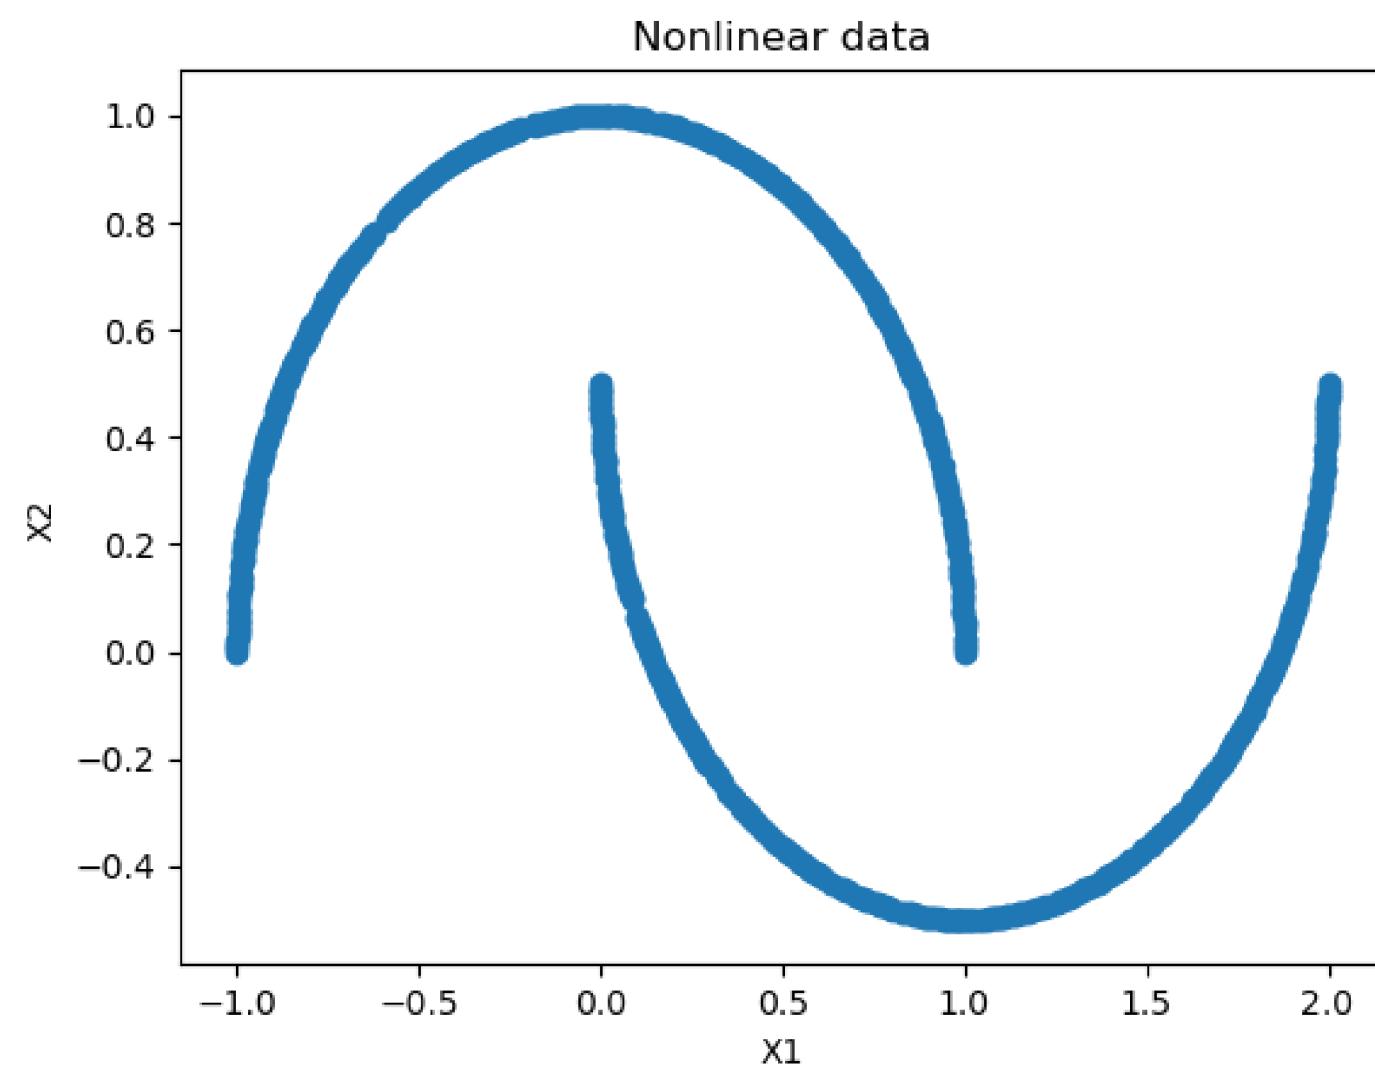
why is non-linearity necessary?



almost all data or problems are nonlinear

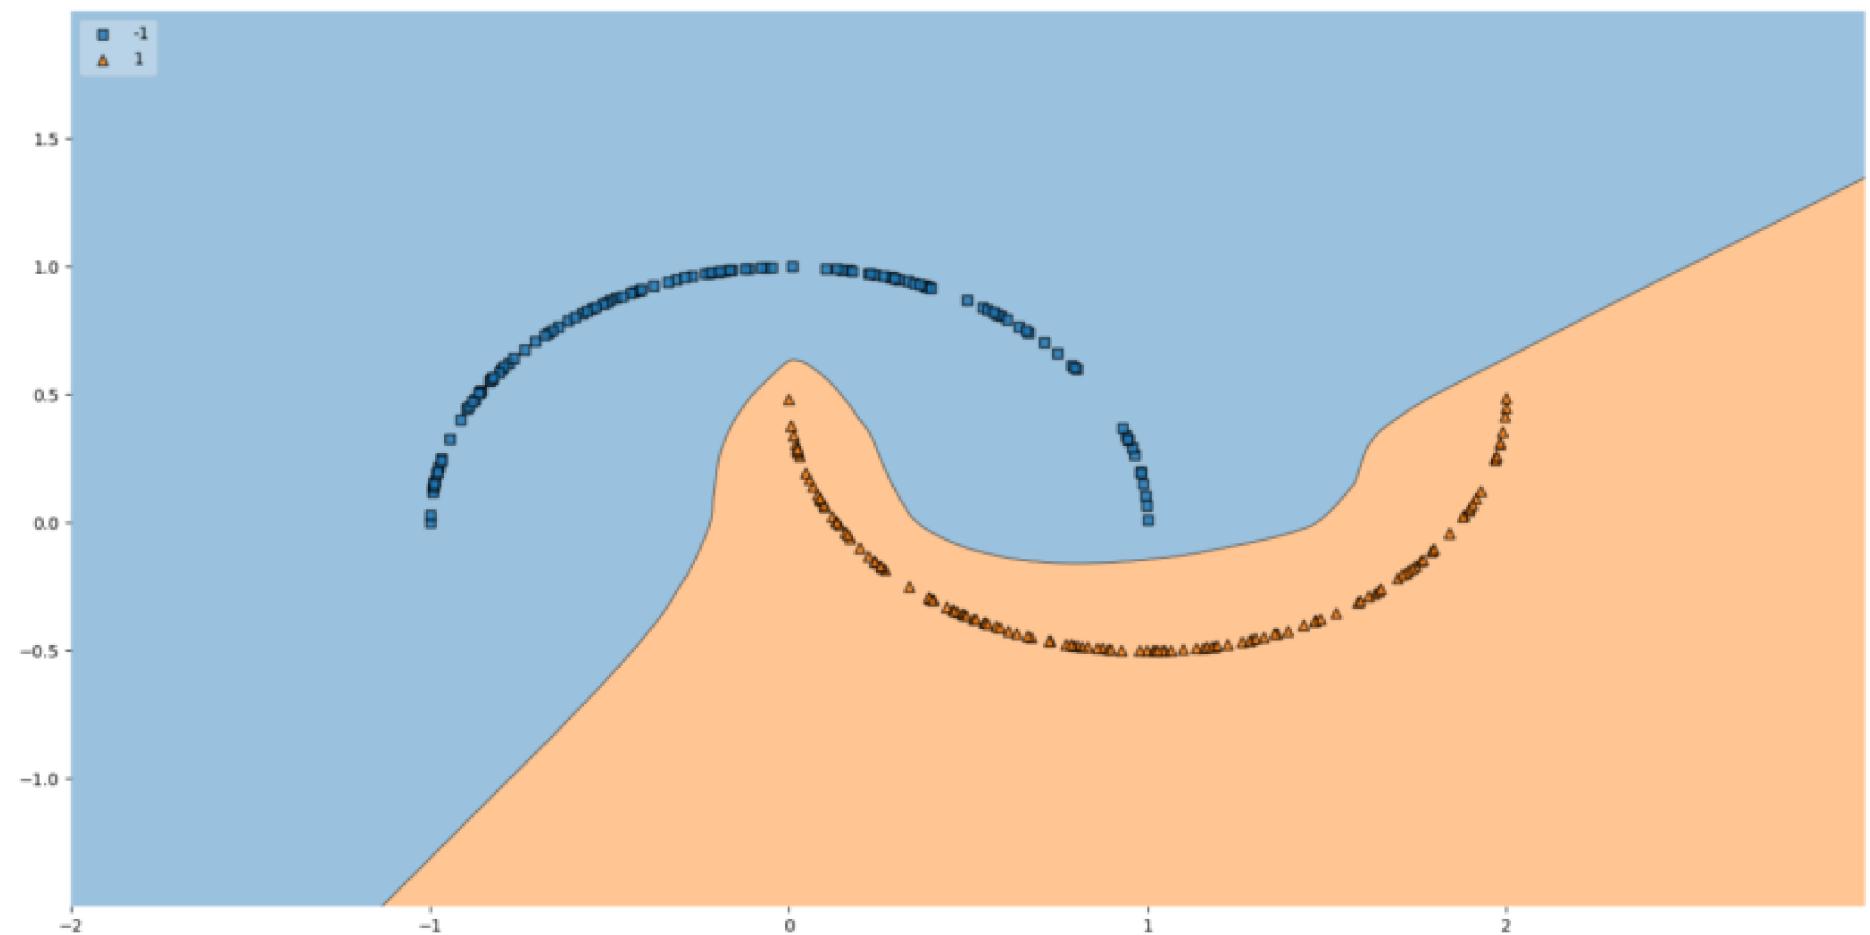
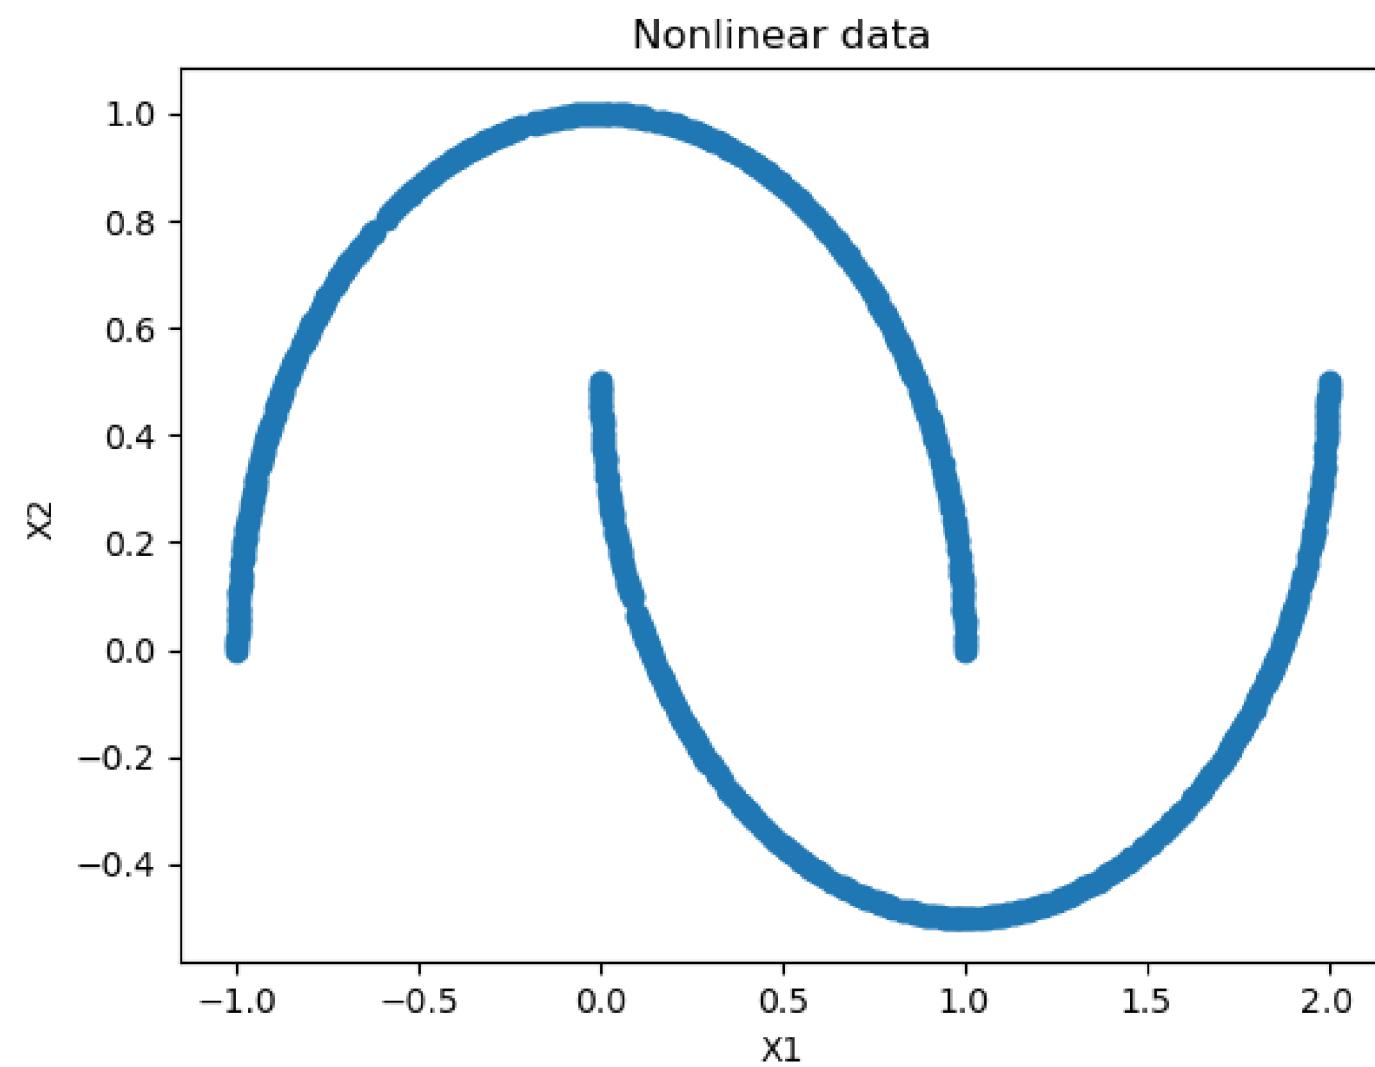
Activation Functions

if we do not use an activation function, we can't solve the problem.



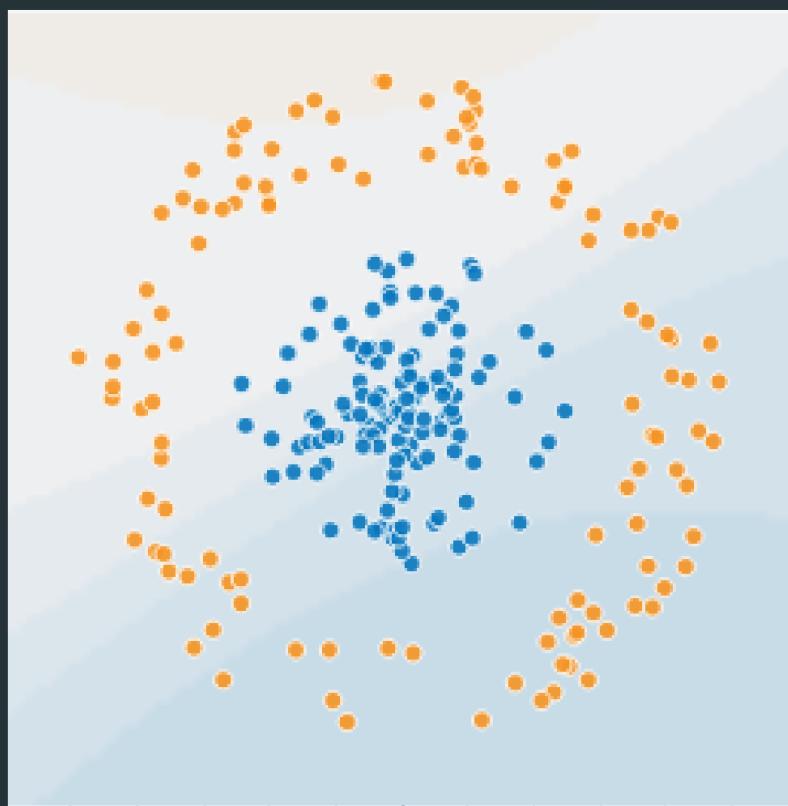
Activation Functions

but by introducing a sigmoid it is quite easy to solve it.



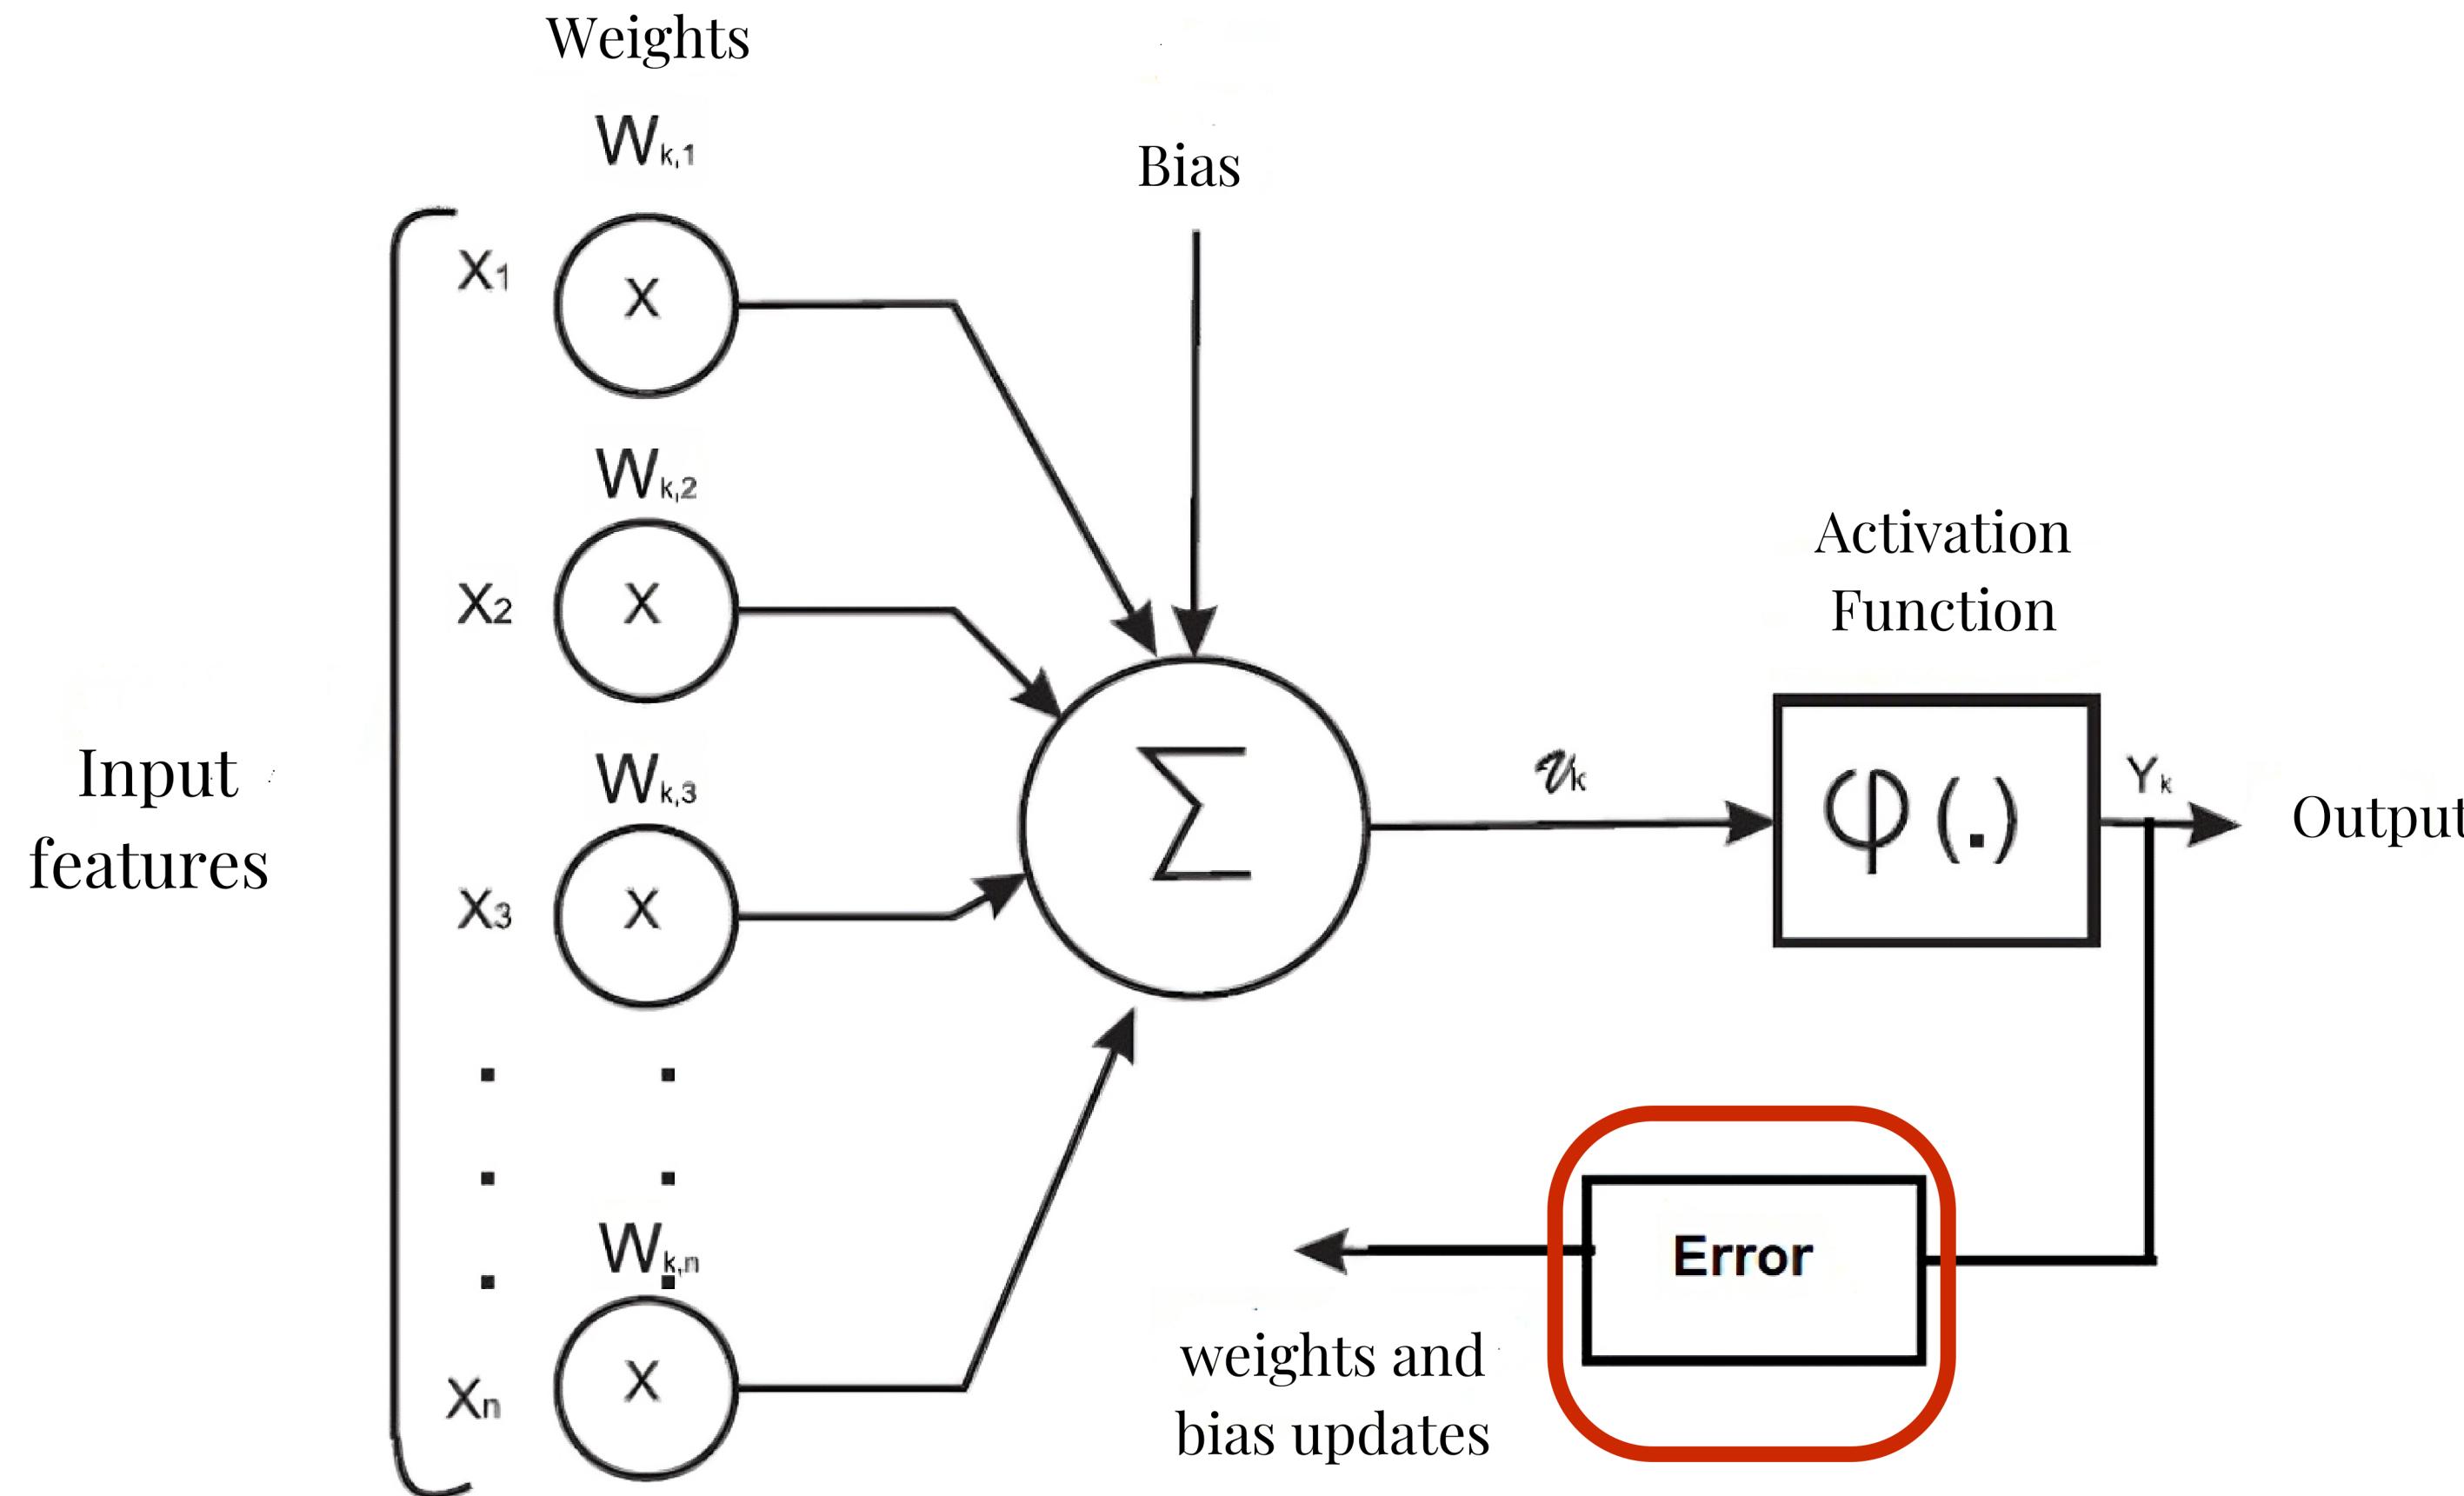
Hands-on

let's solve this problem:



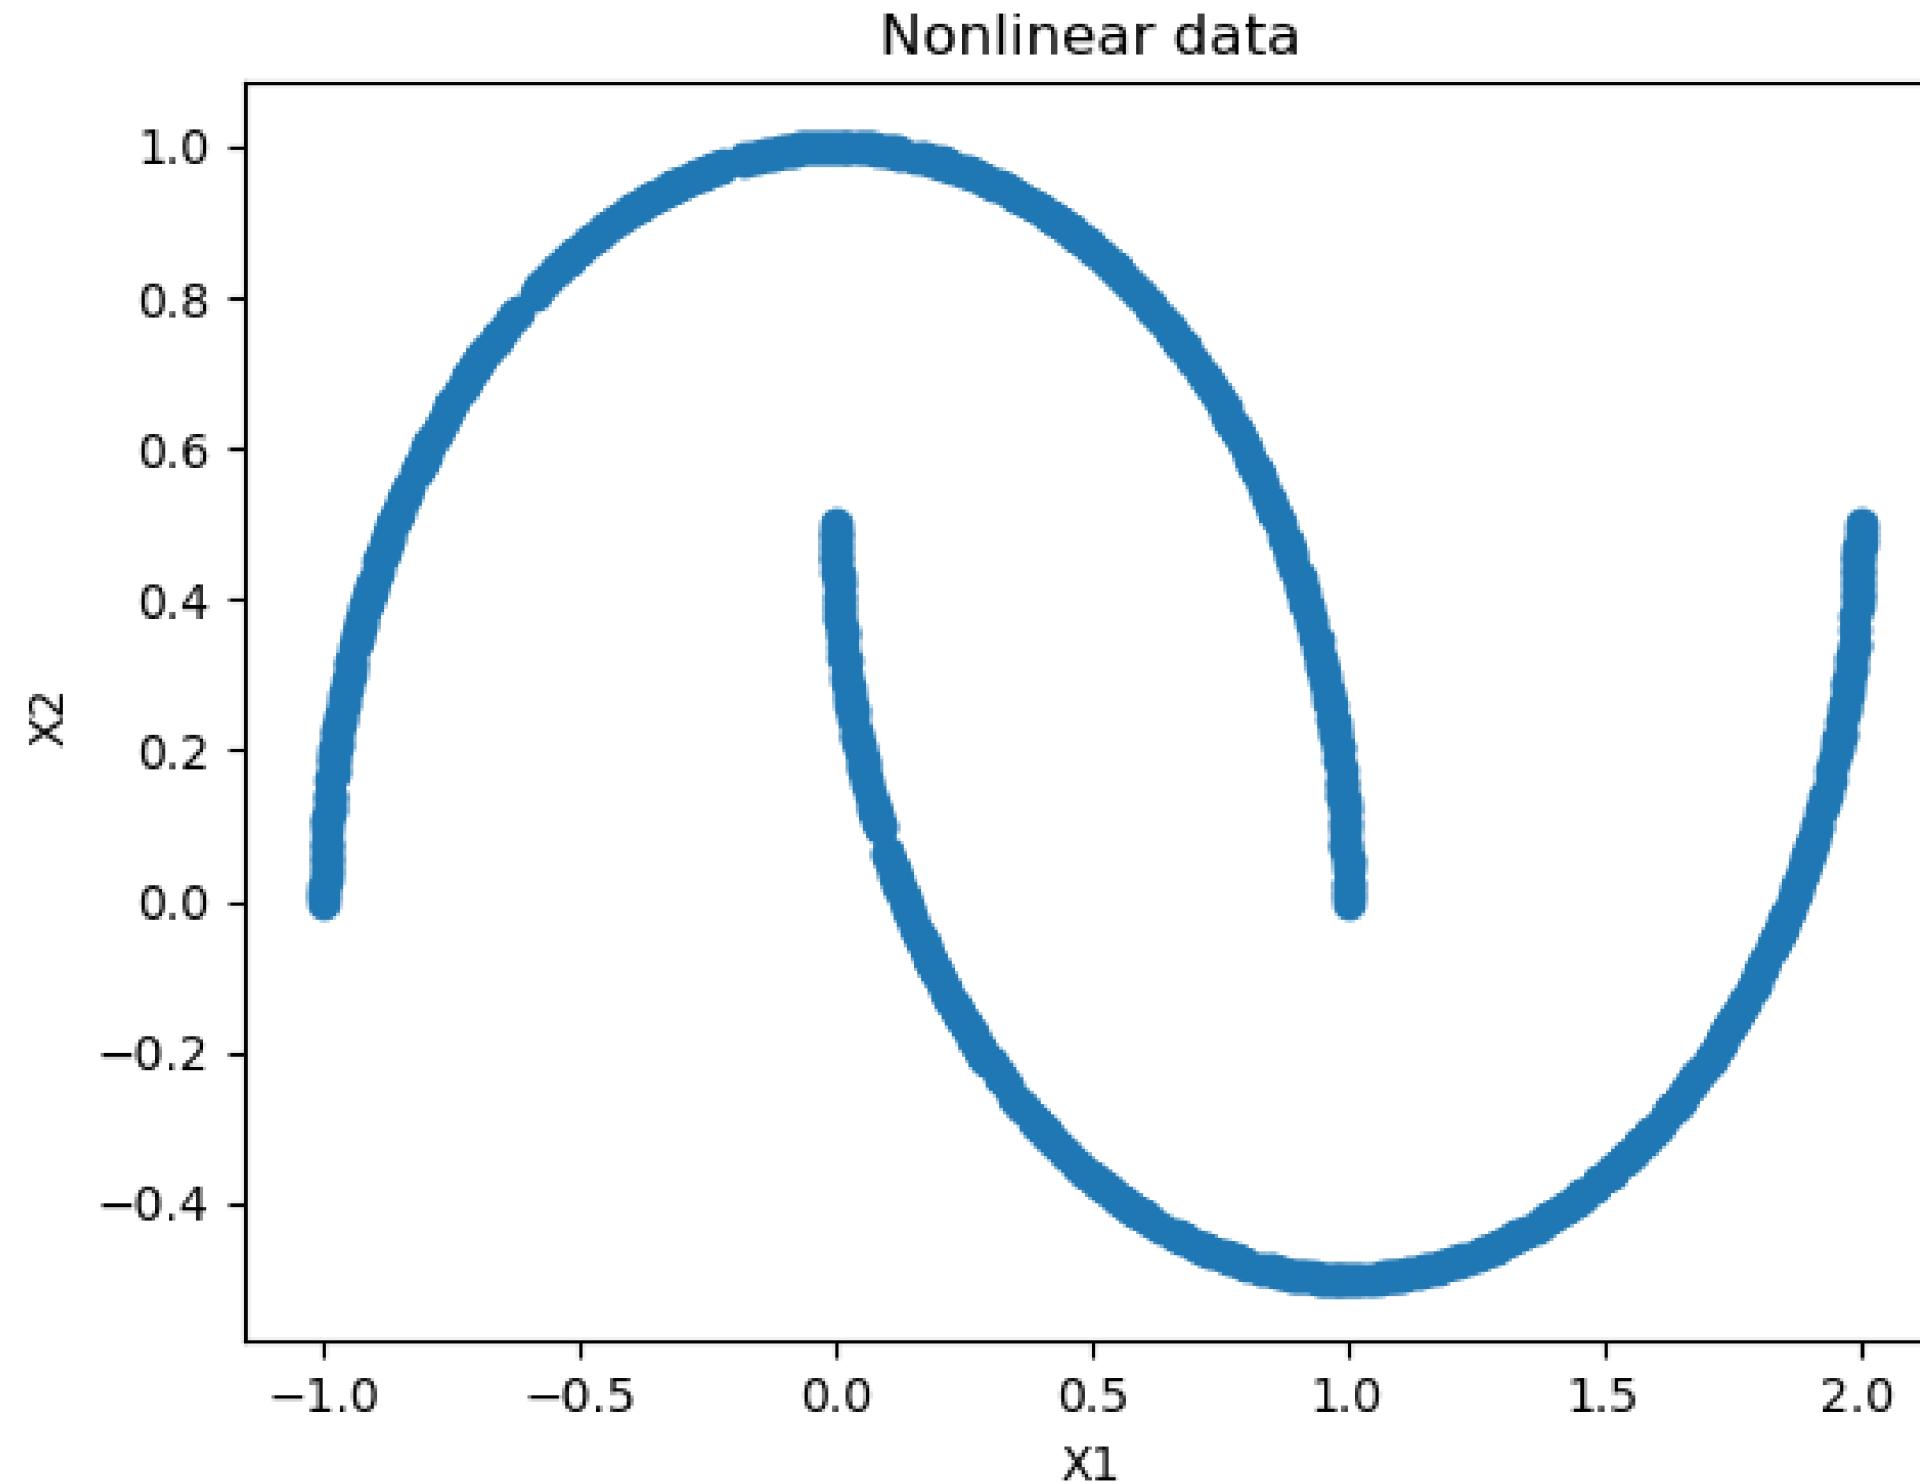
<http://playground.tensorflow.org/>

Diagram of a multilayer perceptron



Cost Functions

how to quantify the loss?



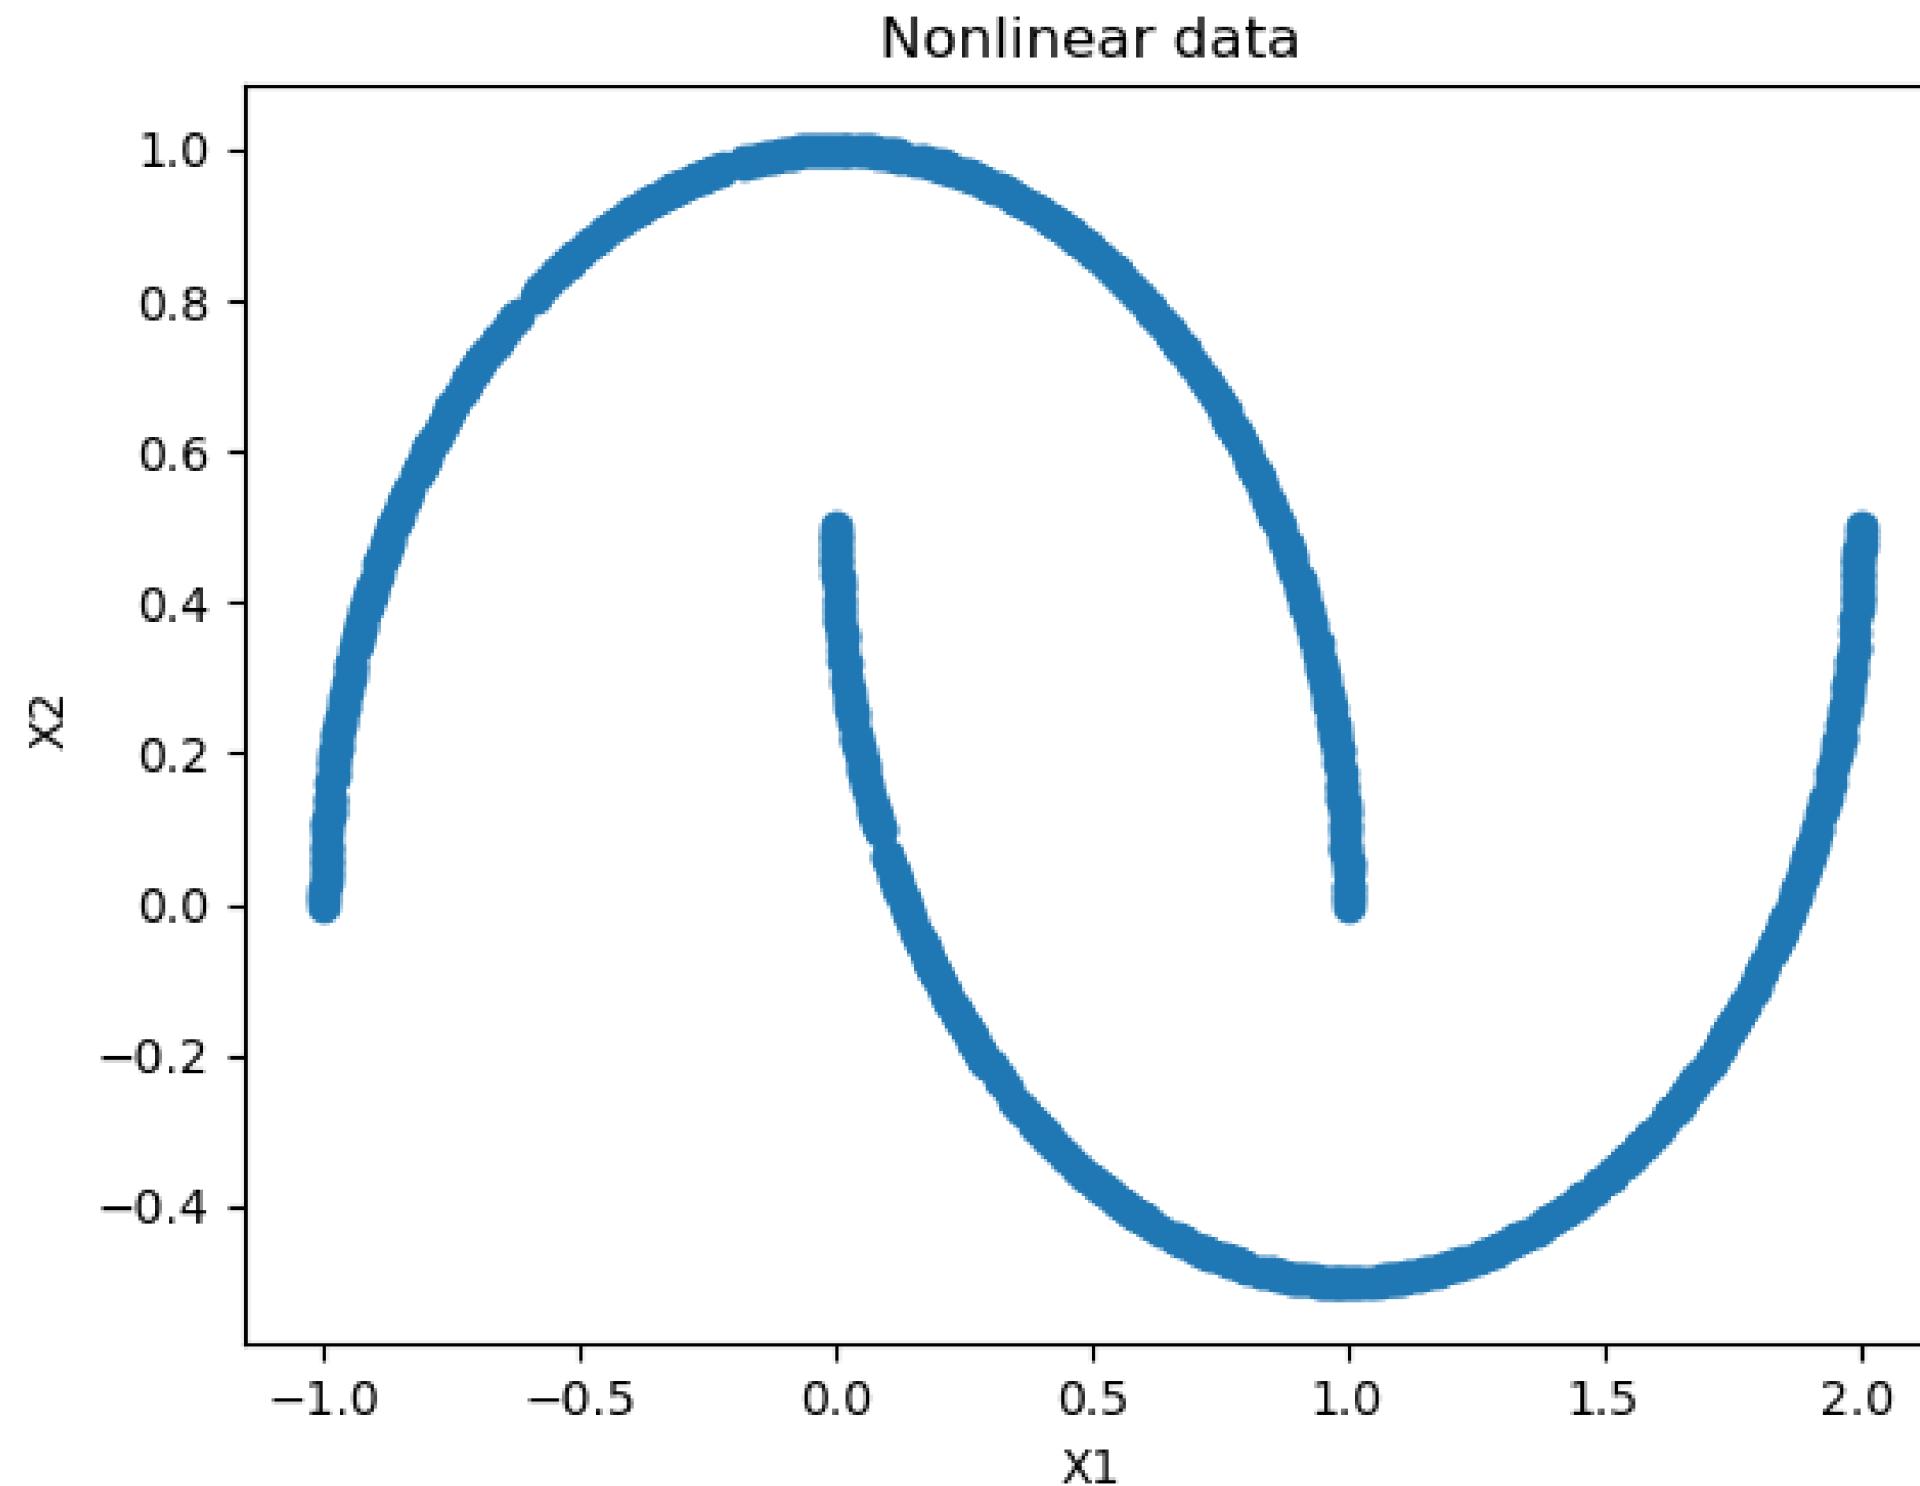
$$x^{(1)} = \begin{bmatrix} 2 \\ 0.4 \end{bmatrix}$$

$$\hat{y} = f(x^{(1)}) = 1$$

$$y = 0$$

Cost Functions

how to quantify the loss?



$$x^{(1)} = \begin{bmatrix} 2 \\ 0.4 \end{bmatrix}$$

$$\hat{y} = f(x^{(1)}) = 1$$

$$y = 0$$

Cost Functions

The loss a function:

$$\mathcal{L}(\hat{y}, y)$$

The function must be
derivable

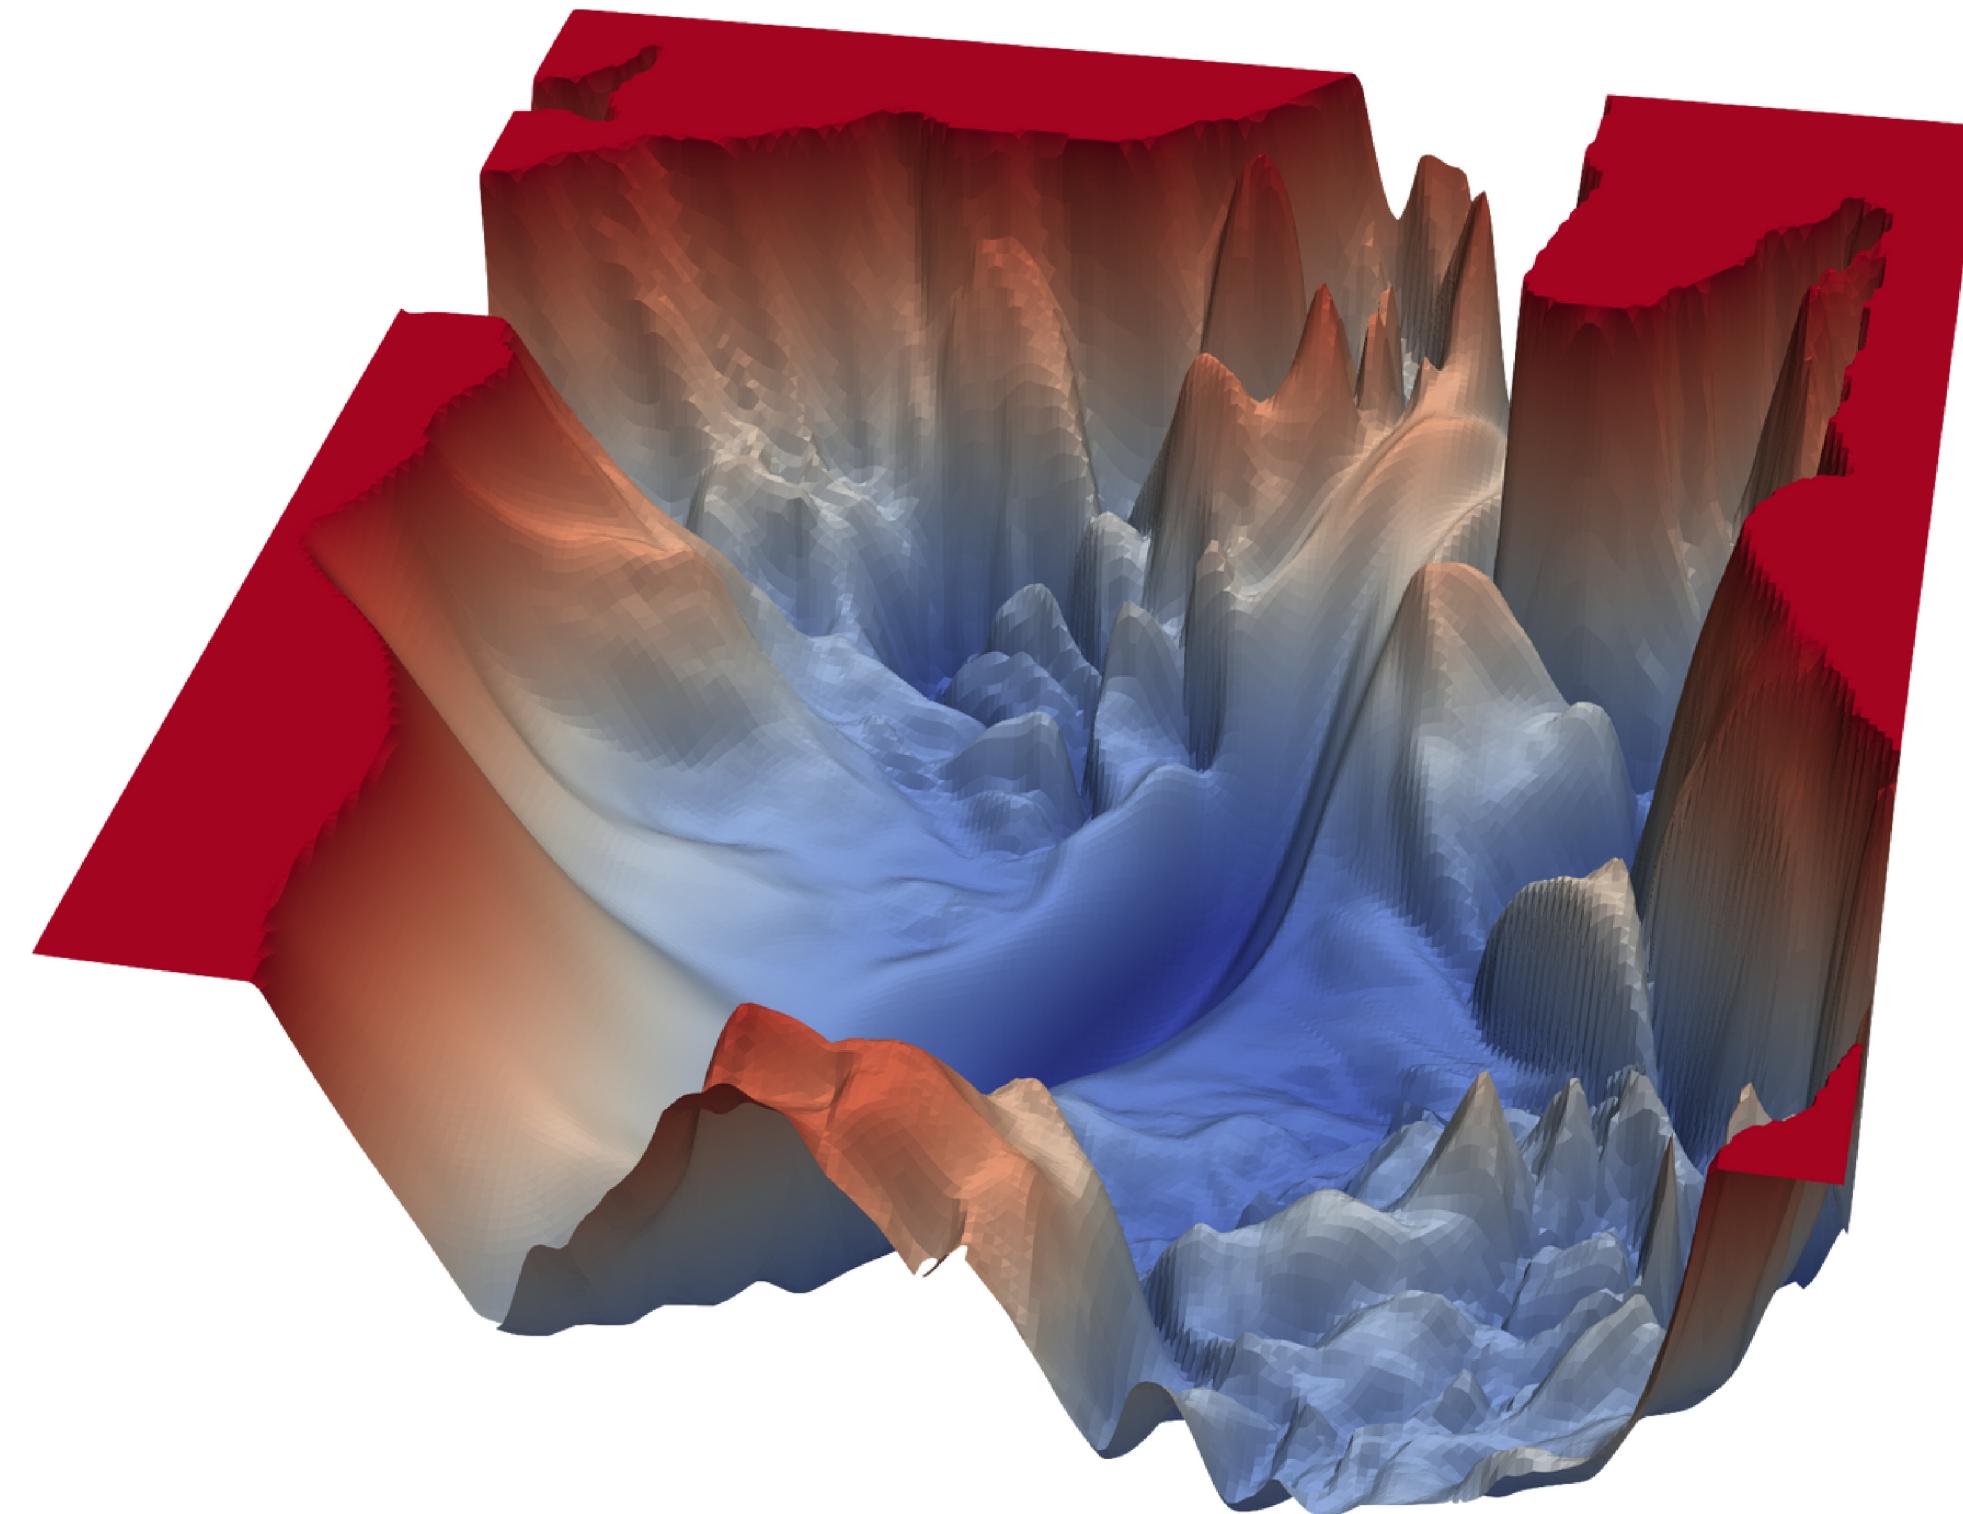
For a batch (a subset $\overline{\text{of data}}$) is expressed as

$$J(W) = \frac{1}{n} \sum_1^n \mathcal{L}(\hat{y}, y)$$



Loss function
Cost function
Objective function

Cost Functions



The Cost function of a VGG-56

Common Cost Functions

Binary Cross Entropy

$$-\frac{1}{N} \sum_{i=1}^N y_i \log(\hat{y}_i) + (1 - y_i) \log(1 - \hat{y}_i)$$

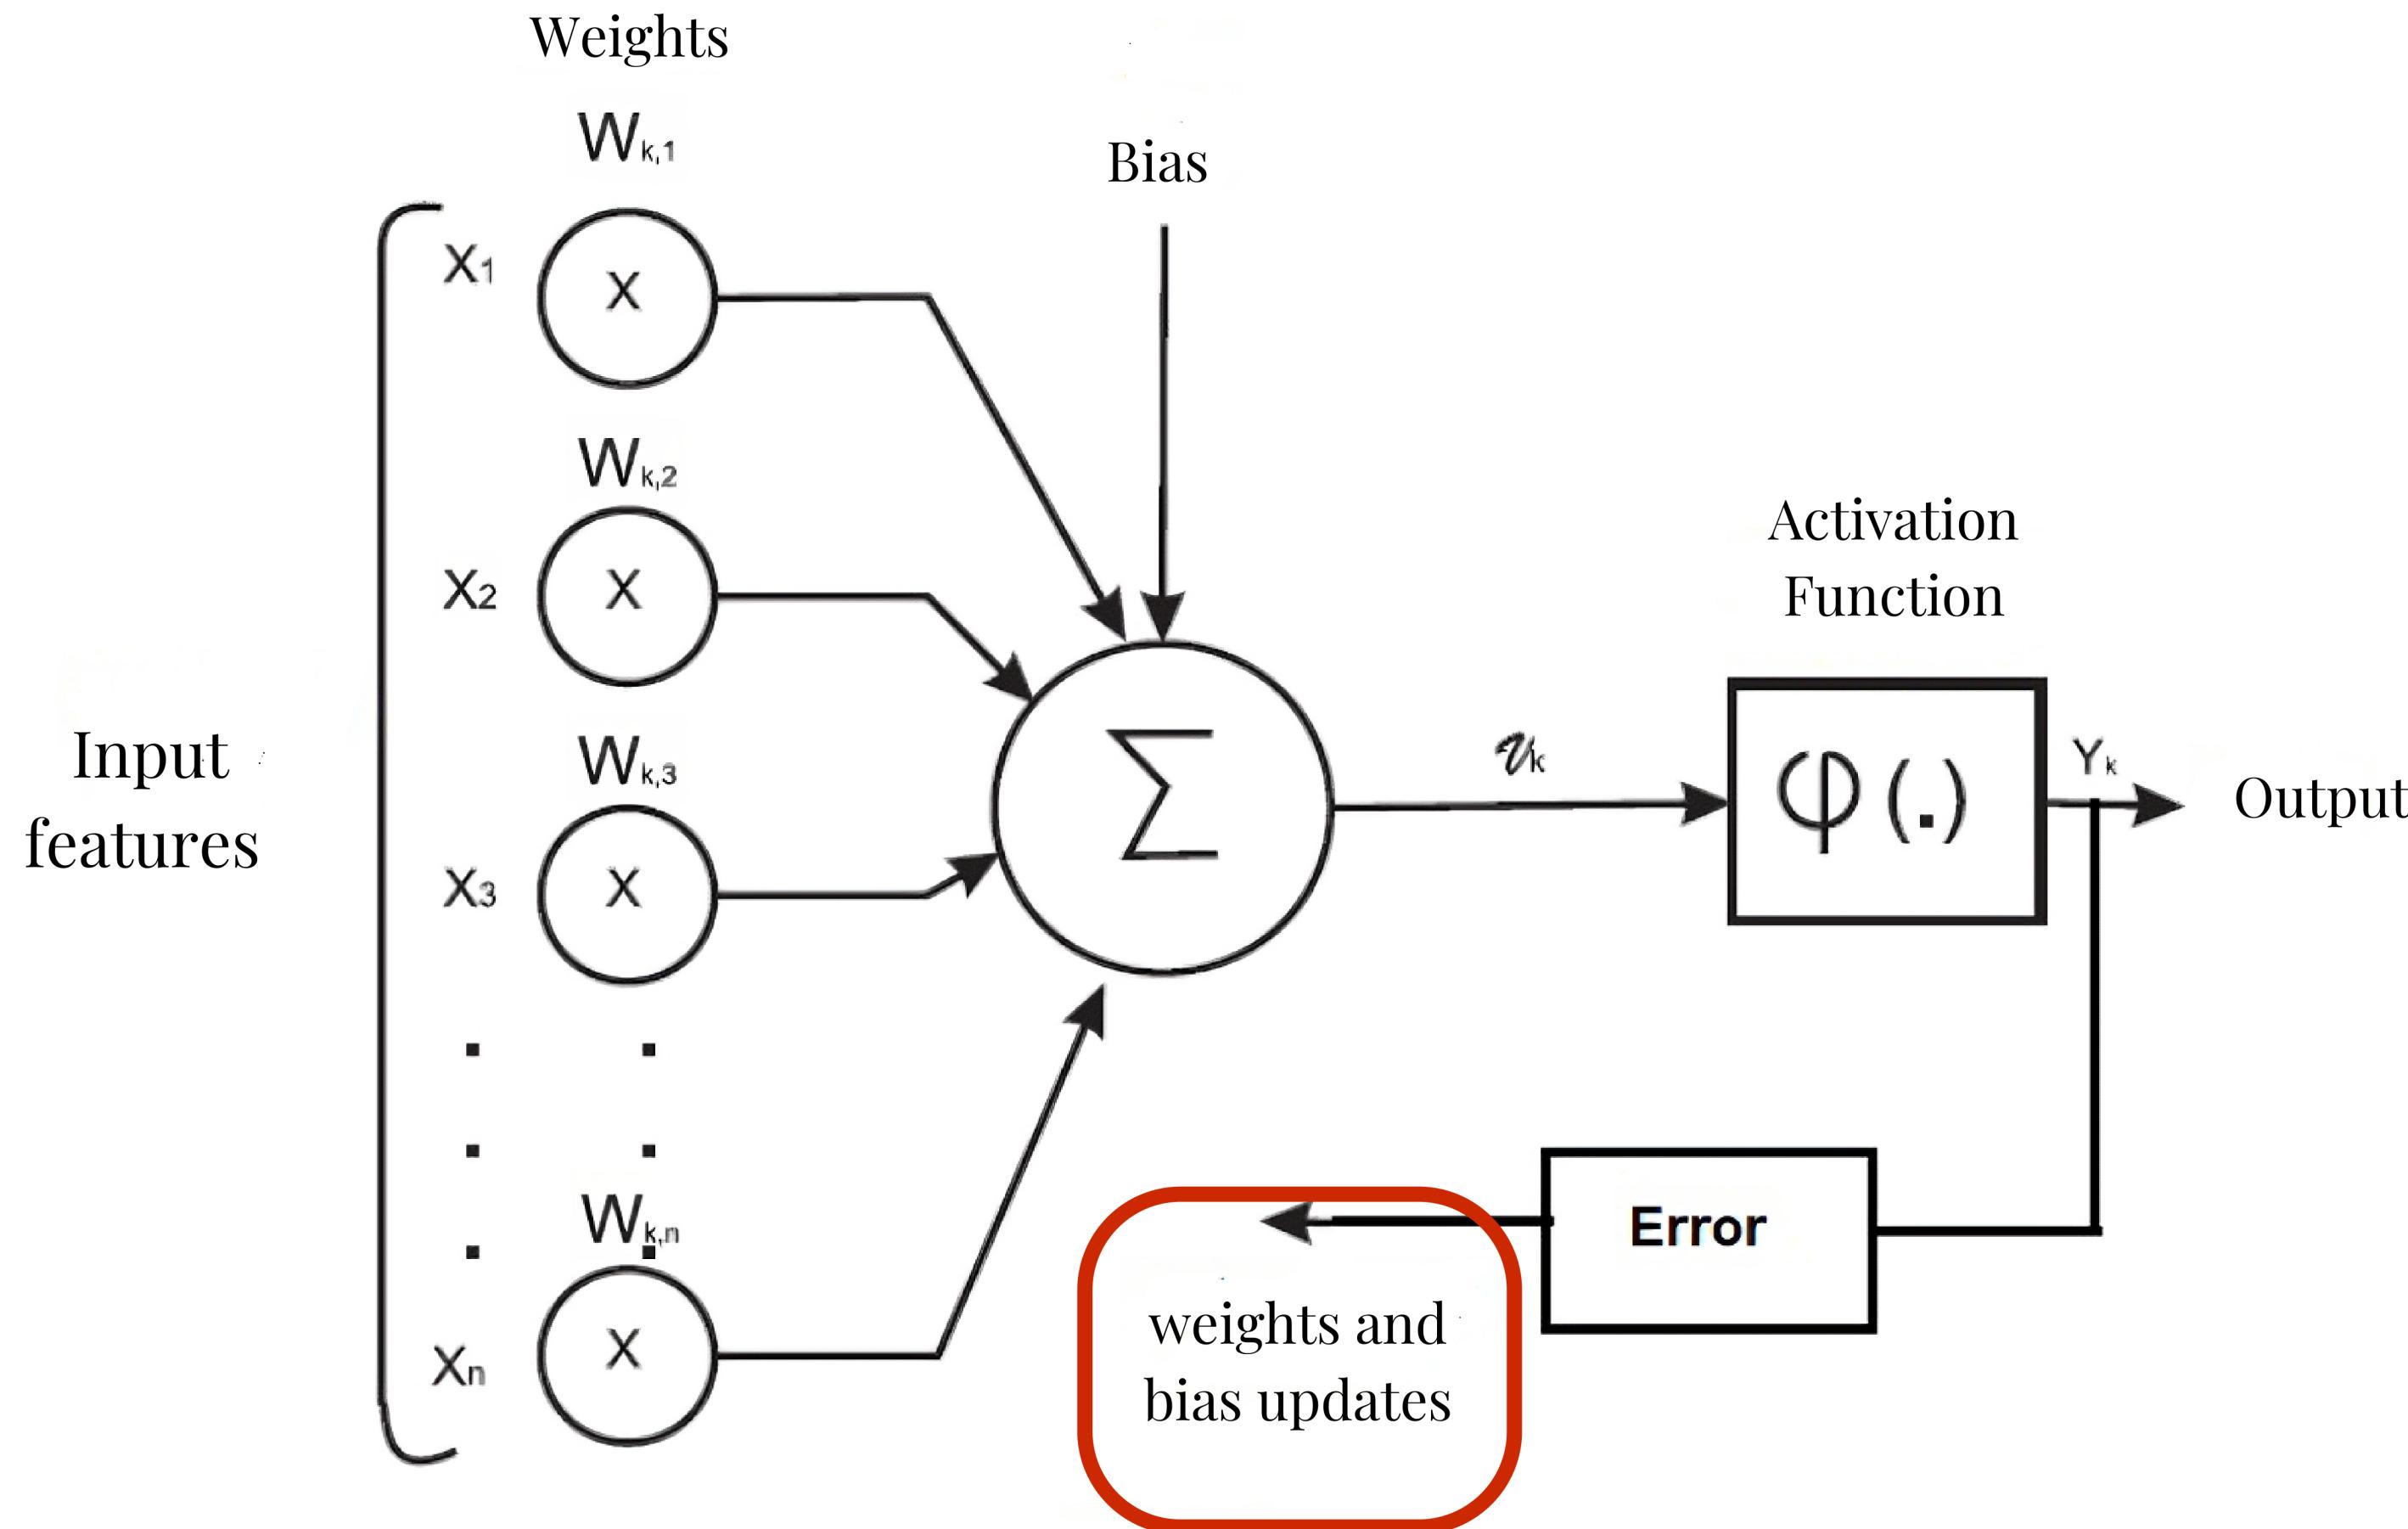
Is used in binary **classification**

Mean Squared Error

$$\frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

Is used in **regression**

Diagram of a multilayer perceptron



Backpropagation

we want to find the net weights at which that achieve the best performance and the lowest loss

$$W^* = \underset{W}{\operatorname{argmin}} f(W, \hat{y}, y)$$



we want to minimize this

Backpropagation

we want to find the net weights at which that achieve the best performance and the lowest loss

$$W^* = \operatorname{argmin}_W \mathcal{L}(W)$$



we want to minimize this

to find the optimal weights

optimizers: sgd

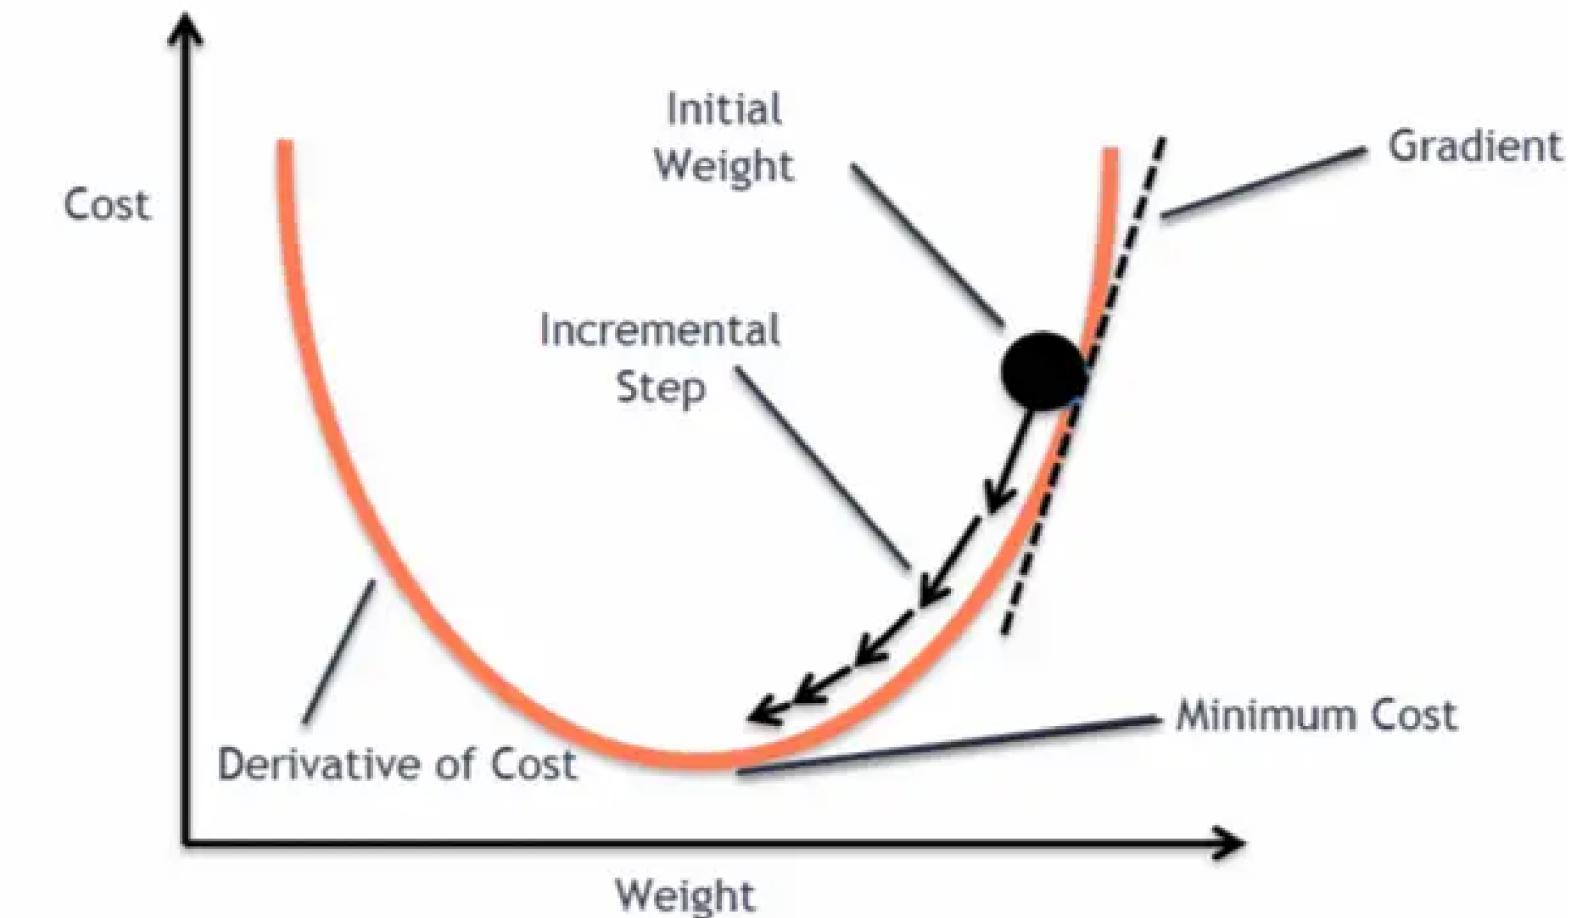
we start with some random weights

and let's find the derivative of

$$\nabla_W = \frac{\partial \mathcal{L}(W)}{\partial W}$$



this is the
gradient



optimizers: sgd

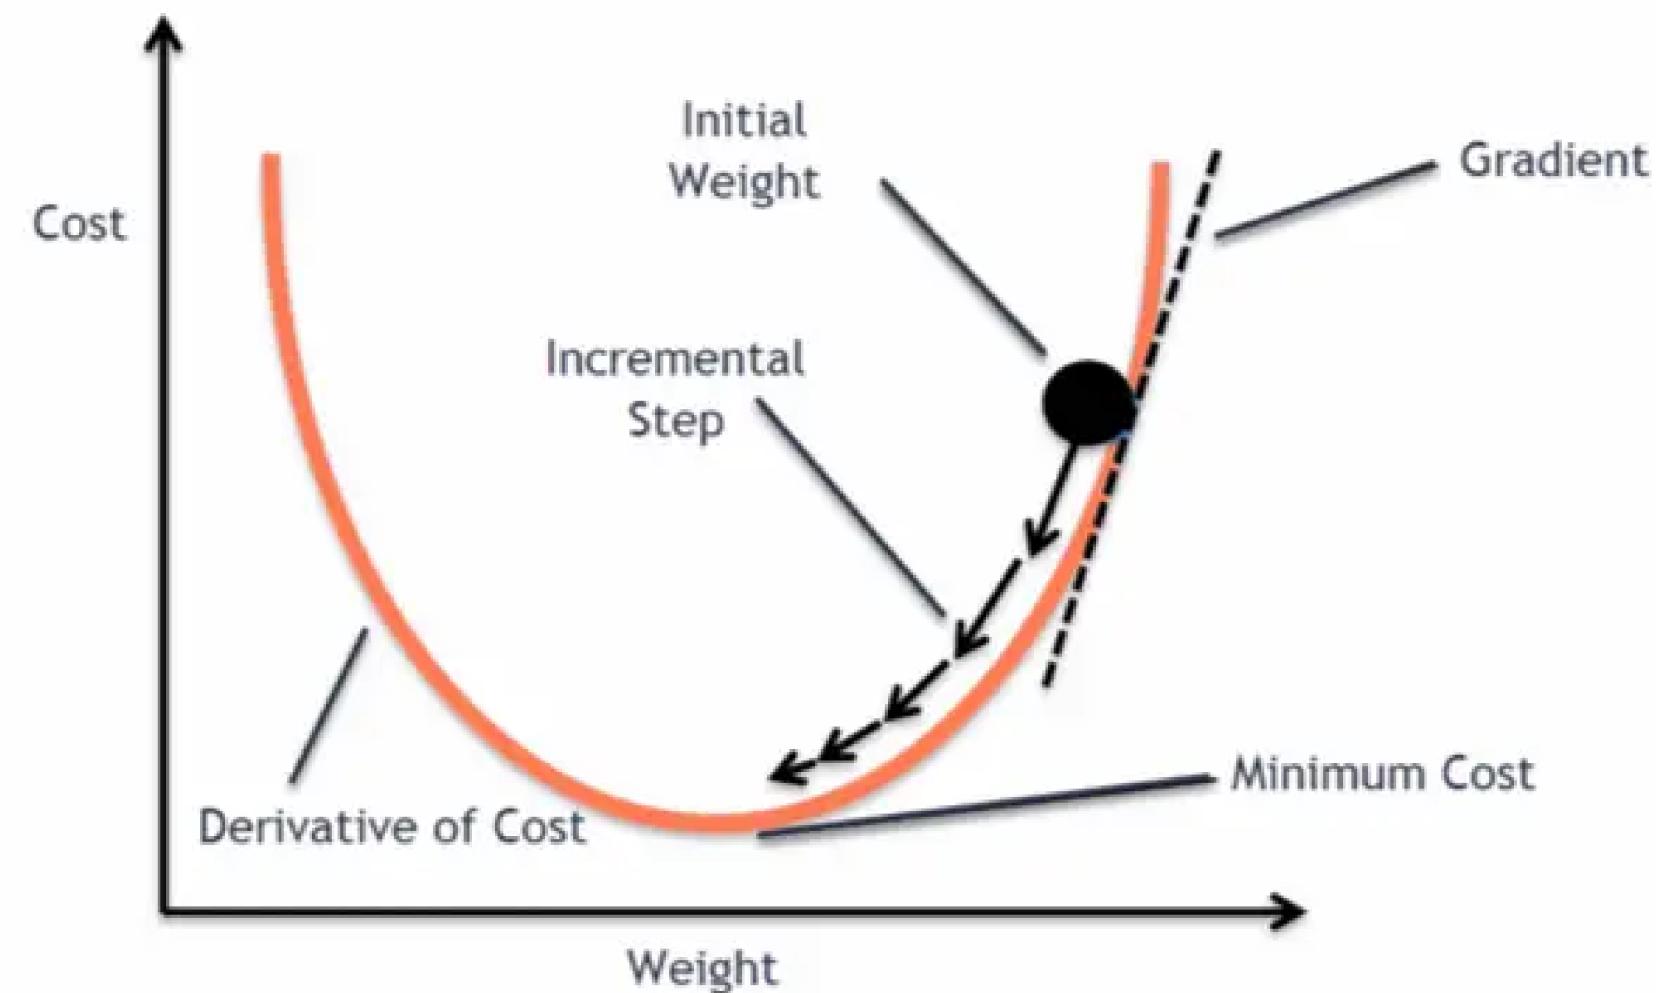
we start with some random weights

and let's find the derivative of

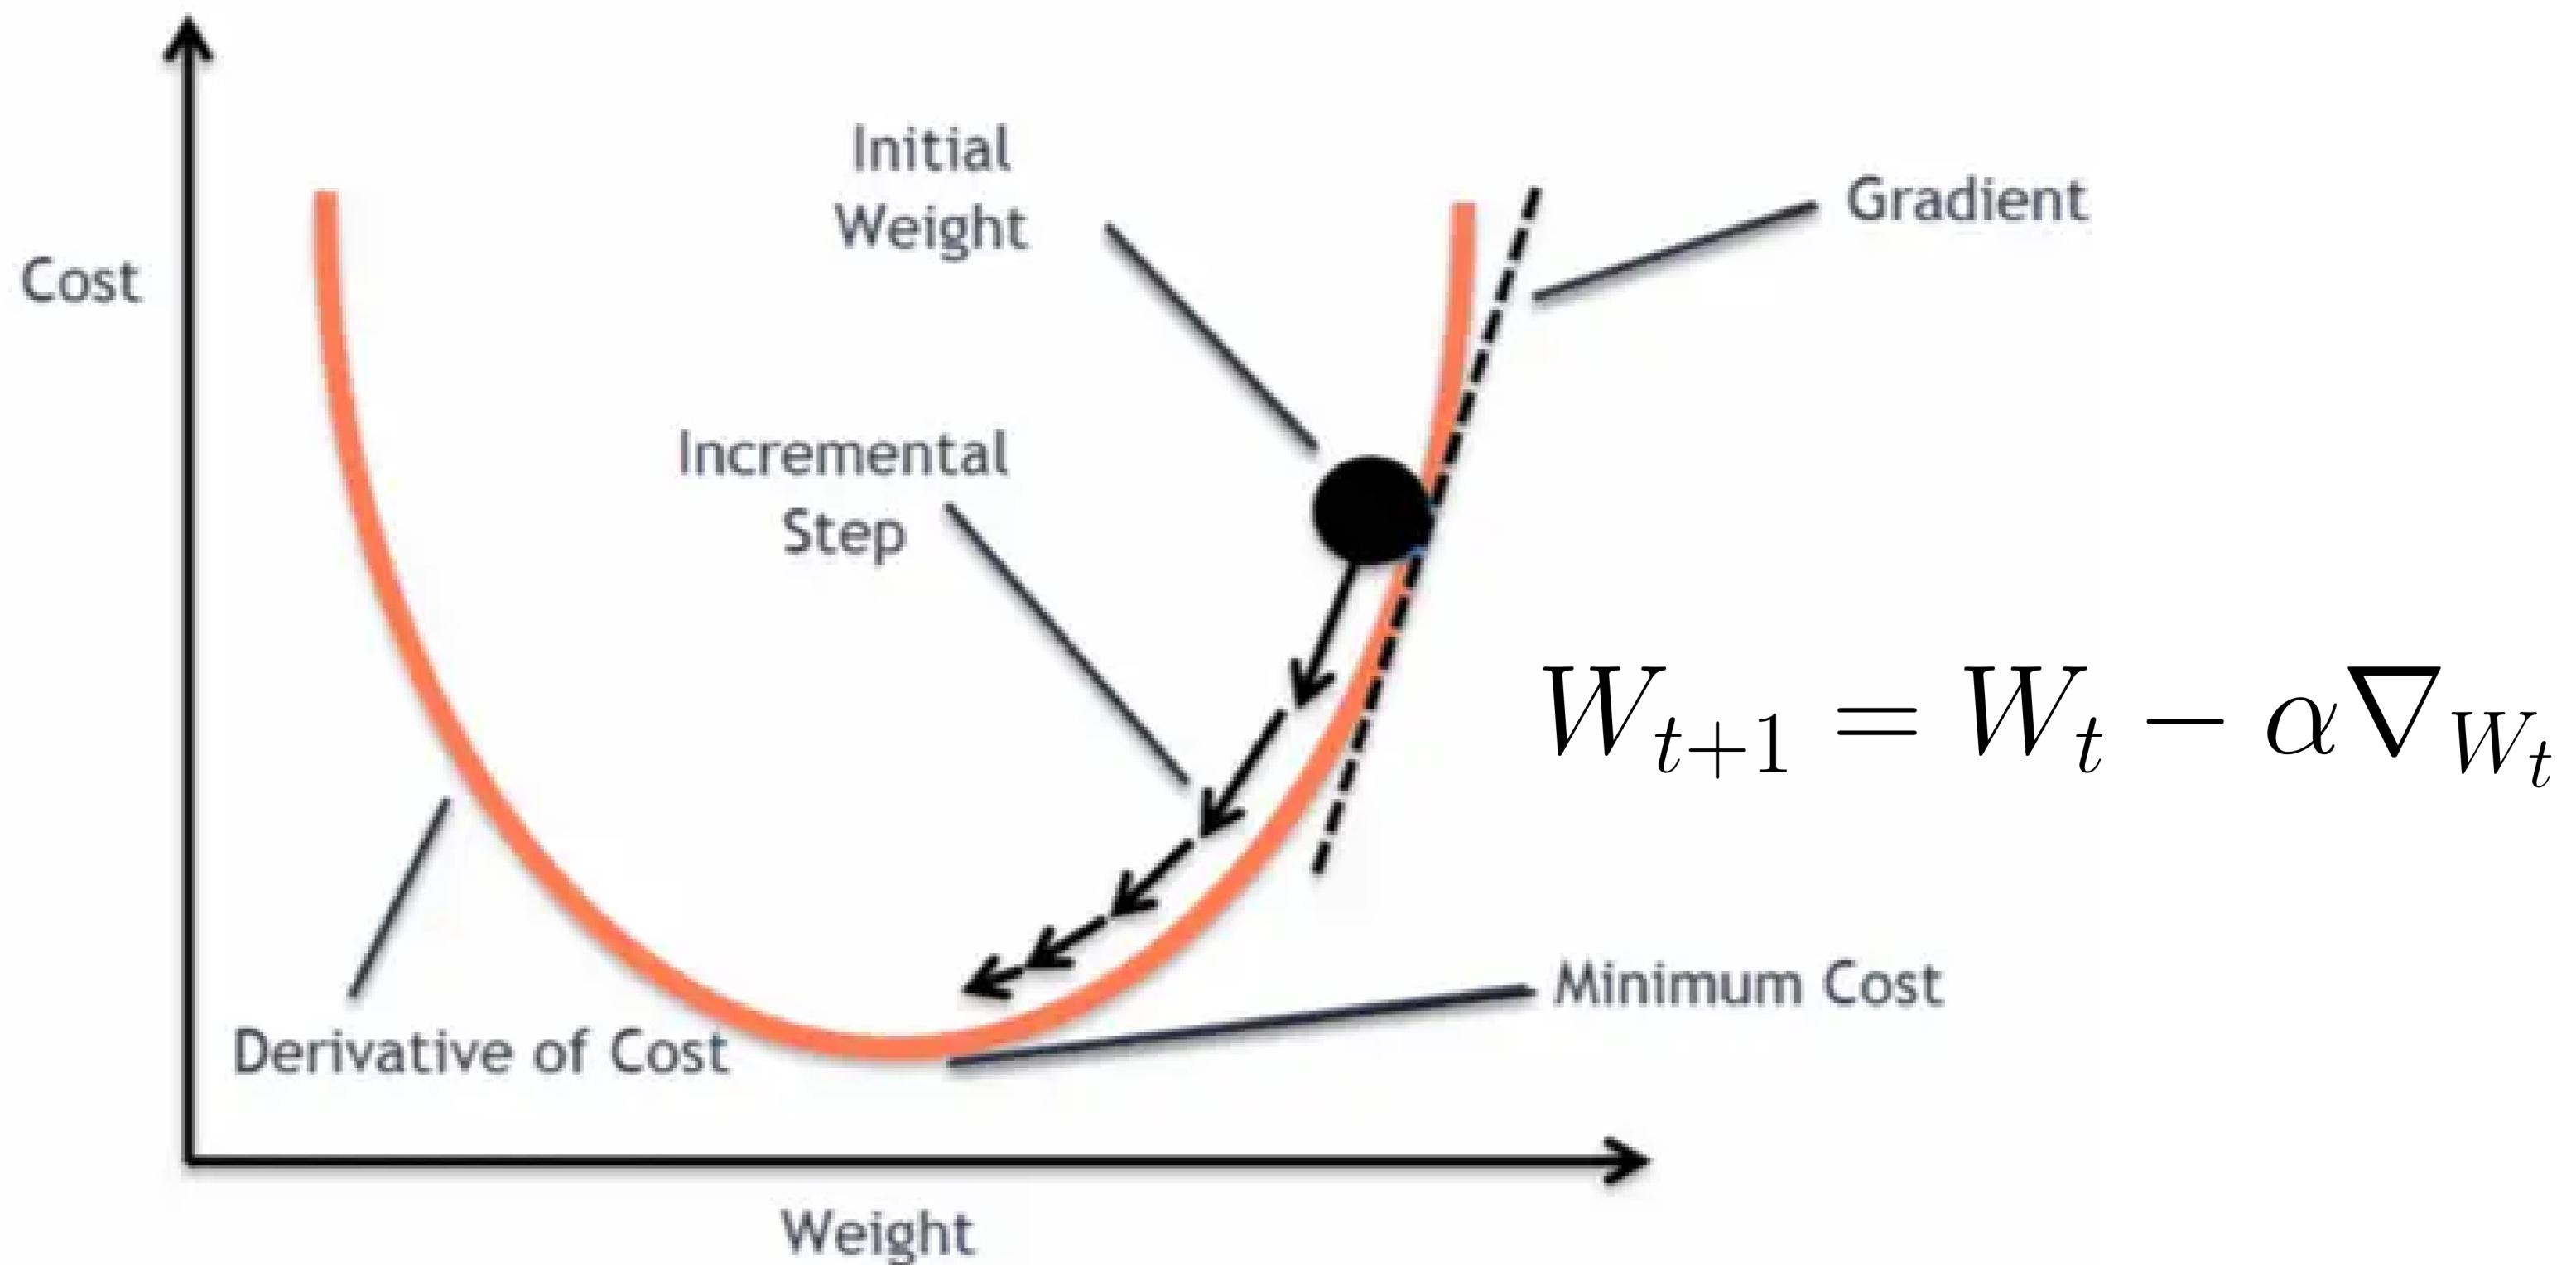
$$\nabla_W = \frac{\partial \mathcal{L}(W)}{\partial W}$$

update weights with this rule:

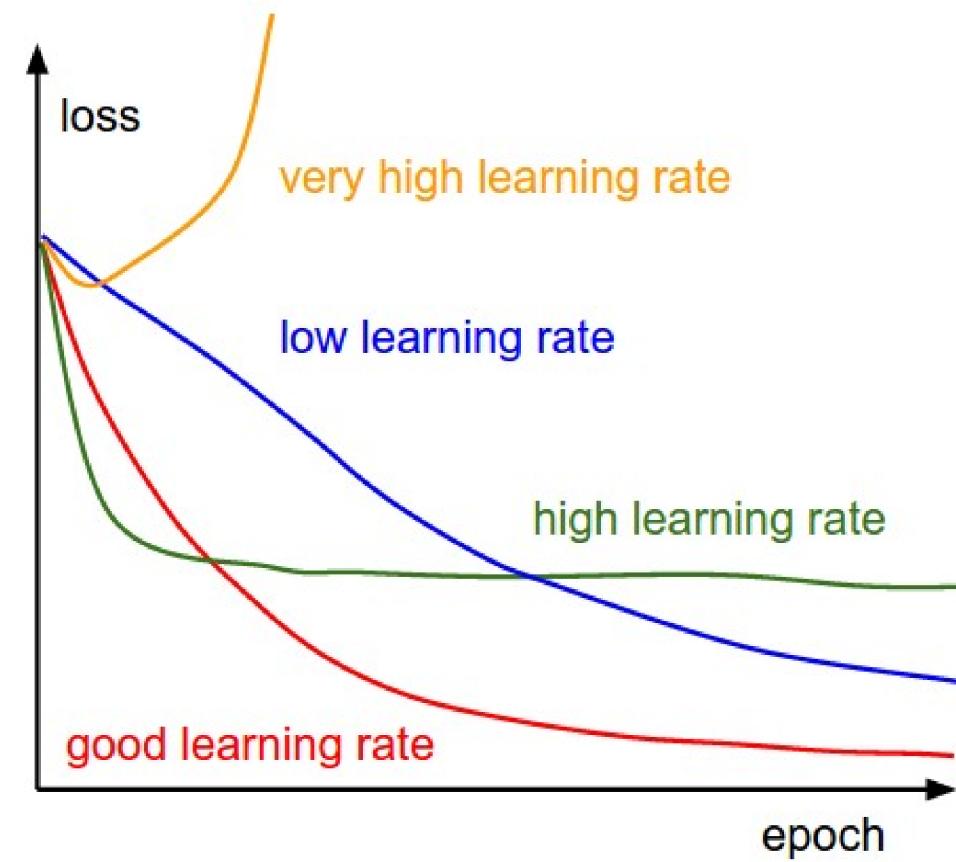
$$W_{t+1} = W_t - \alpha \nabla_{W_t}$$



optimizers: sgd



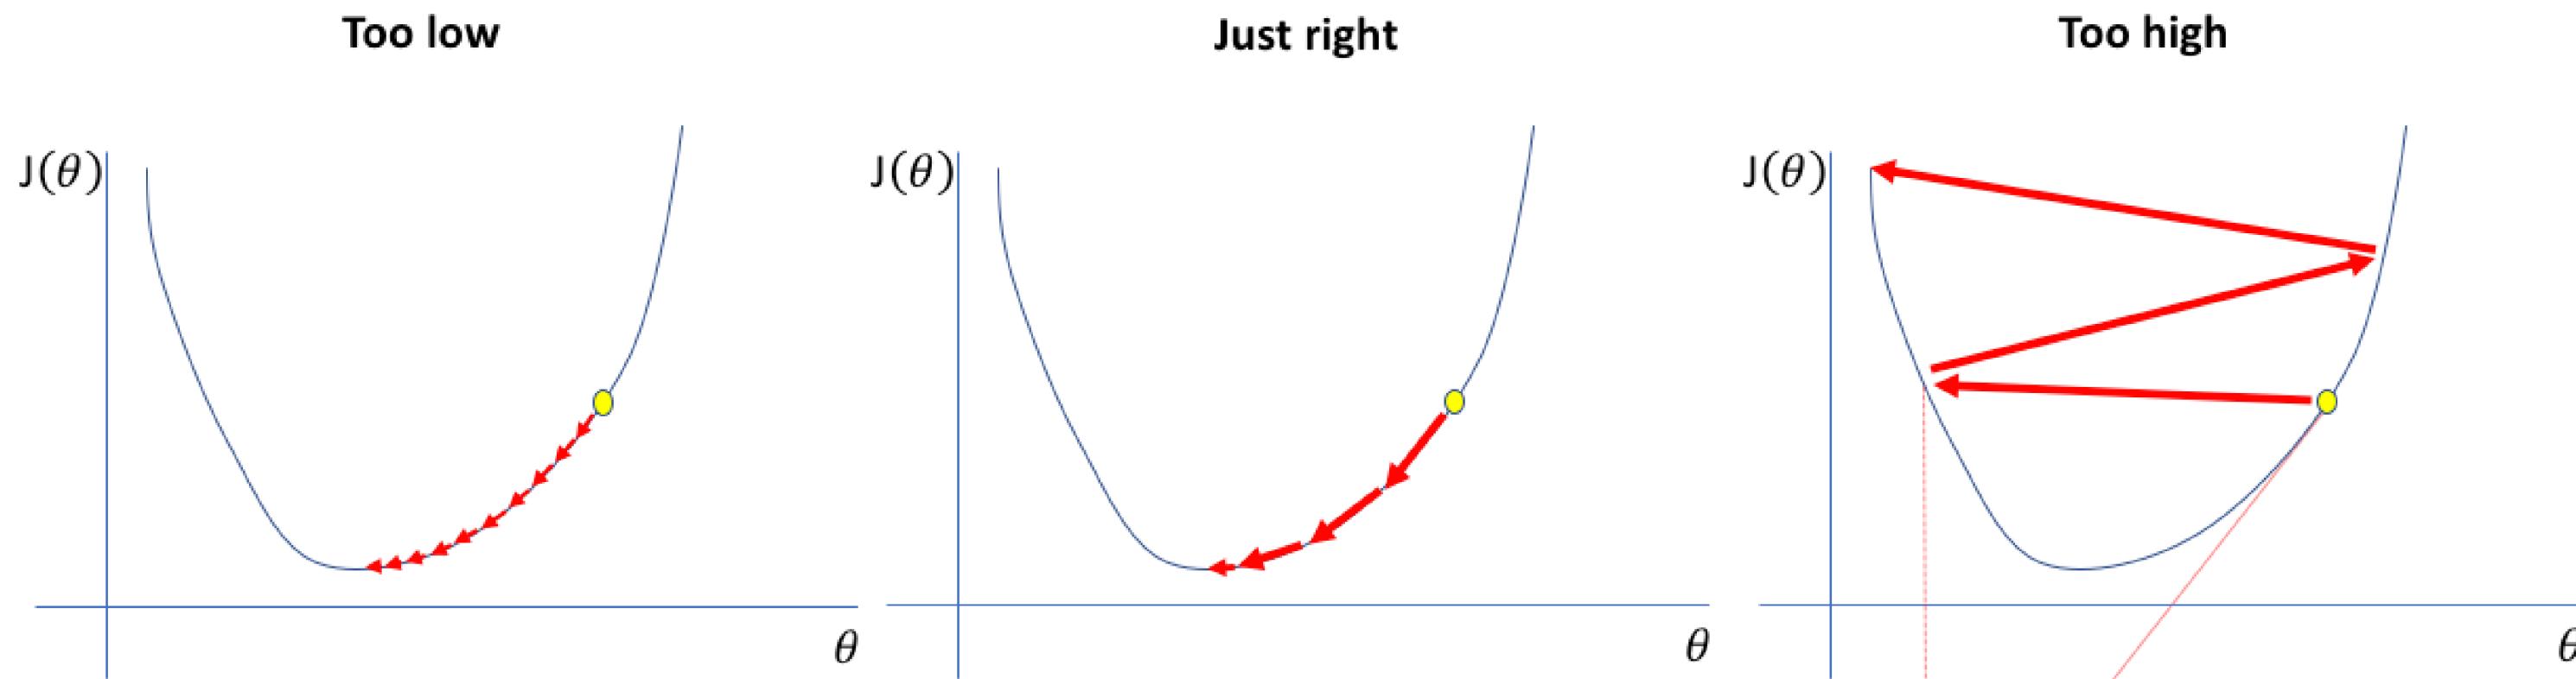
optimizers: learning rate



$$W_{t+1} = W_t - \alpha \nabla W_t$$



$\alpha = \text{learning rate}$



Metrics

metrics are a way to evaluate performance

are used in the test of the model

some common metrics for classification are:

Accuracy

$$\frac{T_p + T_n}{T_p + T_n + F_p + F_n}$$

Precision

$$\frac{T_p}{T_p + F_p}$$

Recall

$$\frac{T_p}{T_p + T_n}$$

Hands-on in colab

Colab link

discussion of results

Thank u <3